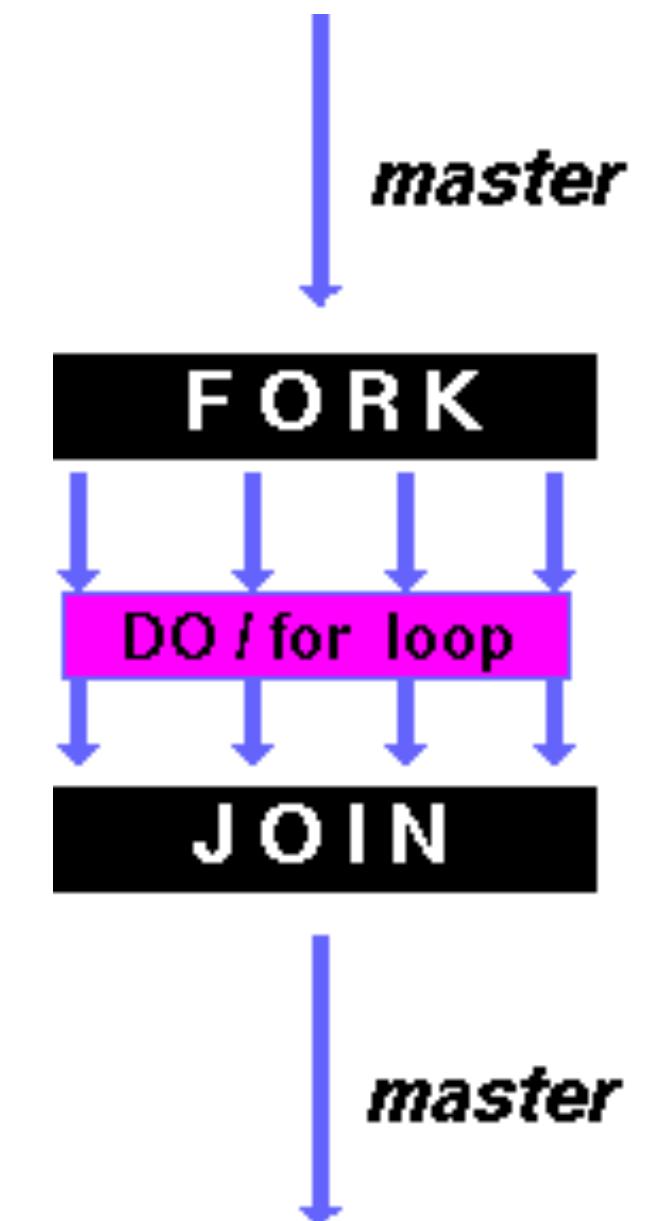


# More On Synchronization And Threading

# OpenMP parallel for pragma

- ```
#pragma omp parallel for
for (i=0; i<max; i++) zero[i] = 0;
```
- Master thread creates additional threads, each with a separate execution context
- All variables declared outside for loop are shared by default, except for loop index which is *implicitly* private per thread
- Implicit “barrier” synchronization at end of for loop
- Divide index regions sequentially per thread
  - Thread 0 gets 0, 1, ..., (max/n)-1;
  - Thread 1 gets max/n, max/n+1, ..., 2\*(max/n)-1

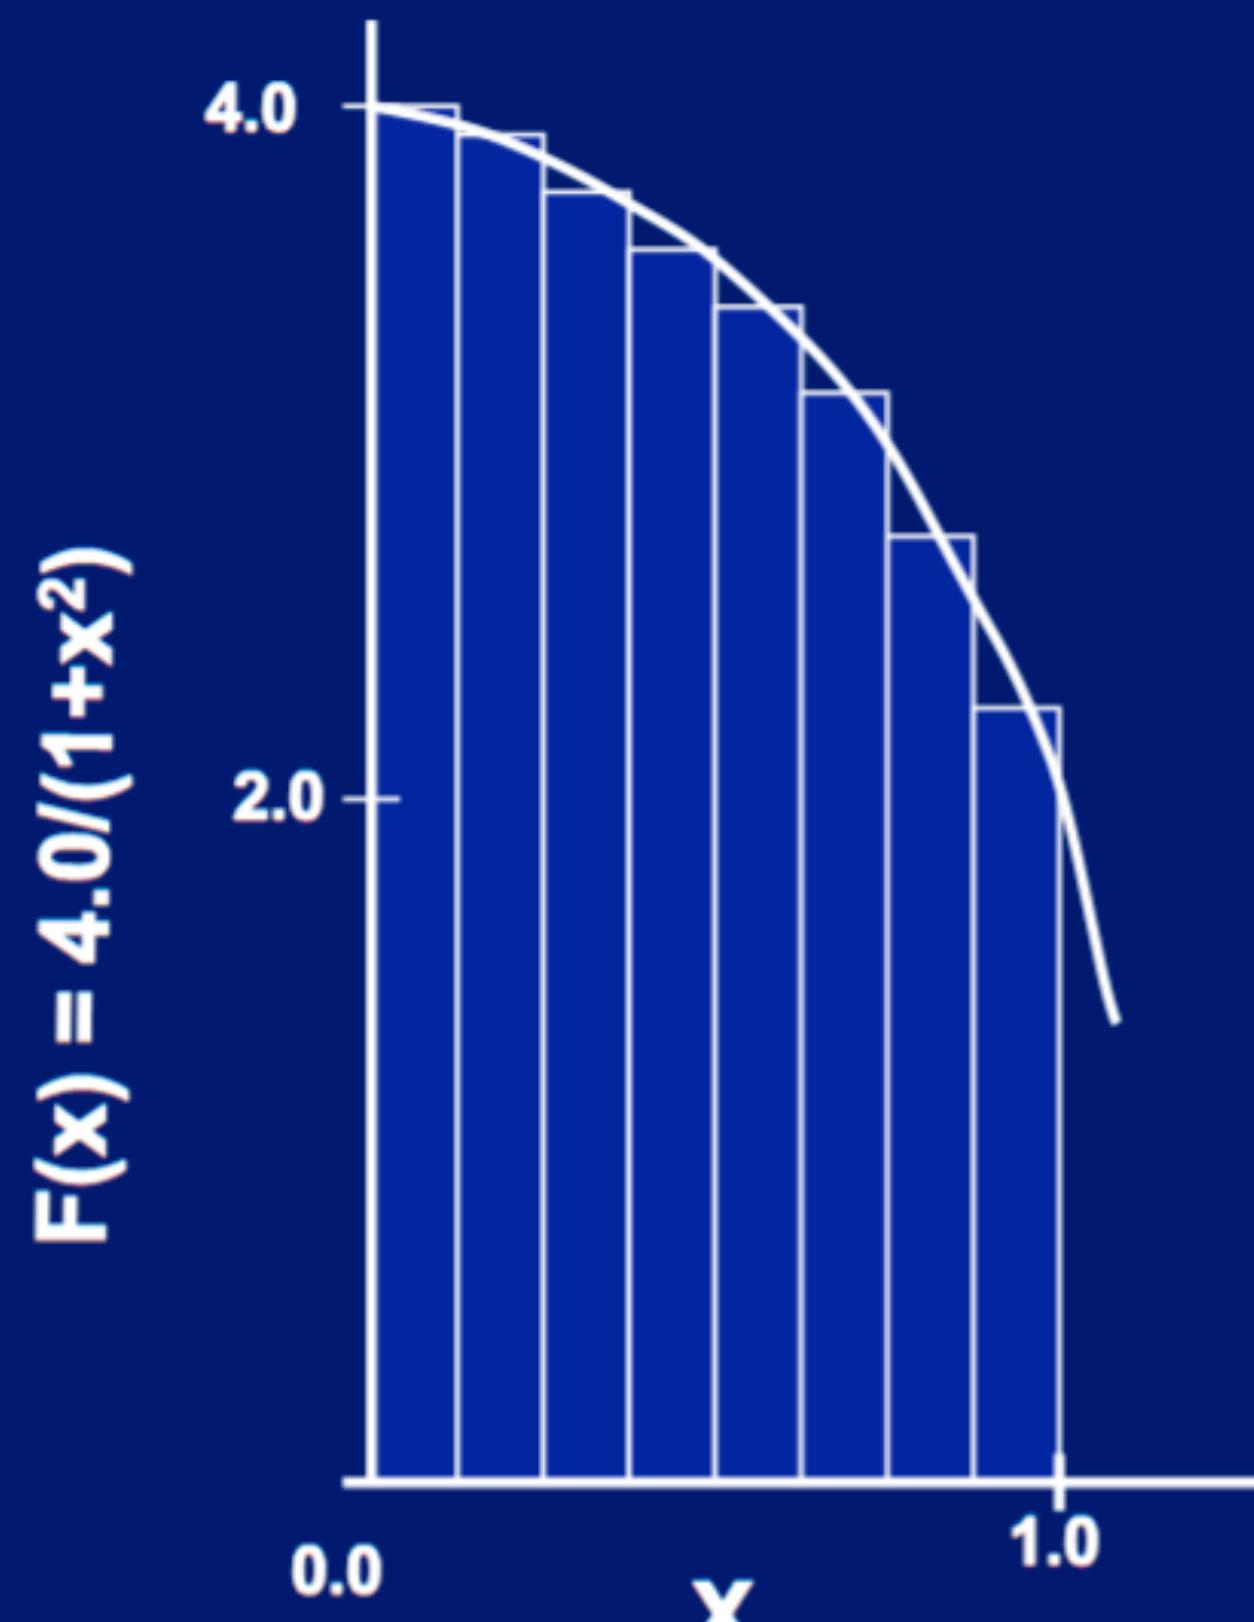


# Example 2: Computing $\pi$

Computer Science 61C

McMahon & Weaver

## Numerical Integration



Mathematically, we know that:

$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

We can approximate the integral as a sum of rectangles:

$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$

Where each rectangle has width  $\Delta x$  and height  $F(x_i)$  at the middle of interval  $i$ .

<http://openmp.org/mp-documents/omp-hands-on-SC08.pdf>

# Working Parallel π without a for loop

```
#include <stdio.h>
#include <omp.h>

Com McMahon & Weaver

void main () {
    const int NUM_THREADS = 4;
    const long num_steps = 10;
    double step = 1.0/((double)num_steps);
    double sum[NUM_THREADS];
    for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
    omp_set_num_threads(NUM_THREADS);

#pragma omp parallel
{
    int id = omp_get_thread_num();
    for (int i=id; i<num_steps; i+=NUM_THREADS) {
        double x = (i+0.5) *step;
        sum[id] += 4.0*step/(1.0+x*x);
        printf("i =%3d,  id =%3d\n", i, id);
    }
}
double pi = 0;
for (int i=0; i<NUM_THREADS; i++) pi += sum[i];
printf ("pi = %6.12f\n", pi);
```

# Trial Run

```
#include <stdio.h>
#include <omp.h>

Com

void main () {
    const int NUM_THREADS = 4;
    const long num_steps = 10;
    double step = 1.0/((double)num_steps);
    double sum[NUM_THREADS];
    for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
    omp_set_num_threads(NUM_THREADS);

#pragma omp parallel
{
    int id = omp_get_thread_num();
    for (int i=id; i<num_steps; i+=NUM_THREADS) {
        double x = (i+0.5) *step;
        sum[id] += 4.0*step/(1.0+x*x);
        printf("i =%3d, id =%3d\n", i, id);
    }
}
double pi = 0;
for (int i=0; i<NUM_THREADS; i++) pi += sum[i];
printf ("pi = %6.12f\n", pi);
```

|                            |
|----------------------------|
| i = 1, id = 1              |
| i = 0, id = 0              |
| i = 2, id = 2              |
| i = 3, id = 3              |
| i = 5, id = 1              |
| i = 4, id = 0              |
| i = 6, id = 2              |
| i = 7, id = 3              |
| i = 9, id = 1              |
| i = 8, id = 0              |
| <b>pi = 3.142425985001</b> |

# Scale up: num\_steps = $10^6$

```
#include <stdio.h>
#include <omp.h>
Comp

void main () {
    const int NUM_THREADS = 4;
    const long num_steps = 1000000;
    double step = 1.0/((double)num_steps);
    double sum[NUM_THREADS];
    for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
    omp_set_num_threads(NUM_THREADS);
#pragma omp parallel
{
    int id = omp_get_thread_num();
    for (int i=id; i<num_steps; i+=NUM_THREADS) {
        double x = (i+0.5) *step;
        sum[id] += 4.0*step/(1.0+x*x);
        // printf("i =%3d, id =%3d\n", i, id);
    }
}
double pi = 0;
for (int i=0; i<NUM_THREADS; i++) pi += sum[i];
printf ("pi = %6.12f\n", pi);
```

pi = 3.141592653590

# Can We Parallelize Computing sum?

```
#include <stdio.h>
#include <omp.h>

void main () {
    const int NUM_THREADS = 1000;
    const long num_steps = 100000;
    double step = 1.0/((double)num_steps);
    double sum[NUM_THREADS];
    for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
    double pi = 0;
    omp_set_num_threads(NUM_THREADS);

#pragma omp parallel
{
    int id = omp_get_thread_num();
    for (int i=id; i<num_steps; i+=NUM_THREADS) {
        double x = (i+0.5) *step;
        sum[id] += 4.0*step/(1.0+x*x);
    }
    pi += sum[id];
}
printf ("pi = %6.12f\n", pi);
```

McMahon & Weaver

Always looking for ways to beat Amdahl's Law ...

Summation inside parallel section

- Insignificant speedup in this example, but ..
- **pi = 3.138450662641**
- Wrong! And value changes between runs?!
- What's going on?

# What's Going On?

```
#include <stdio.h>
#include <omp.h>

void main () {
    const int NUM_THREADS = 1000;
    const long num_steps = 100000;
    double step = 1.0/((double)num_steps);
    double sum[NUM_THREADS];
    for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
    double pi = 0;
    omp_set_num_threads(NUM_THREADS);
#pragma omp parallel
{
    int id = omp_get_thread_num();
    for (int i=id; i<num_steps; i+=NUM_THREADS) {
        double x = (i+0.5) *step;
        sum[id] += 4.0*step/(1.0+x*x);
    }
    pi += sum[id]; ←
}
printf ("pi = %6.12f\n", pi);
```

- Operation is really  
 $\text{pi} = \text{pi} + \text{sum}[id]$
- What if  $>1$  threads reads current (same) value of **pi**, computes the sum, stores the result back to **pi**?
- Each processor reads same intermediate value of **pi**!
- Result depends on who gets there when
  - A “race” → **result is not deterministic but if we locked this we'd lose almost all speedup**

# OpenMP Reduction

- ```
double avg, sum=0.0, A[MAX]; int i;
#pragma omp parallel for private ( sum )
for (i = 0; i <= MAX ; i++) {sum += A[i];}
avg = sum/MAX; // bug, we only get the master thread's sum
```
- Problem is that we really want sum over all threads!
- **Reduction:** specifies that 1 or more variables that are private to each thread are subject of reduction operation at end of parallel region:  
`reduction(operation:var) where`
  - Operation: operator to perform on the variables (var) at the end of the parallel region
  - Var: One or more variables on which to perform scalar reduction: private than combined
- ```
double avg, sum=0.0, A[MAX]; int i;
#pragma omp for reduction(+ : sum)
for (i = 0; i <= MAX ; i++) {sum += A[i];}
avg = sum/MAX;
```

# Calculating $\pi$ Simple Version

```
#include <omp.h>
#define NUM_THREADS 4
static long num_steps = 100000; double step;

void main () {
    int i;      double x, pi, sum[NUM_THREADS];
    step = 1.0/(double) num_steps;
#pragma omp parallel private ( i, x )
{
    int id = omp_get_thread_num();
    for (i=id, sum[id]=0.0; i<num_steps; i=i+NUM_THREADS)
    {
        x = (i+0.5)*step;
        sum[id] += 4.0/(1.0+x*x);
    }
}
for(i=1; i<NUM_THREADS; i++)
    sum[0] += sum[i]; pi = sum[0];
printf ("pi = %6.12f\n", pi);
}
```

# Version 2: parallel for, reduction

```
#include <omp.h>
#include <stdio.h>
static long num_steps = 100000;
double step;
void main ()
{
    int i;      double x, pi, sum = 0.0;
    step = 1.0/(double) num_steps;
    #pragma omp parallel for private(x) reduction(+:sum)
    for (i=1; i<= num_steps; i++) {
        x = (i-0.5)*step;
        sum = sum + 4.0/(1.0+x*x);
    }
    pi = sum;
    printf ("pi = %6.8f\n", pi);
}
```

# Reduction Options...

- Arithmetic
  - + \* -
- Comparison
  - **min max**
- Logical
  - & && | || ^
- And now you know why RISC-V has the atomic memory operations:
  - **amoadd, amoand, amoor, amoxor, amomax, amomin**
    - All but - and \* as the reduction can be implemented as single instructions

# OpenMP Timing

- Elapsed wall clock time:

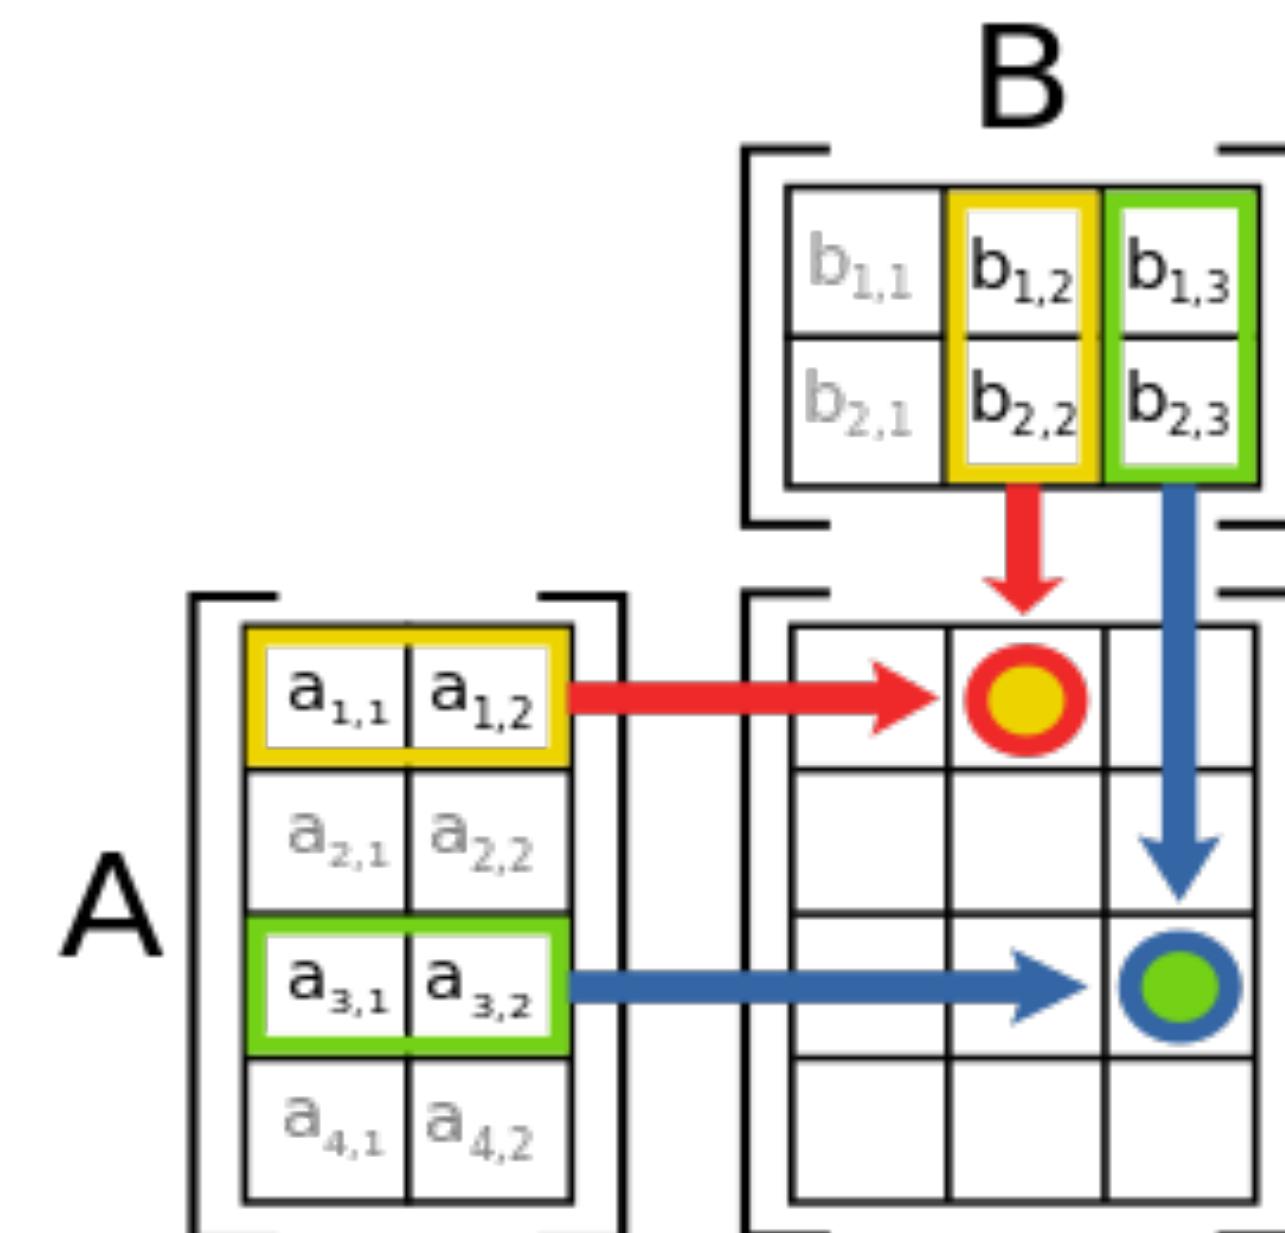
```
double omp_get_wtime(void);
```

- Returns elapsed wall clock time in seconds
- Time is measured per thread, no guarantee can be made that two distinct threads measure the same time
- Time is measured from “some time in the past,” so subtract results of two calls to **omp\_get\_wtime** to get elapsed time

# Matrix Multiply in OpenMP

```
// C[M][N] = A[M][P] × B[P][N]
start_time = omp_get_wtime();
#pragma omp parallel for private(tmp, j, k)
for (i=0; i<M; i++) {
    for (j=0; j<N; j++) {
        tmp = 0.0;
        for (k=0; k<P; k++) {
            /* C(i,j) = sum(over k) A(i,k) * B(k,j) */
            tmp += A[i][k] * B[k][j];
        }
        C[i][j] = tmp; // doing this to tmp prevents
                       // potential cache issues from multiple
                       // writers & "false sharing"
    }
}
run_time = omp_get_wtime() - start_time;
```

Outer loop spread across N threads;  
inner loops inside a single thread



# Matrix Multiply in Open MP

- More performance optimizations available:
  - Higher compiler optimization (-O2, -O3) to reduce number of instructions executed
  - Cache blocking to improve memory performance
  - Using SIMD AVX instructions to raise floating point computation rate (DLP)
- Gives a guide for Project 4:
  - 1: Get it working period
  - 2: OpenMP Parallel **for** the outer loop
    - Make sure the different threads are writing to different parts of memory (CRITICAL!)
  - 3: Make the inner loop SIMD vectorized
  - 4: Make the inner loop unrolled x4
  - 5: Add blocking (~32 or 64 sounds like a good number)

# Data Races and Synchronization

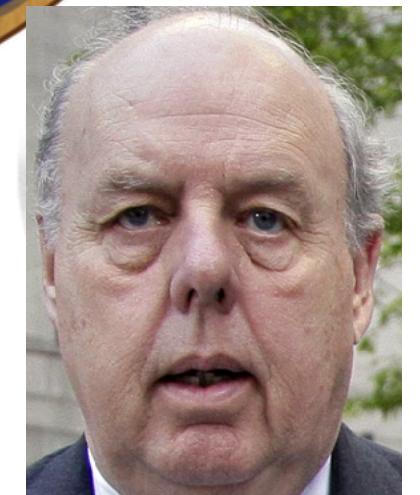
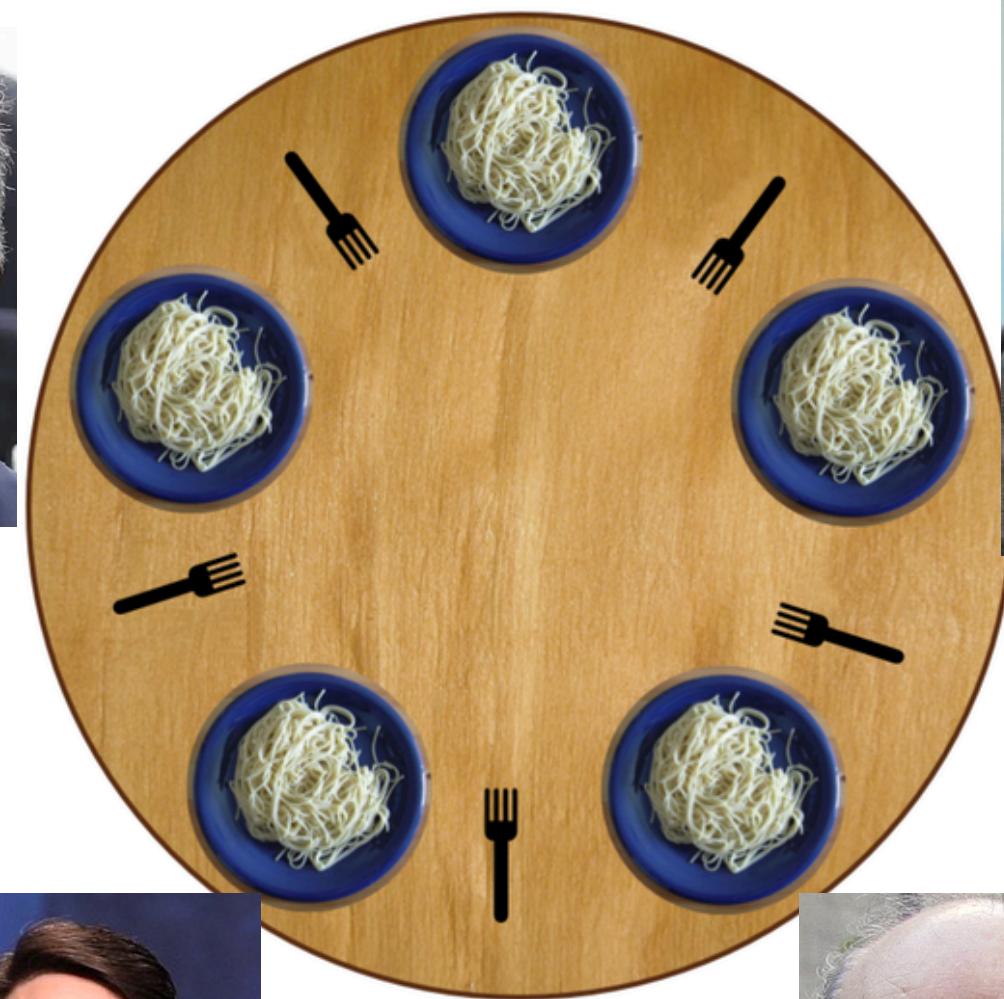
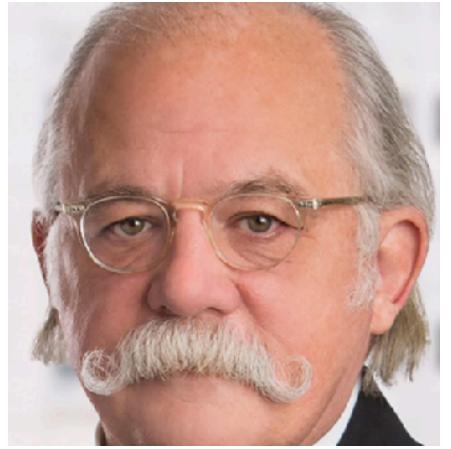
- Two memory accesses form a *data race* if from different threads access same location, at least one is a write, and they occur one after another
- If there is a data race, result of program varies depending on chance (which thread first?)
- Avoid data races by synchronizing writing and reading to get *deterministic* behavior
- Synchronization done by user-level routines that rely on hardware synchronization instructions

# Reminder: Locks

- Computers use locks to control access to shared resources
  - Serves purpose of microphone in example
  - Also referred to as “semaphore”
    - Although “semaphores” have slightly different semantics
- Usually implemented with a variable
  - In most object-oriented languages its an object that you can call
  - Under the hood its usually just an integer that has atomic updates:  
0 == unlocked, 1 == locked
- Two operations:
  - lock.acquire() // blocks this thread until the lock is unlocked
  - lock.release() // Unlocks the lock
- OpenMP also has a lock
  - `#pragma omp critical`  
{  
... Only one thread is running at a time here  
}

# Deadlock

- Deadlock: a system state in which no progress is possible because everything is locked waiting for something else
- Dining ~~Philosopher's~~ Lawyers Problem:
  - Pontificate until the left fork is available; when it is, pick it up
  - Pontificate until the right fork is available; when it is, pick it up
  - When both forks are held, eat for a fixed amount of time
  - Then, put the right fork down
  - Then, put the left fork down
  - Repeat from the beginning
- Solution?

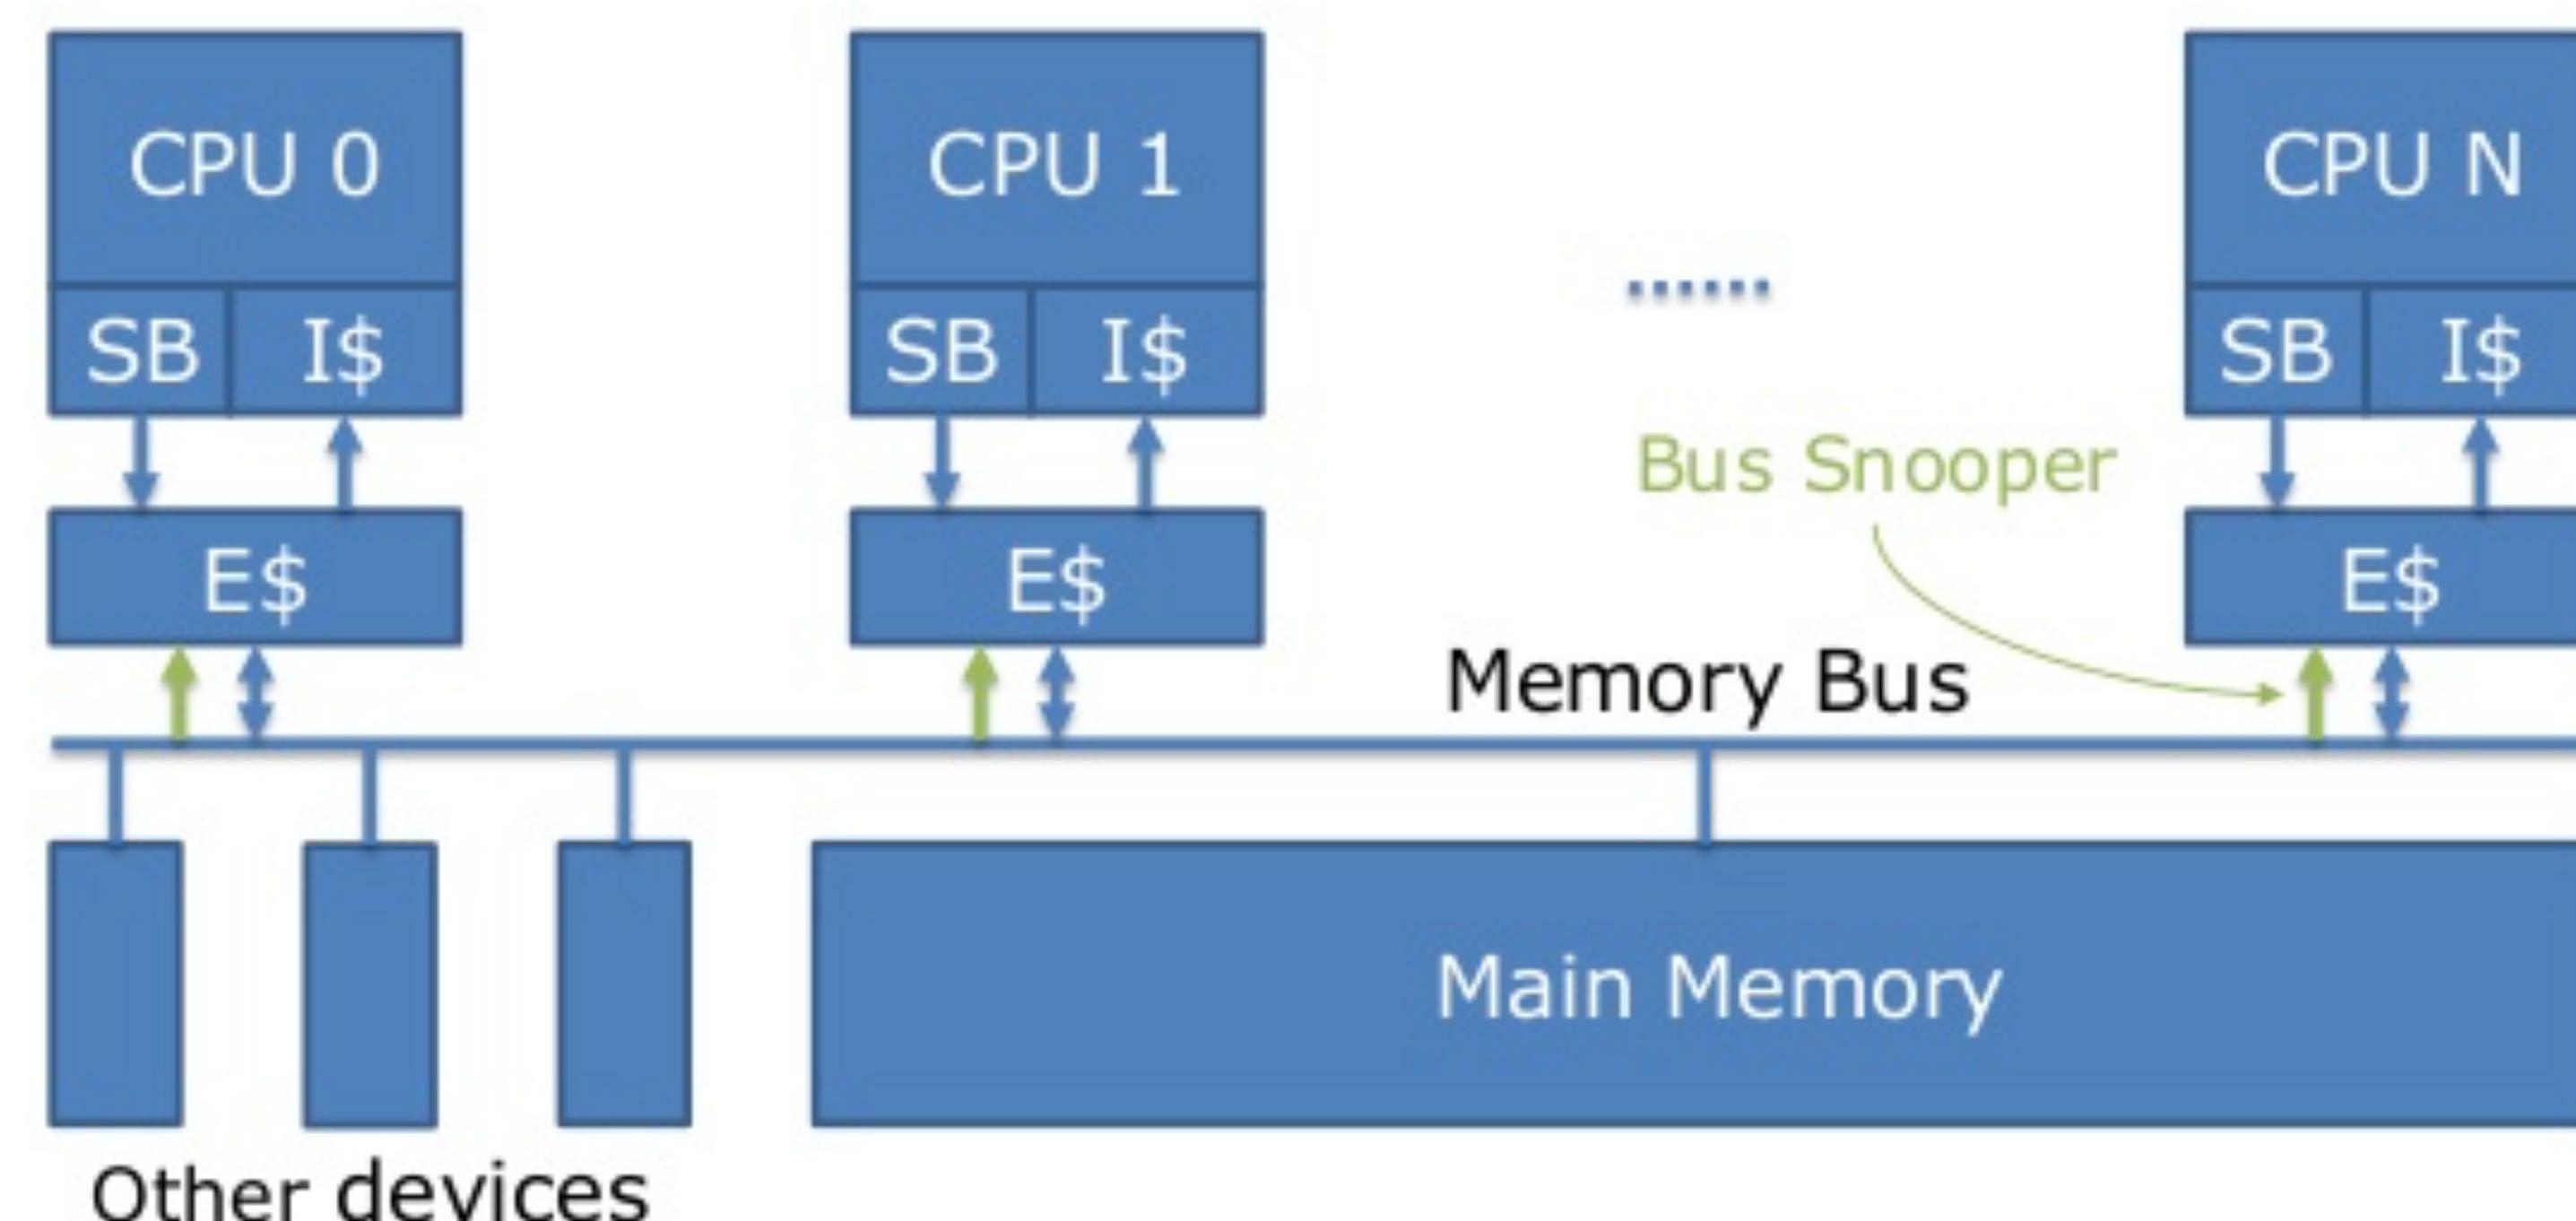


# Limiting Parallelism

- Locks act to inhibit parallelism
  - Innately sequential regions -> Amdahl's law problem...
  - Python is "threaded" but not really:  
A global interpreter lock means other threads can be waiting on I/O, but (mostly) only one compute thread at a time
- Can try to limit the locks
  - Rather than locking everything have a different lock for each region...
    - But be careful, then its much easier to get into deadlock situations
- Or try to eliminate locks altogether
  - Thus, e.g. why Go's parallelism is not focused around locks, and RISC-V has the **lrfsc** instructions

# (Chip) Multicore Multiprocessor

- SMP: (Shared Memory) Symmetric Multiprocessor
  - Two or more identical CPUs/Cores
  - Single shared coherent memory



# Multiprocessor Key Questions

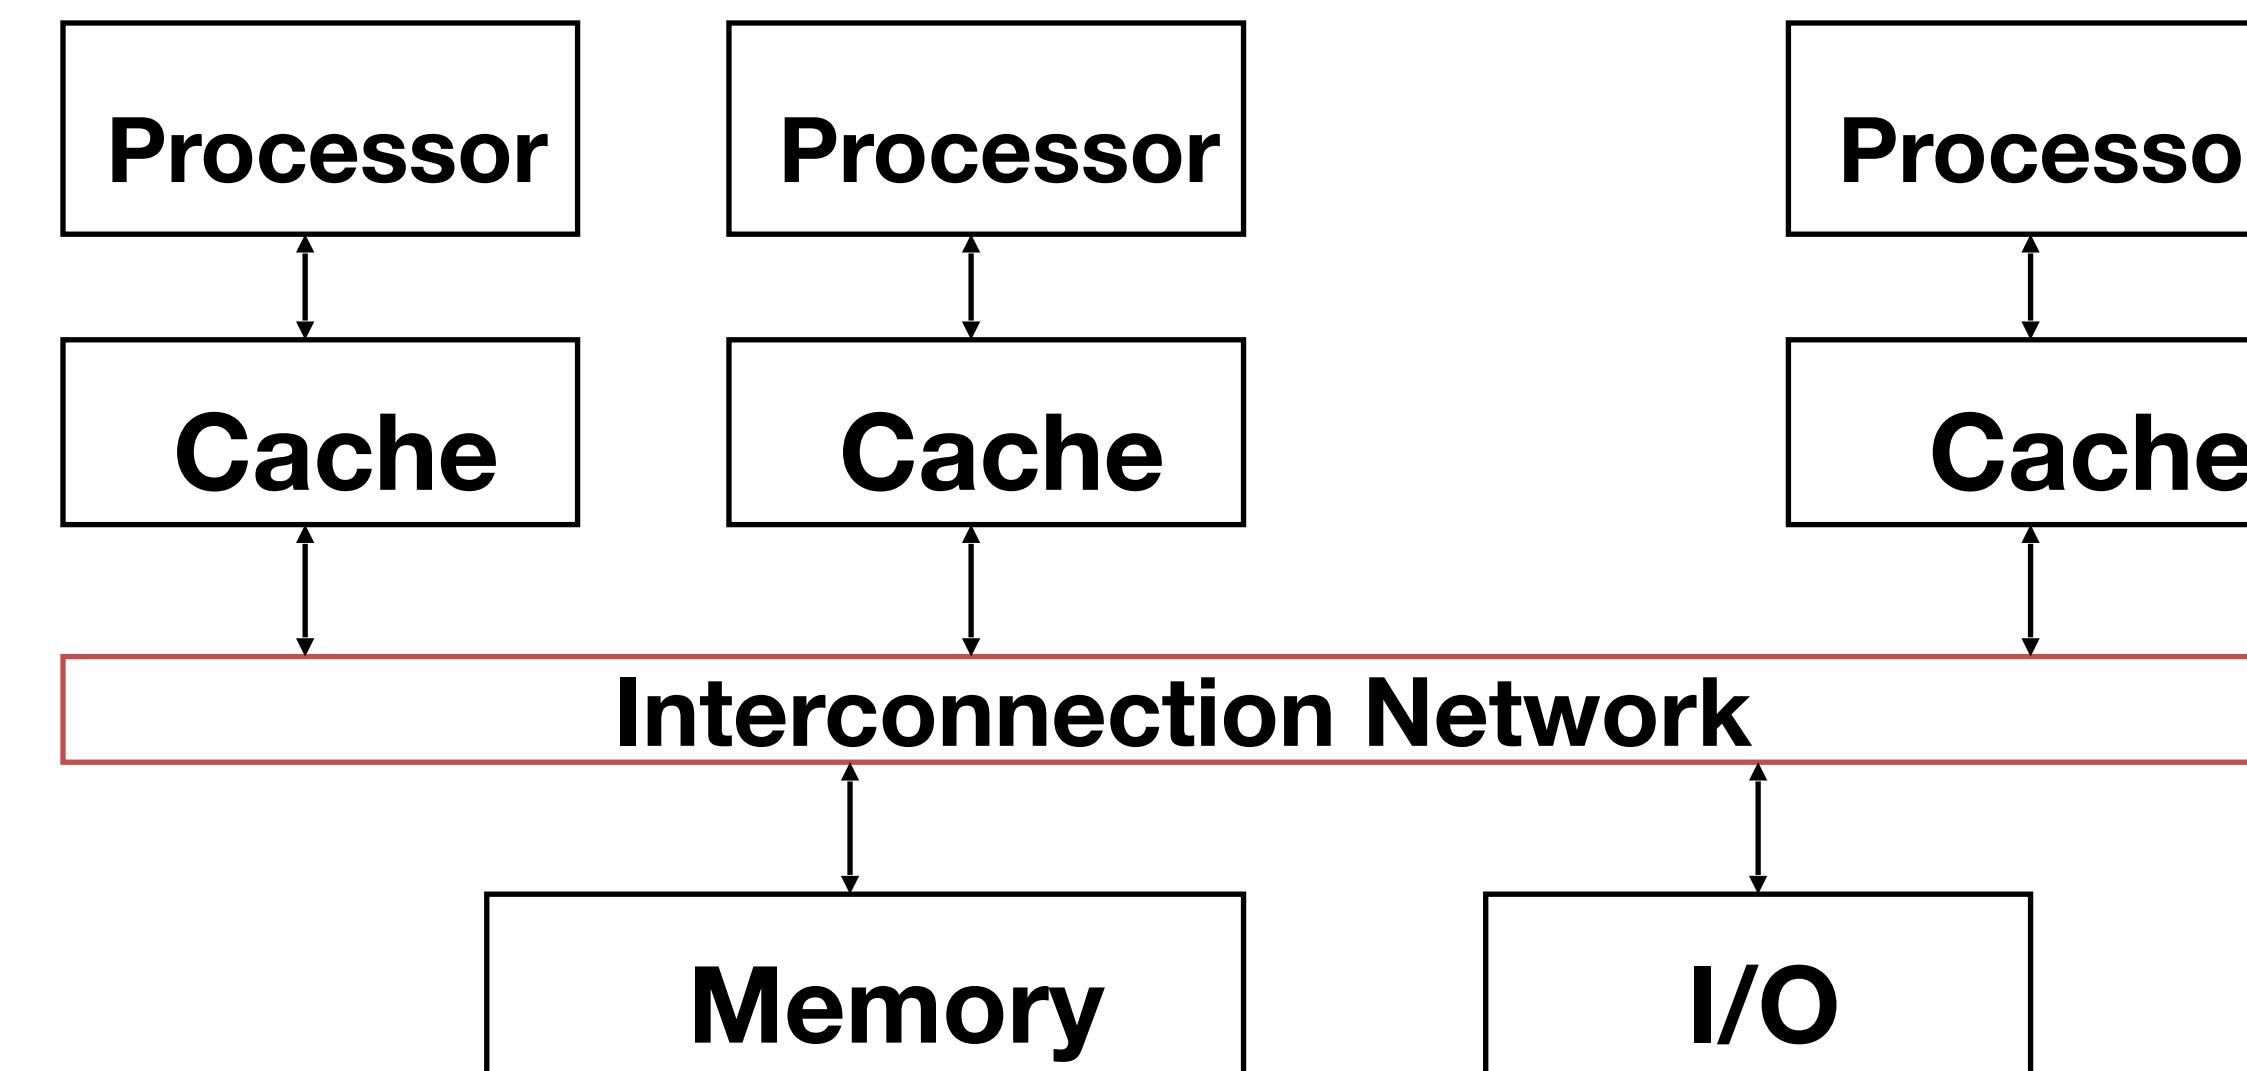
- Q1 – How do they share data?
- Q2 – How do they coordinate?
- Q3 – How many processors can be supported?

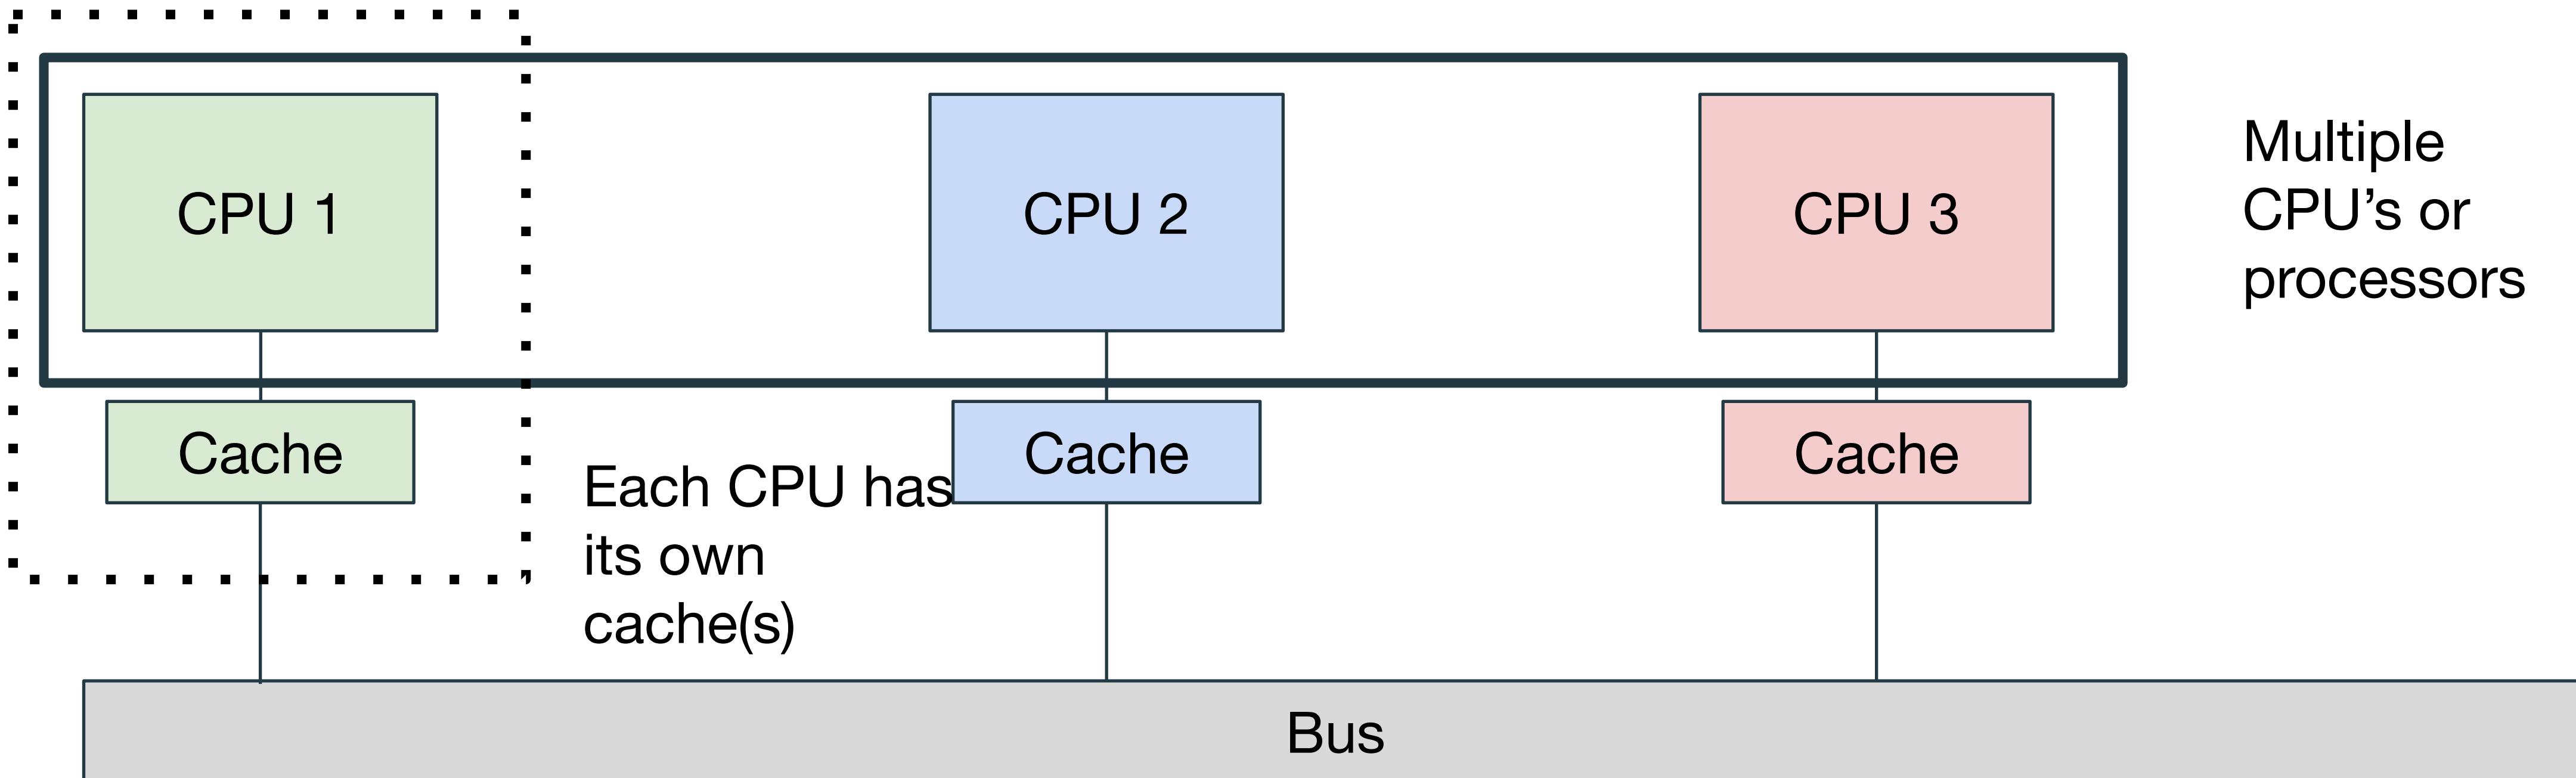
# Shared Memory Multiprocessor (SMP)

- Q1 – Single physical address space shared by all processors/cores
- Q2 – Processors coordinate/communicate through shared variables in memory (via loads and stores)
  - Use of shared data must be coordinated via synchronization primitives (locks) that allow access to data to only one processor at a time
  - Effectively all multicore computers today are SMP:  
Only difference is modern SMPs will use something that looks more like a network to build the shared "bus" (and shared L3\$)
- Q3 - Depends on the workload!
  - Most systems go with "Best available single core within constraints, duplicate that"
  - Power-critical systems (e.g. phones) go "Some of the best available single cores, some of the most power efficient single cores"

# Multiprocessor Caches

- Memory is a performance bottleneck even with one processor
- Use caches to reduce bandwidth demands on main memory
- Each core has a local private cache holding data it has accessed recently
- Only cache misses have to access the shared common memory



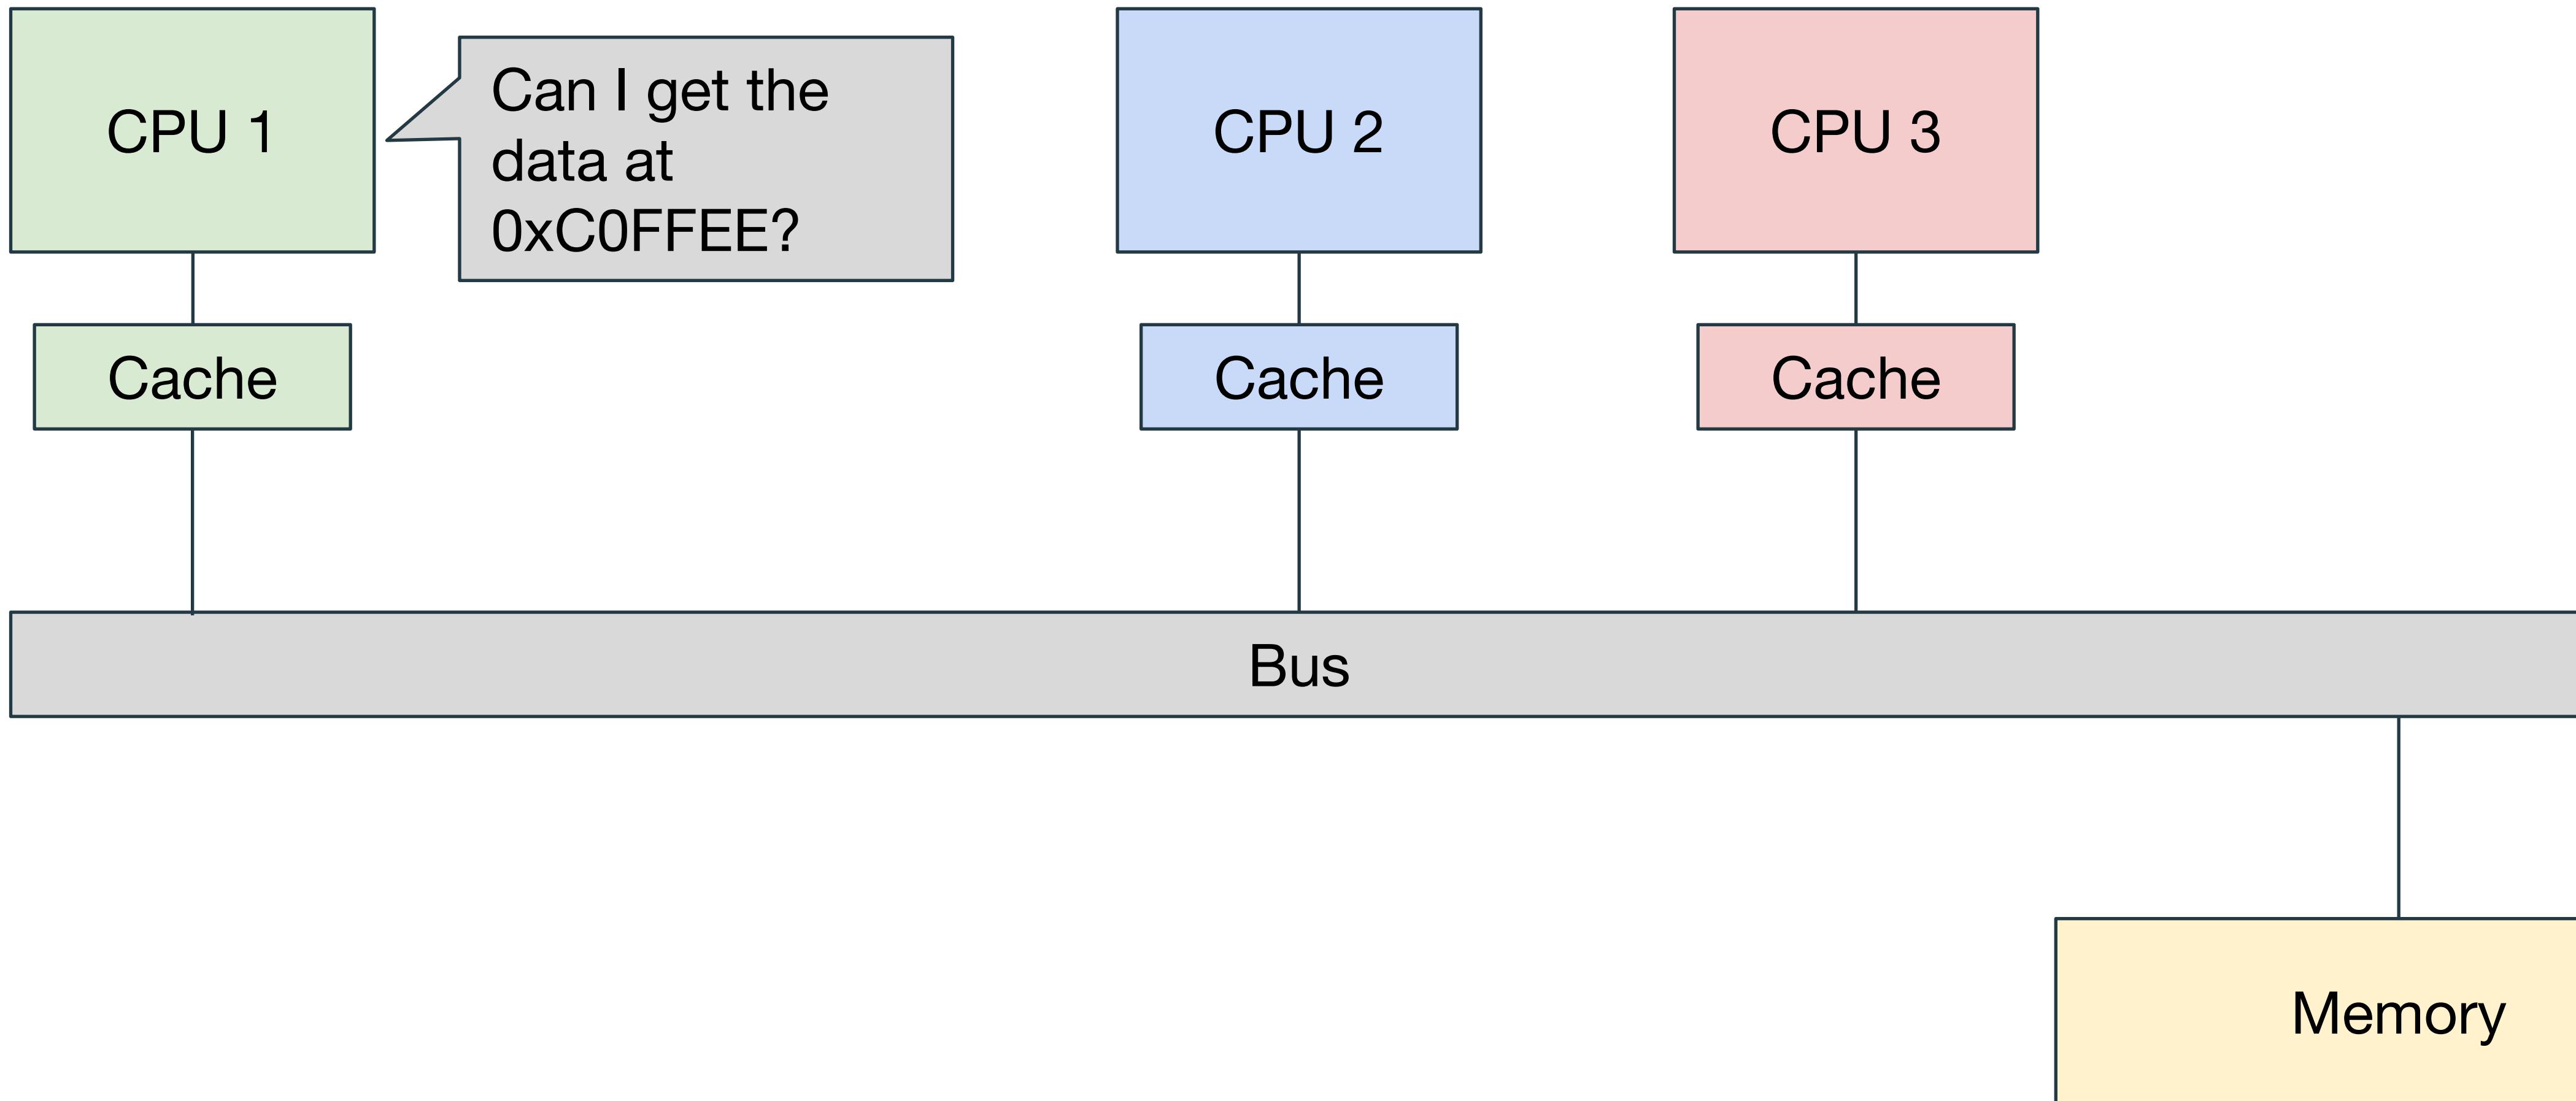


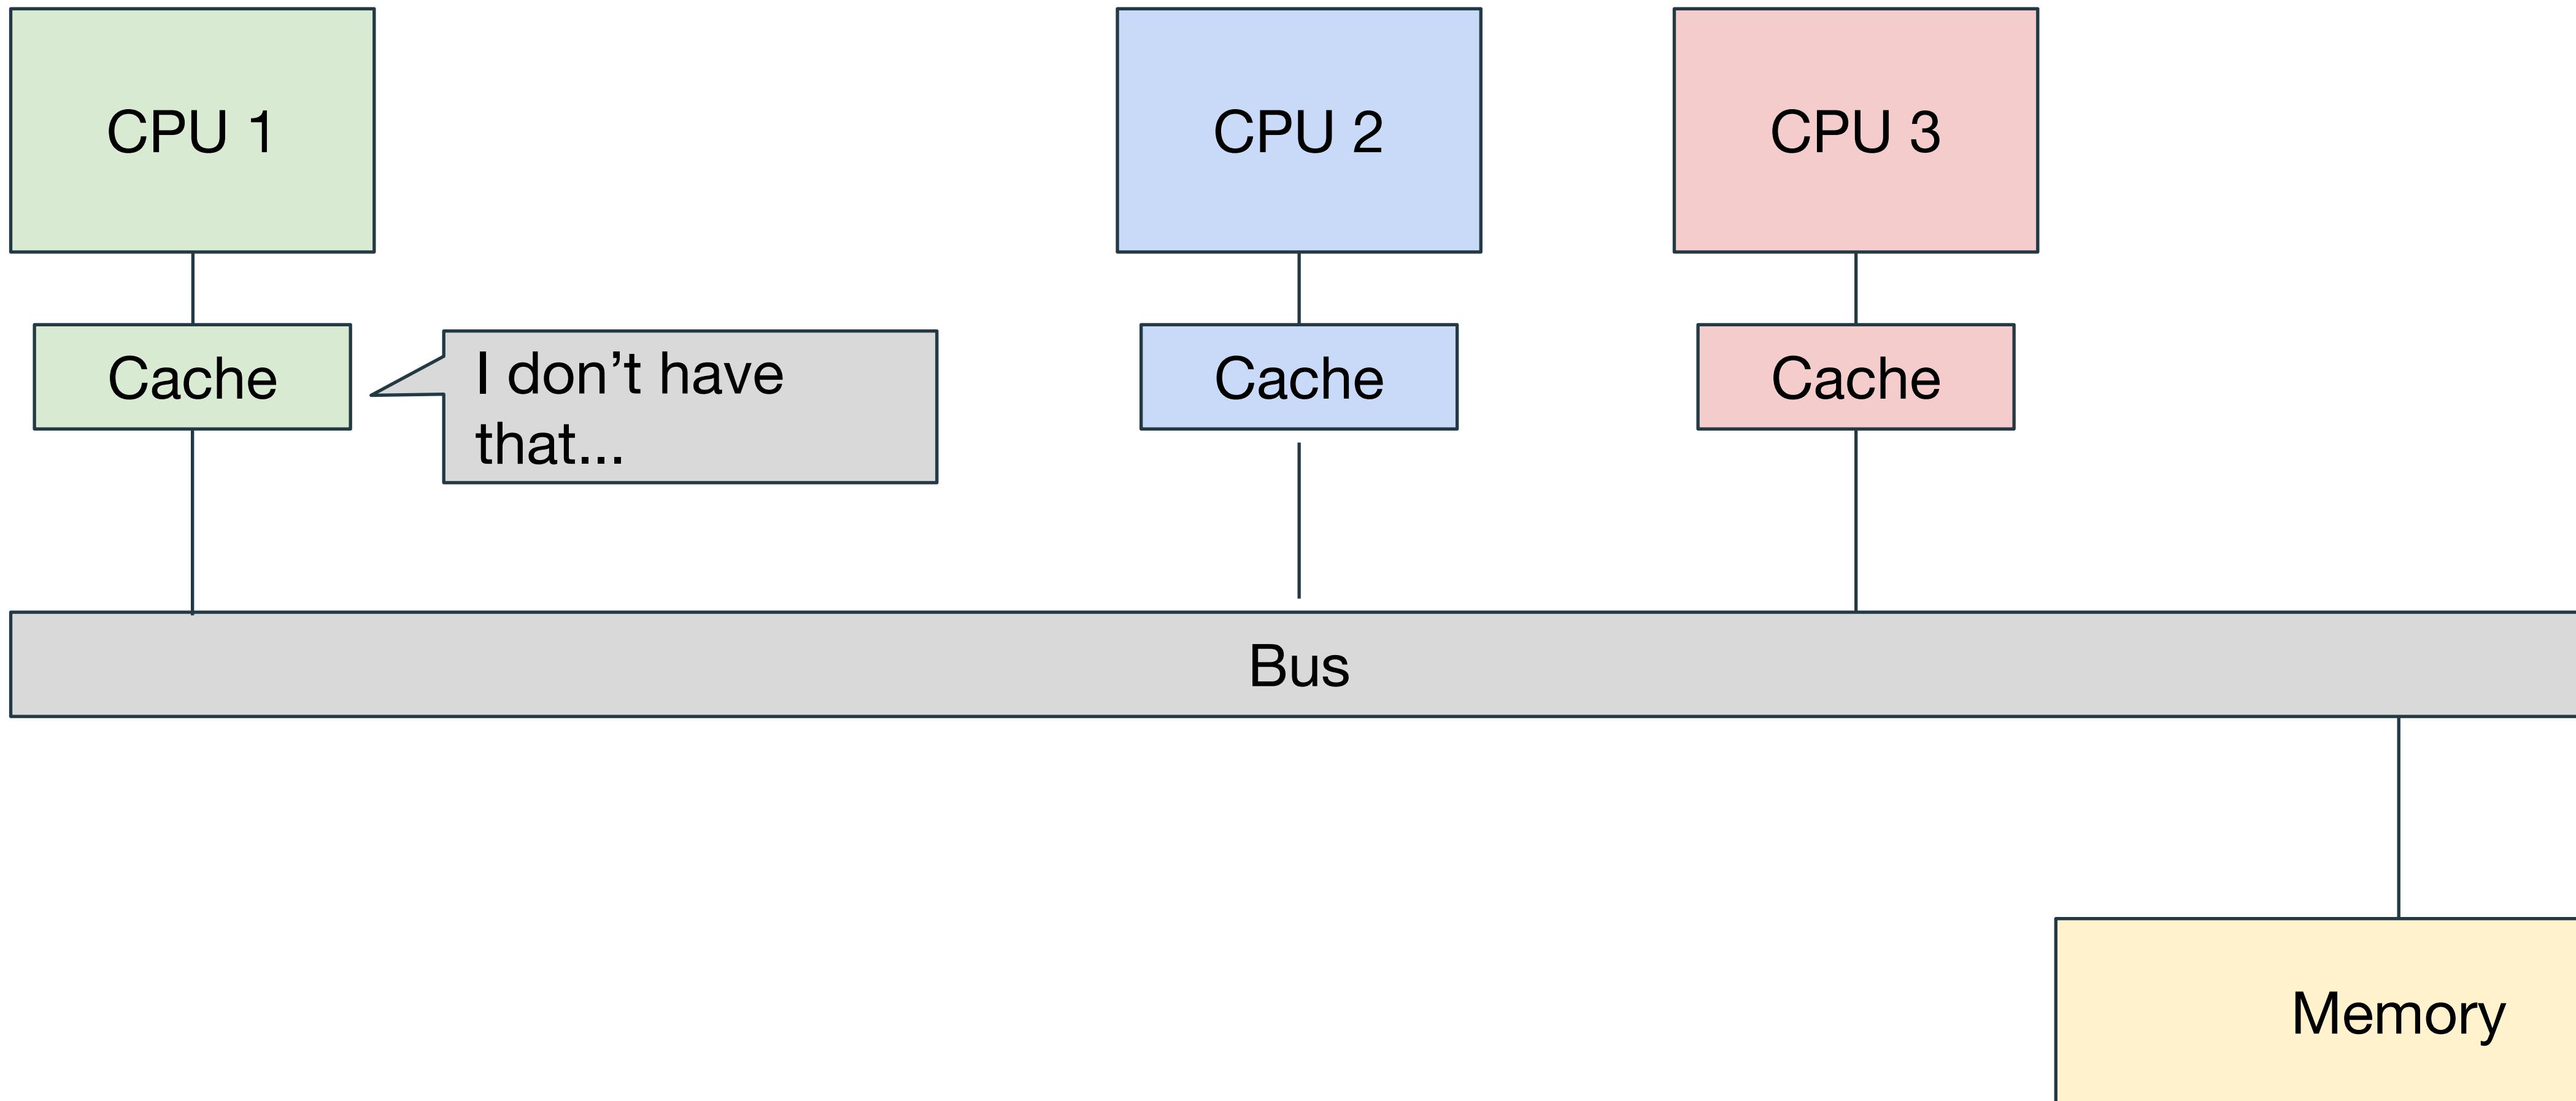
All CPU's communicate with each other and memory through a bus

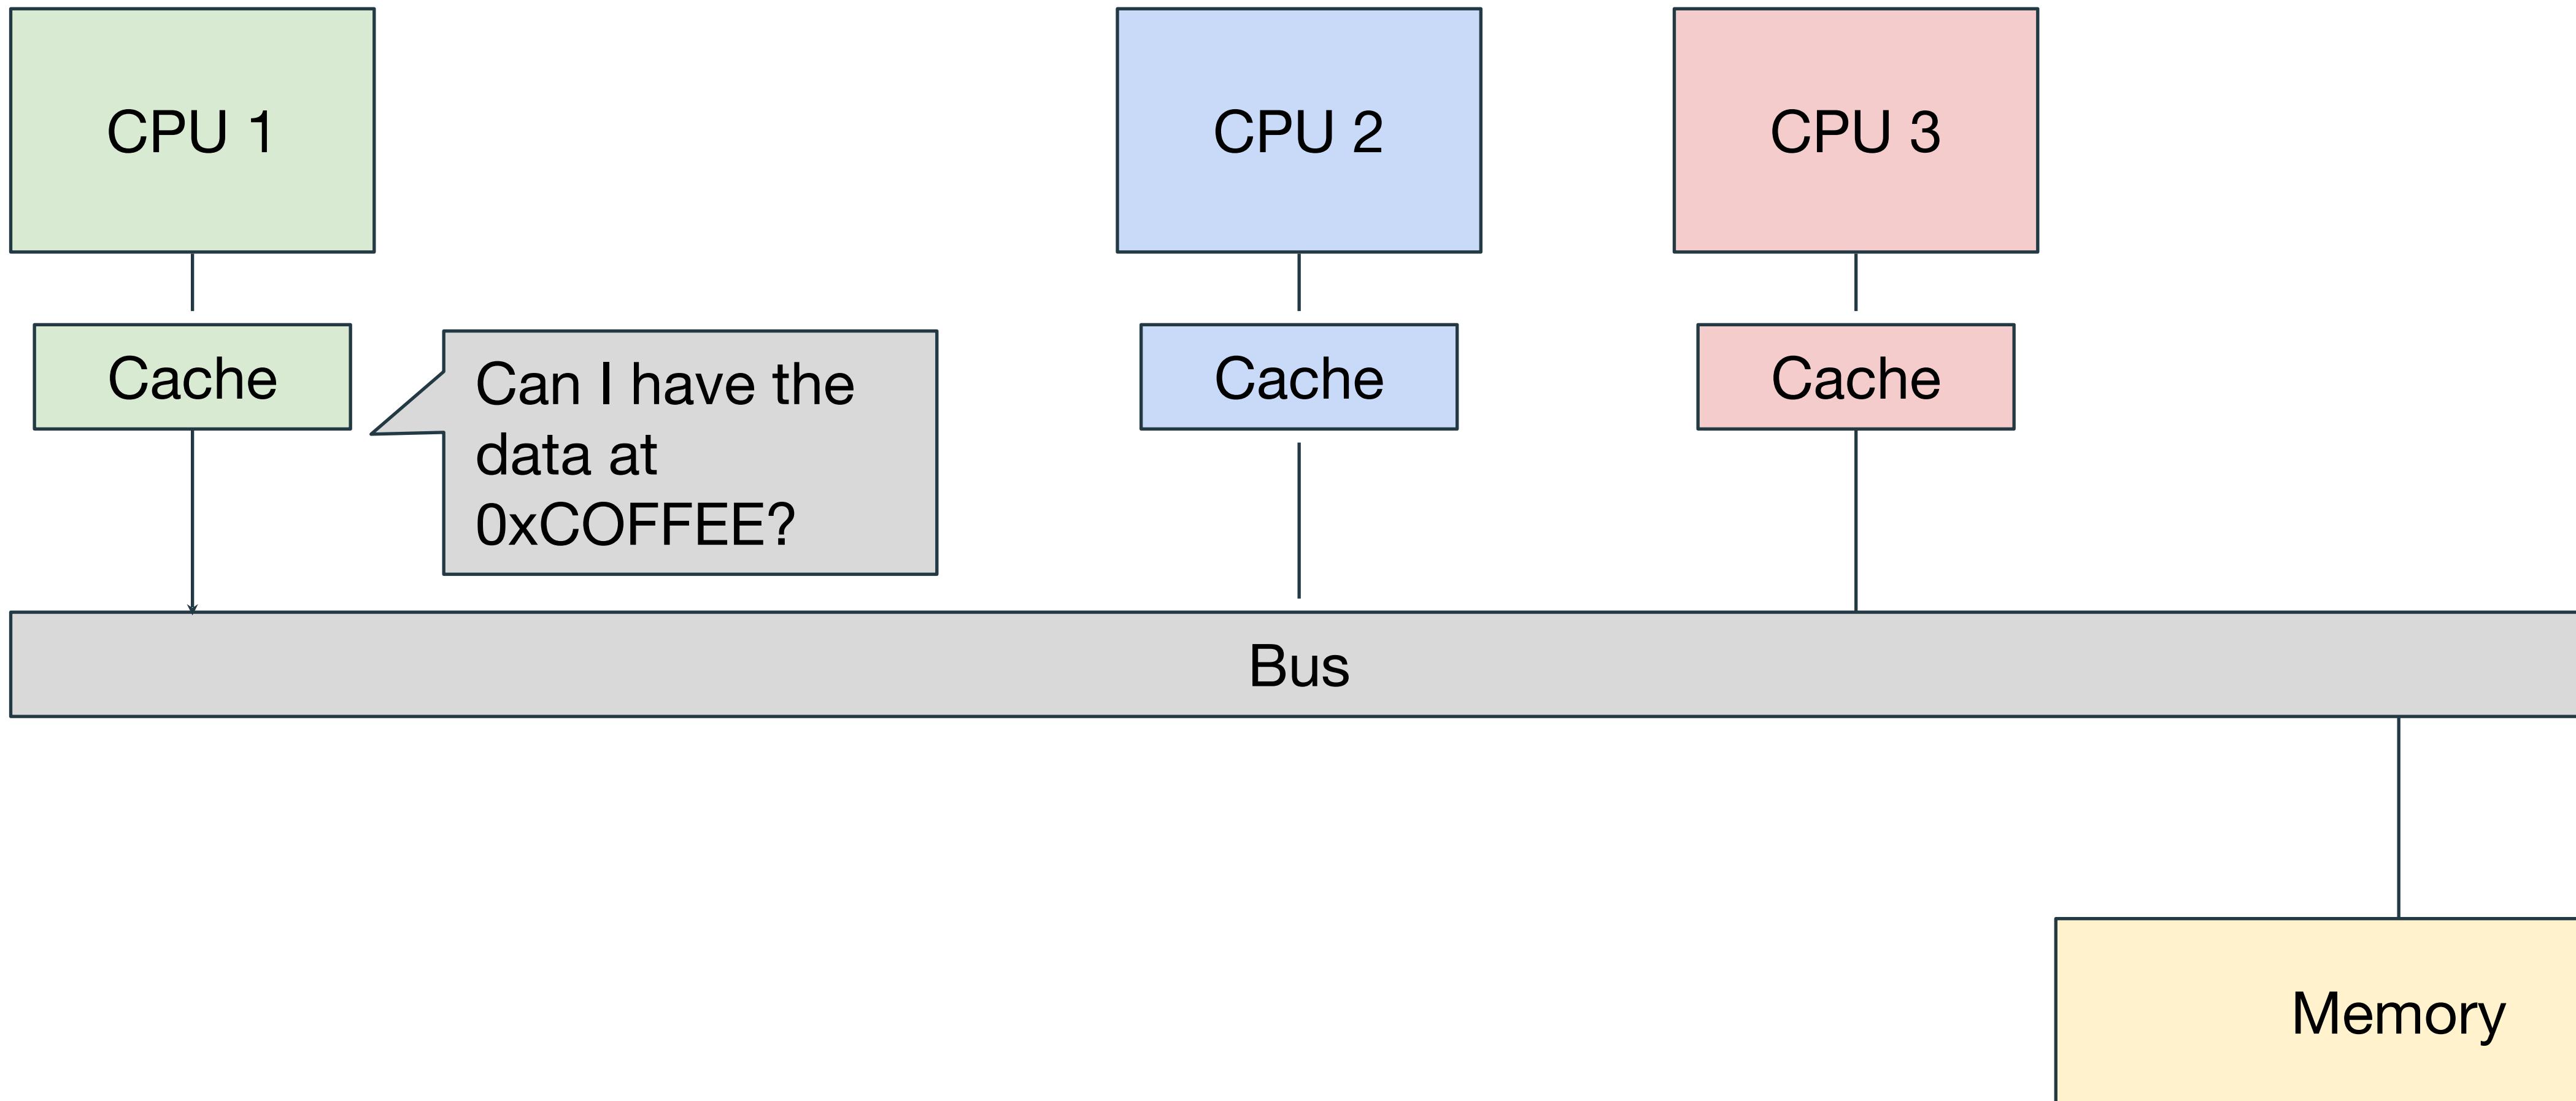
Let's go through an example to motivate why we need cache coherence!

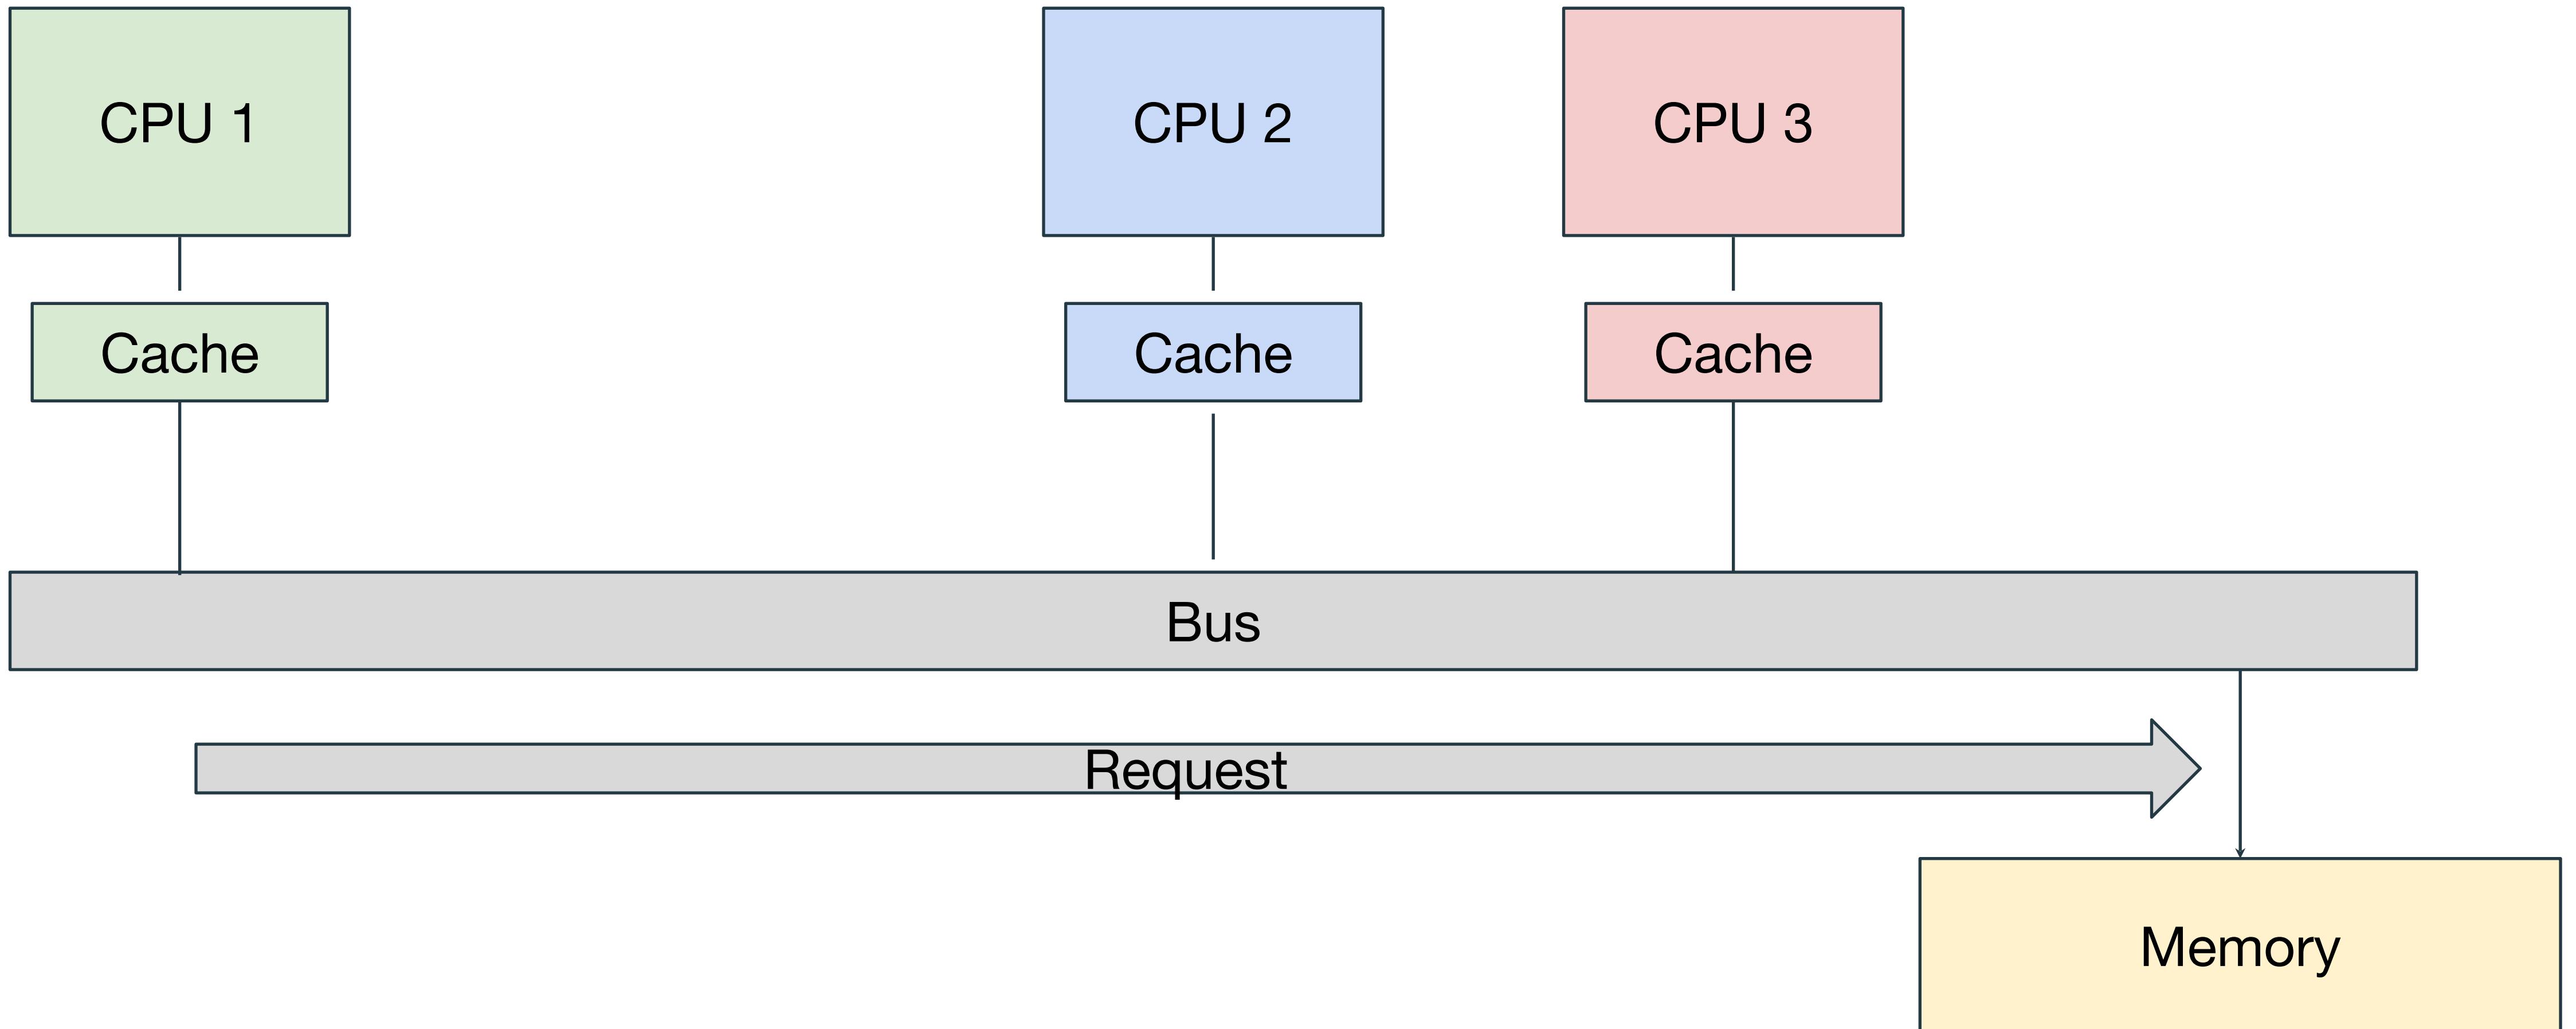
One bank of memory is shared by all CPU's

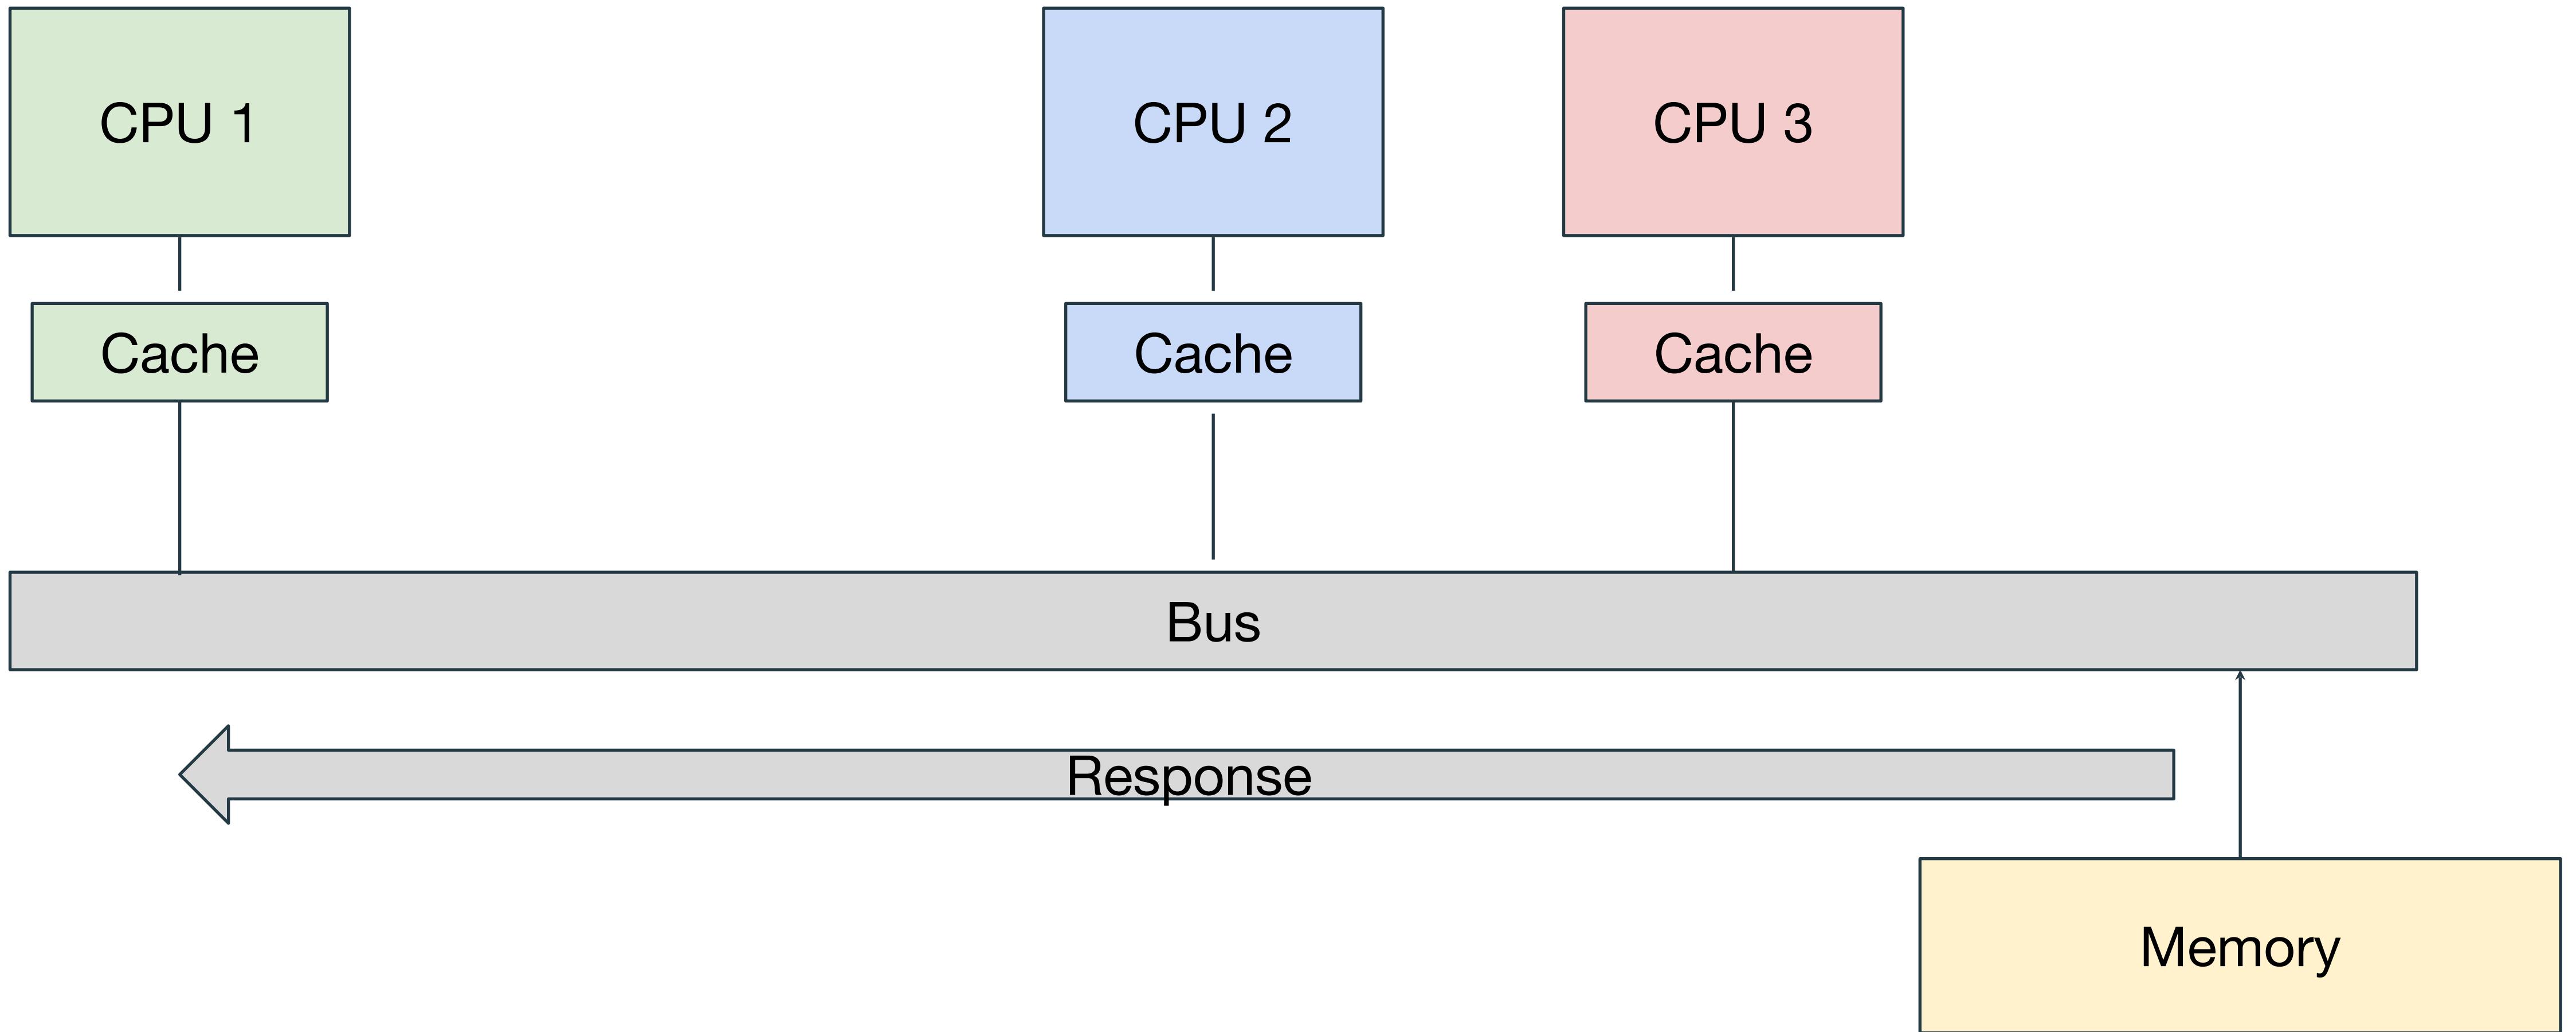
Memory

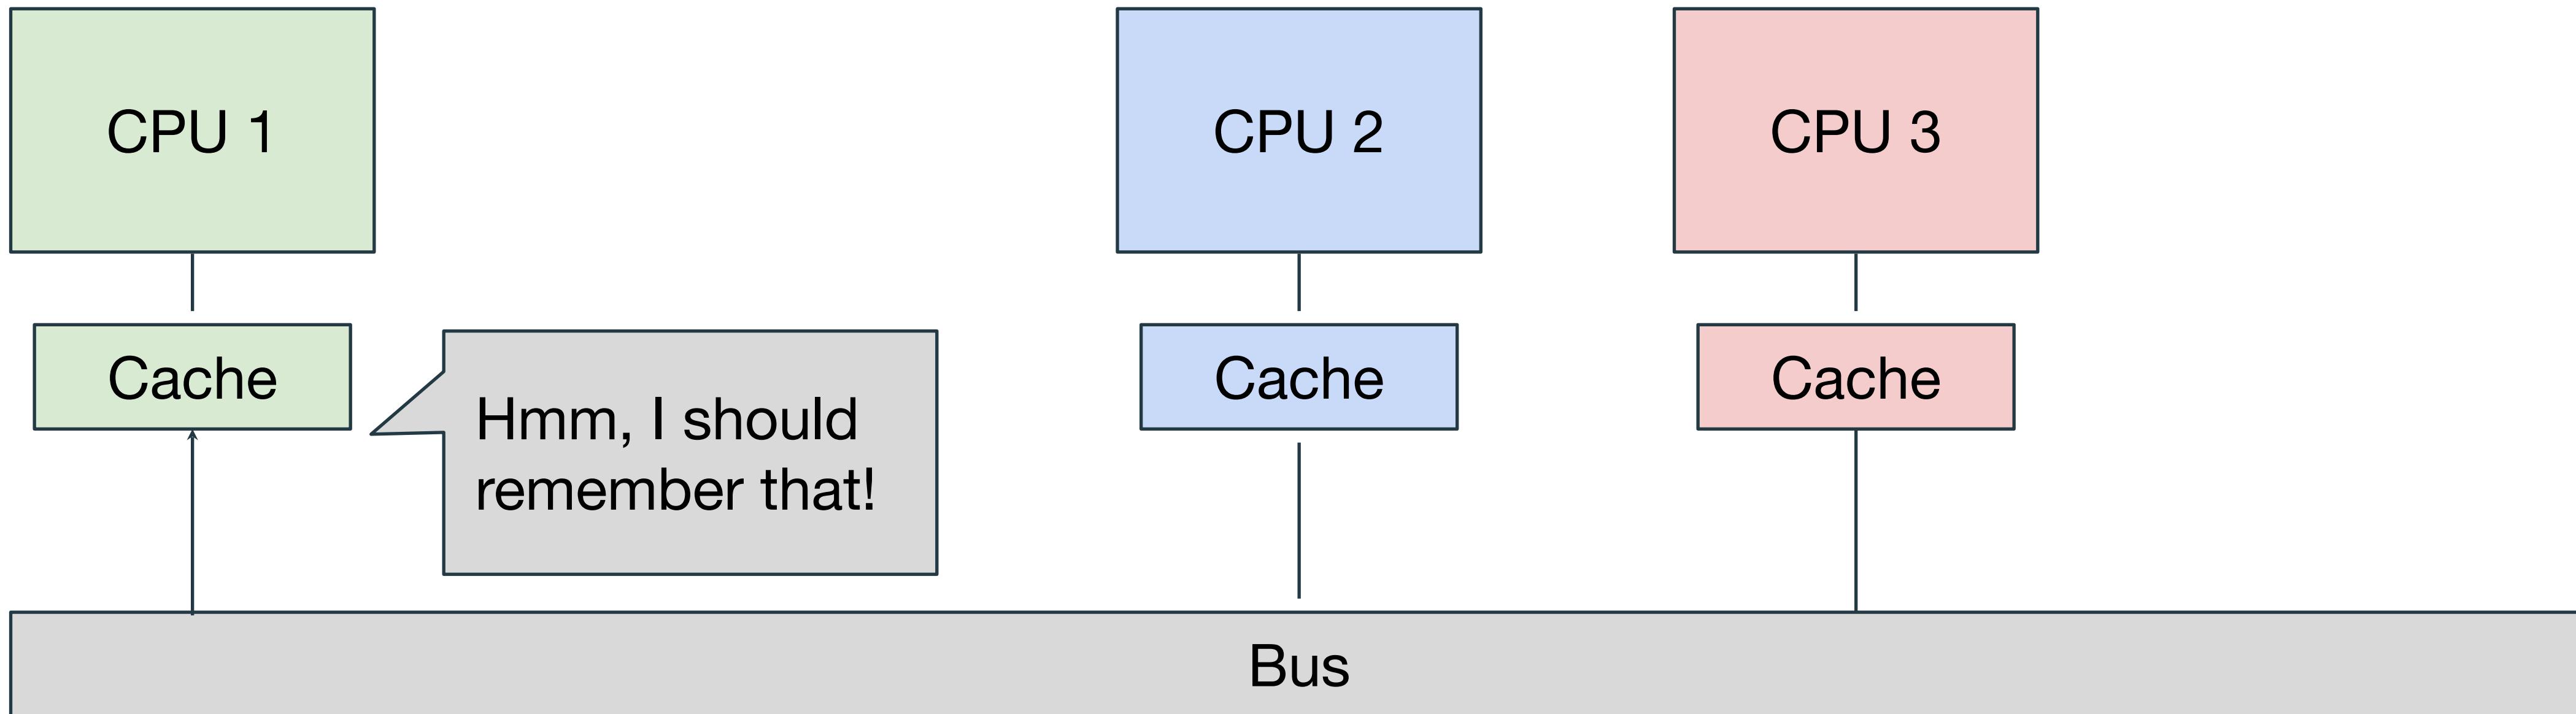




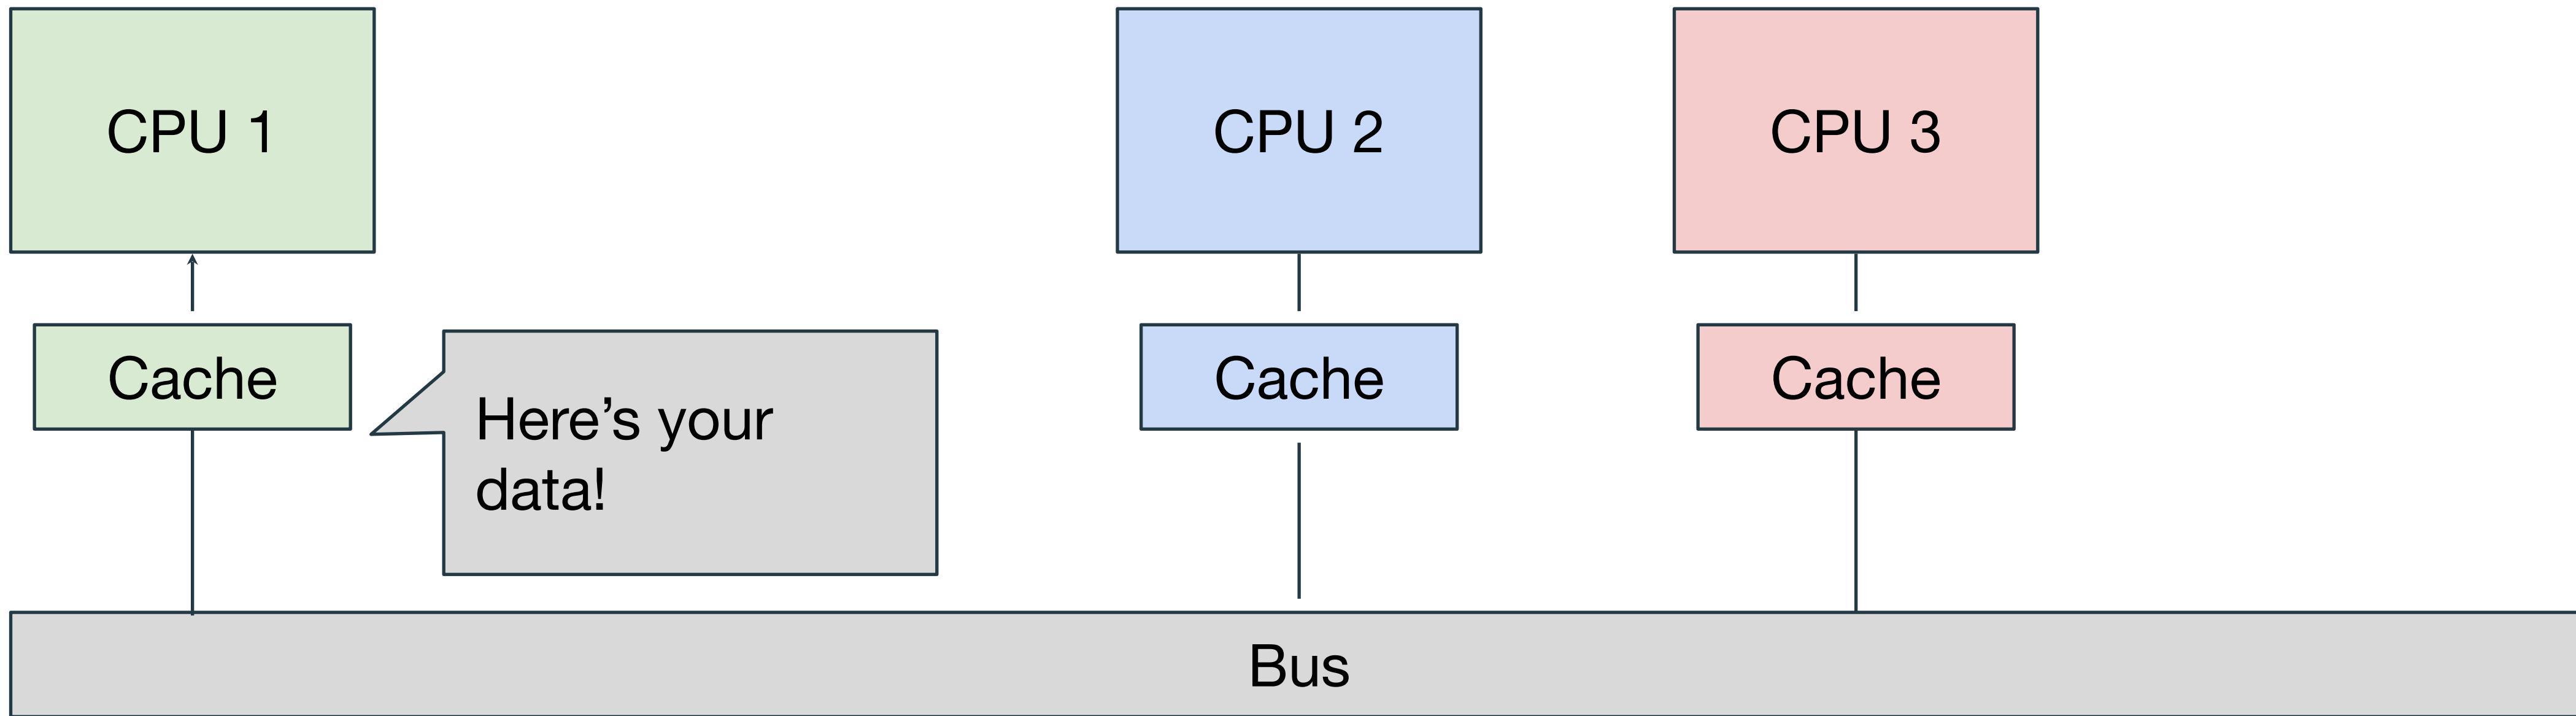




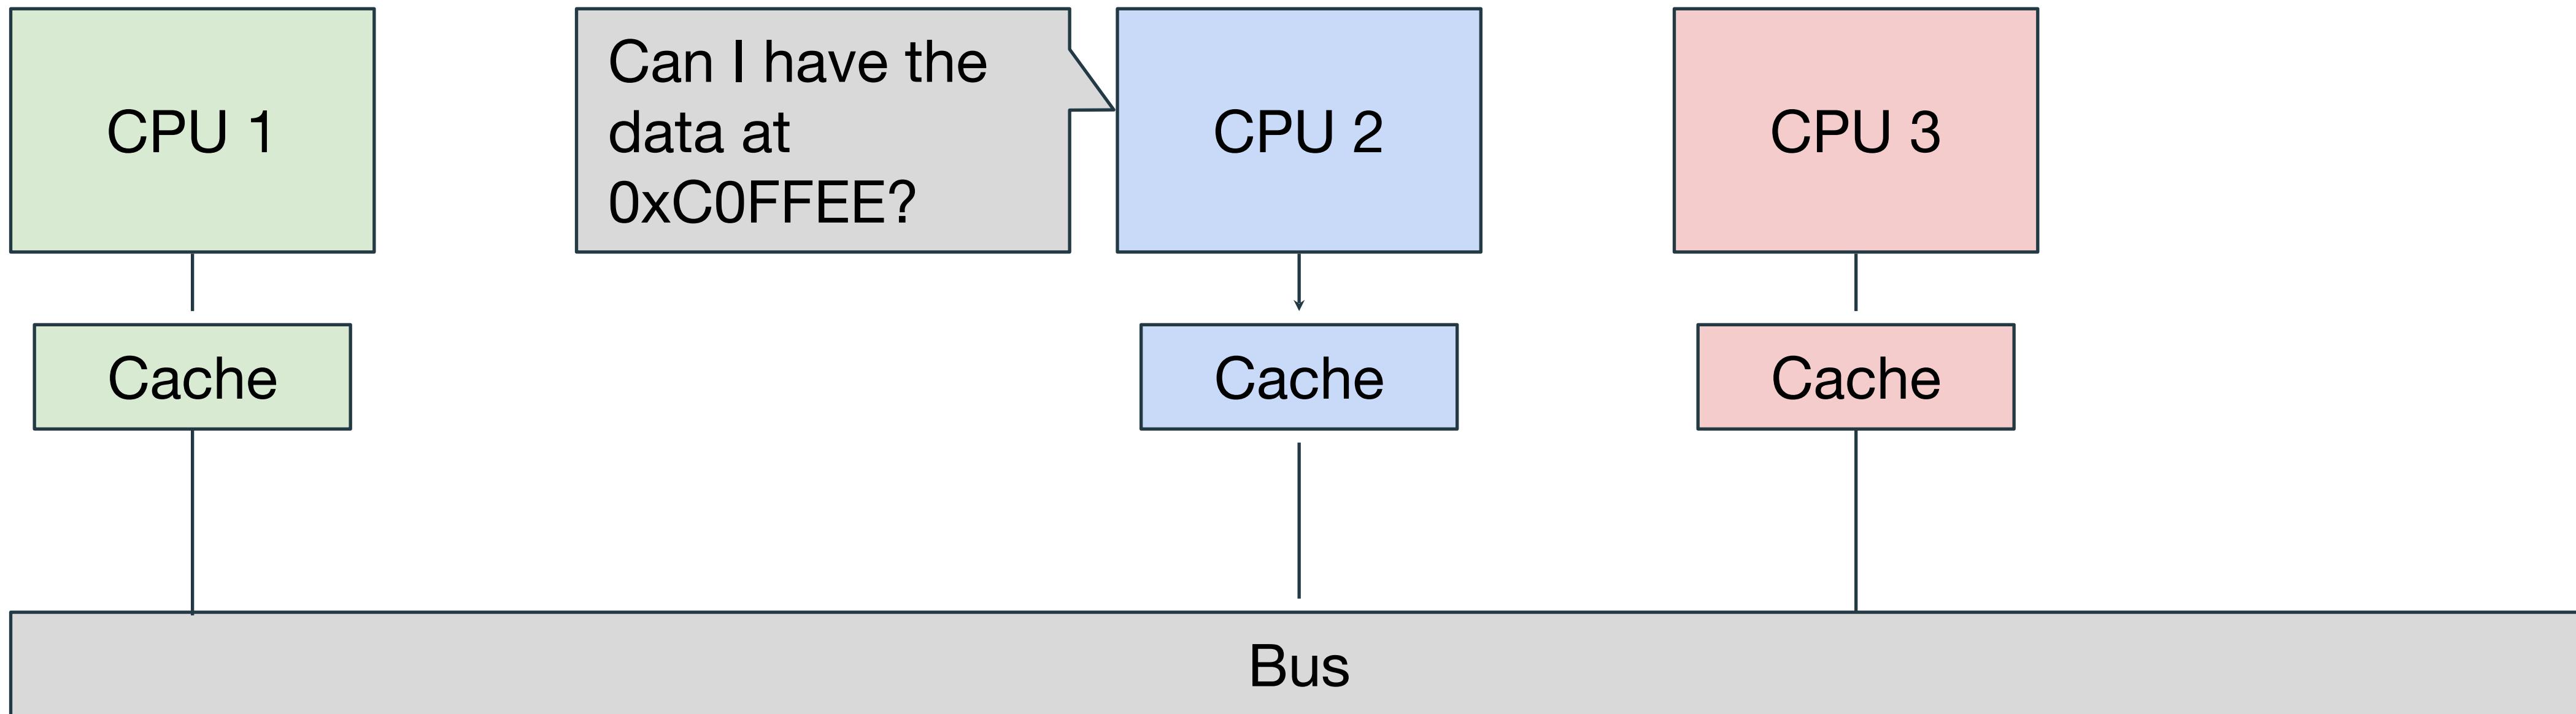




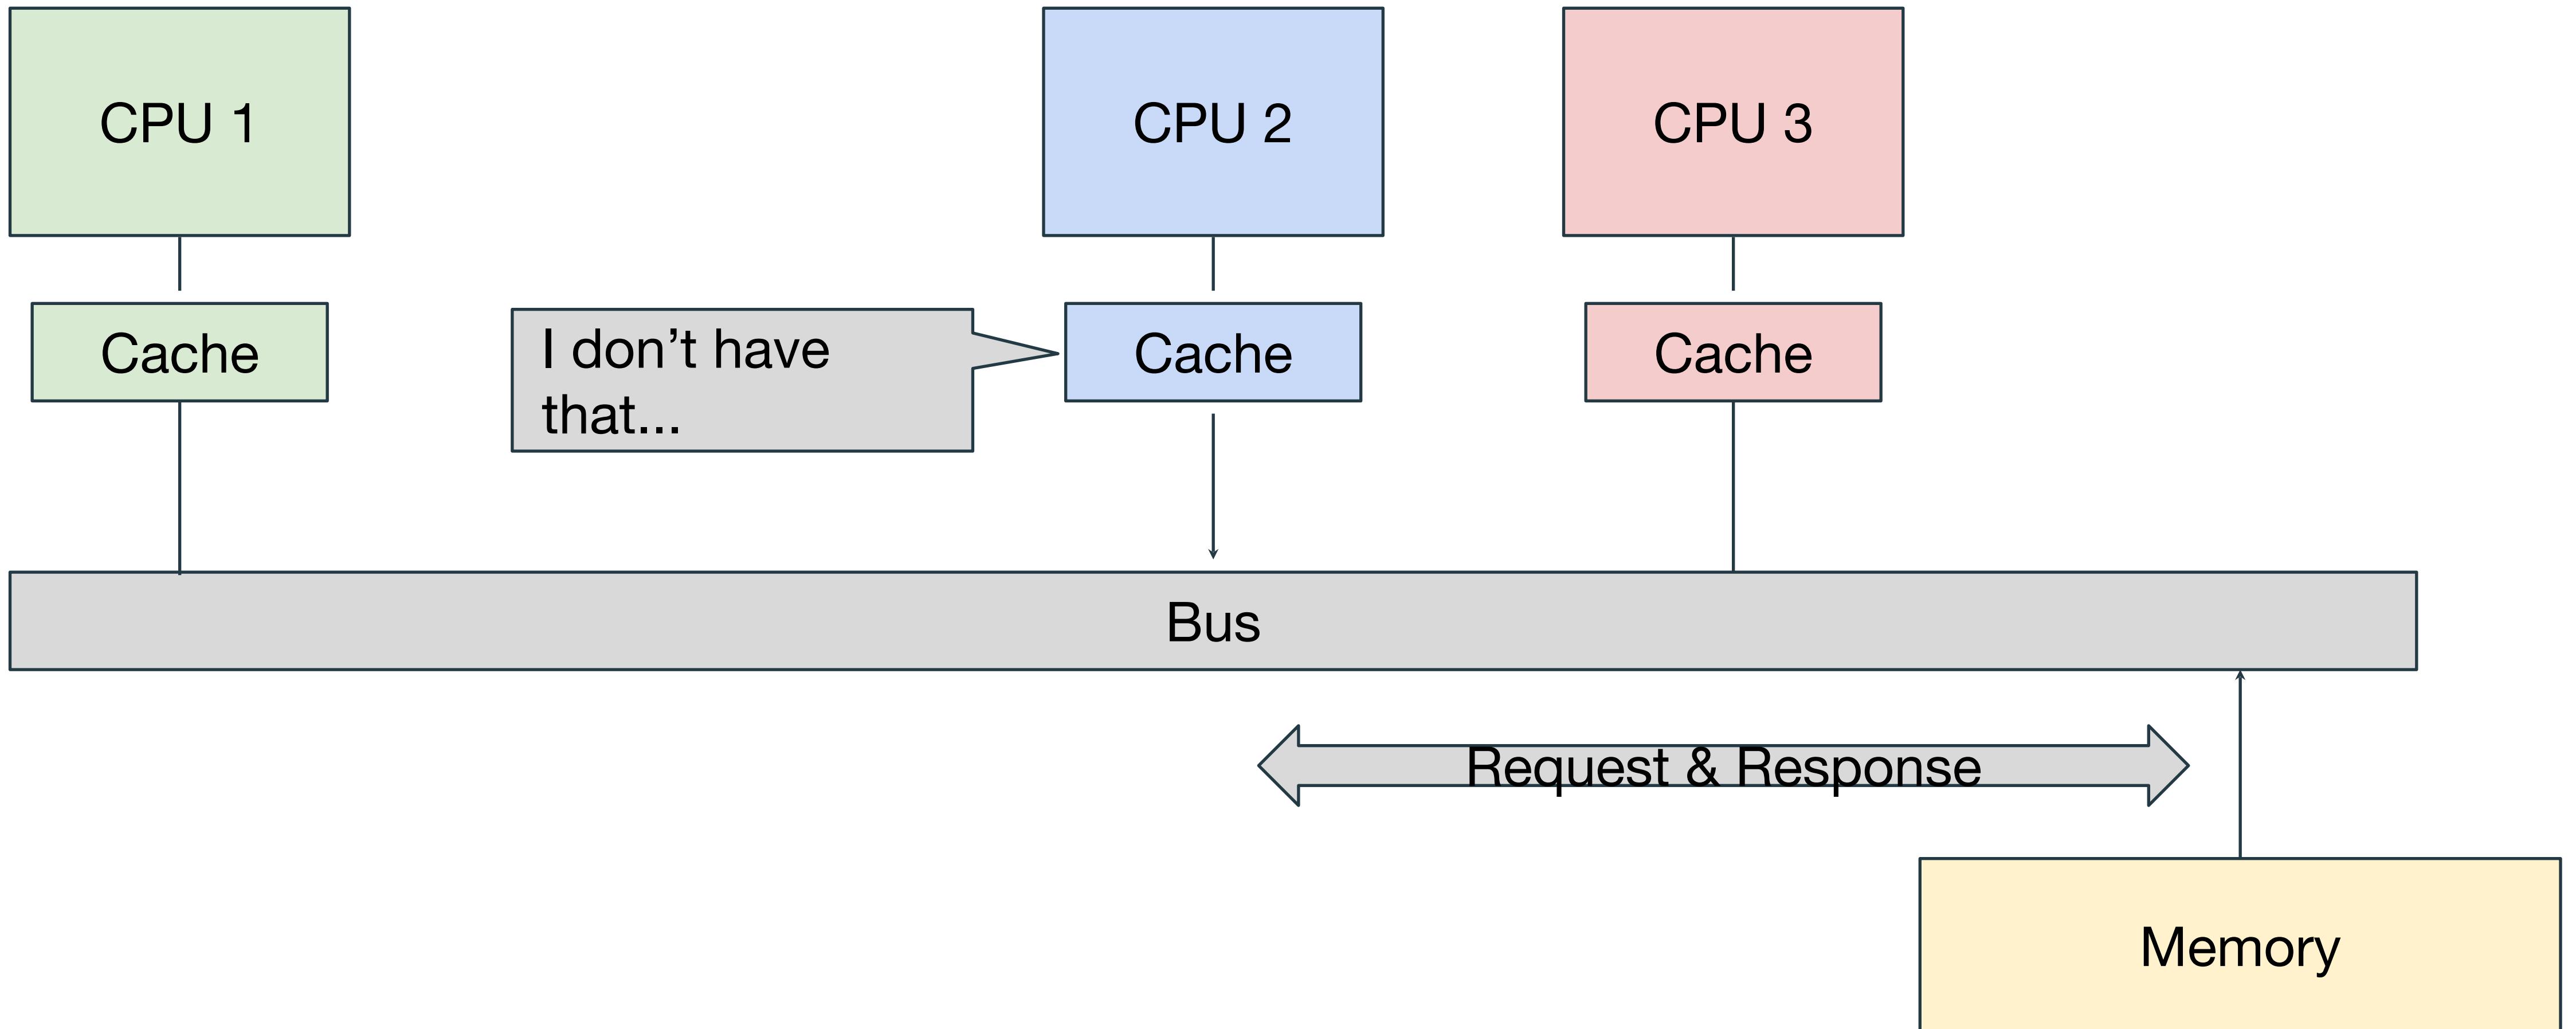
0xC0FFEE = False



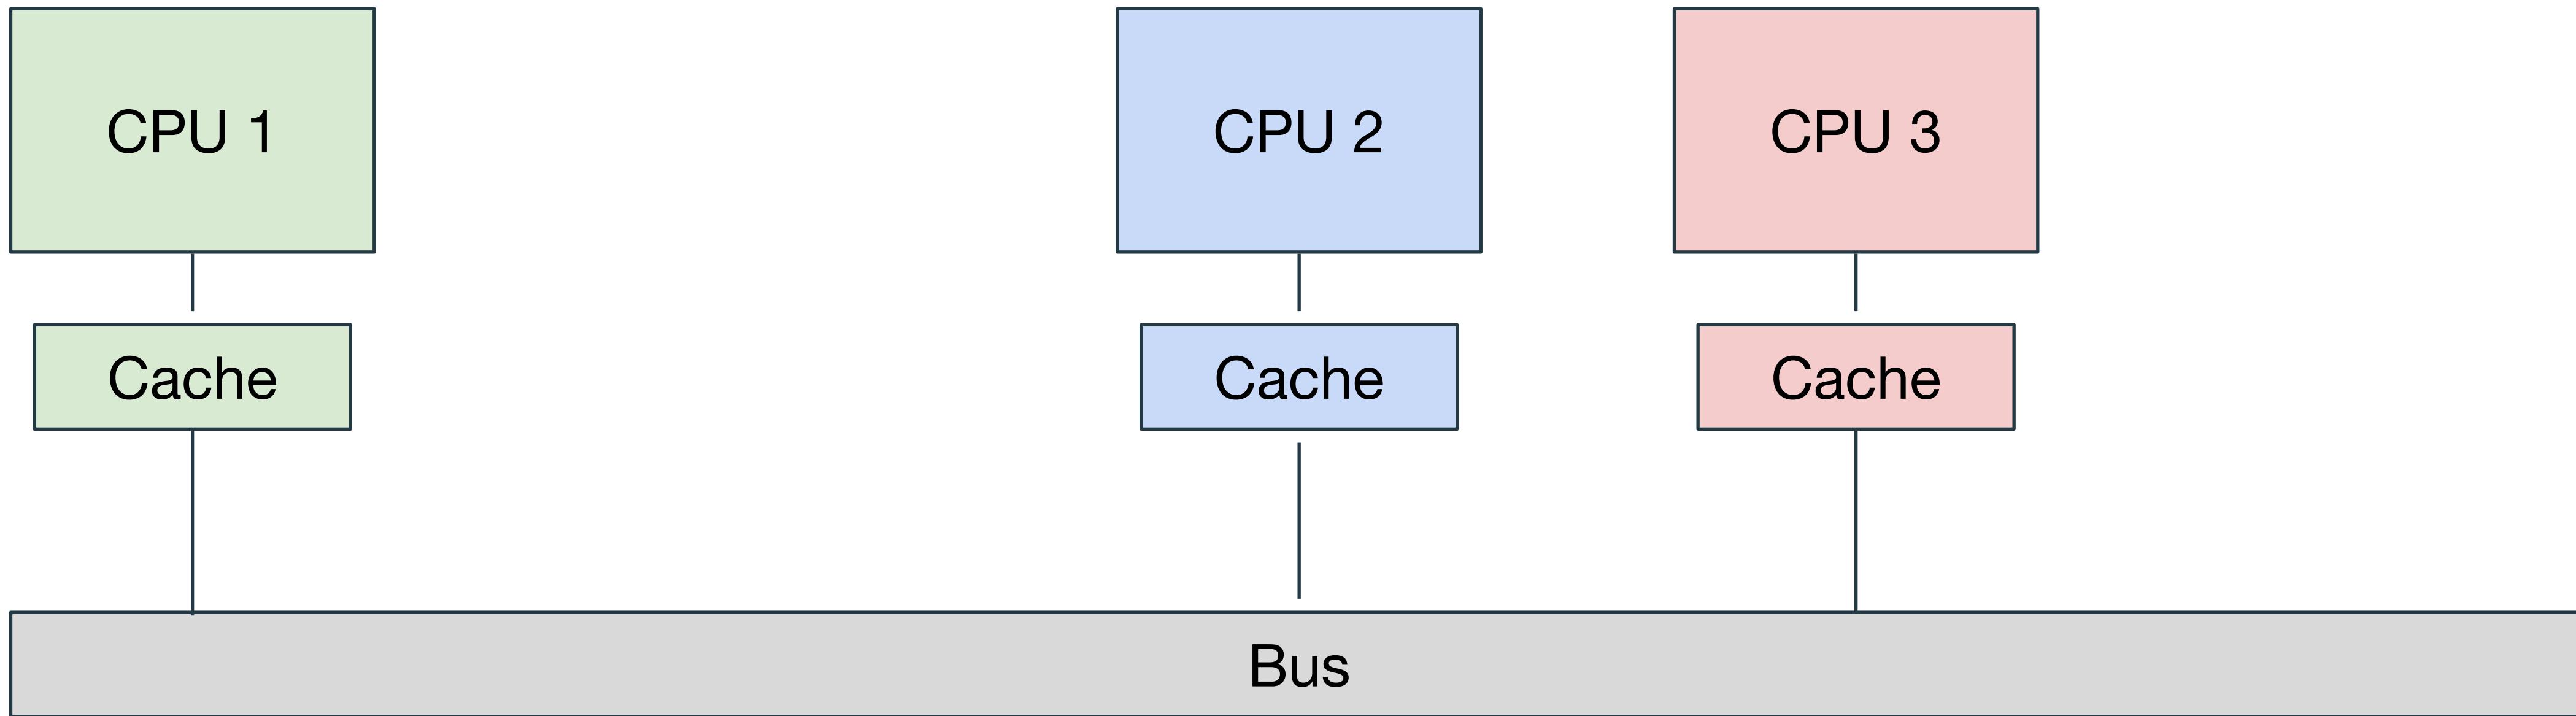
0xC0FFEE = False



0xC0FFEE = False

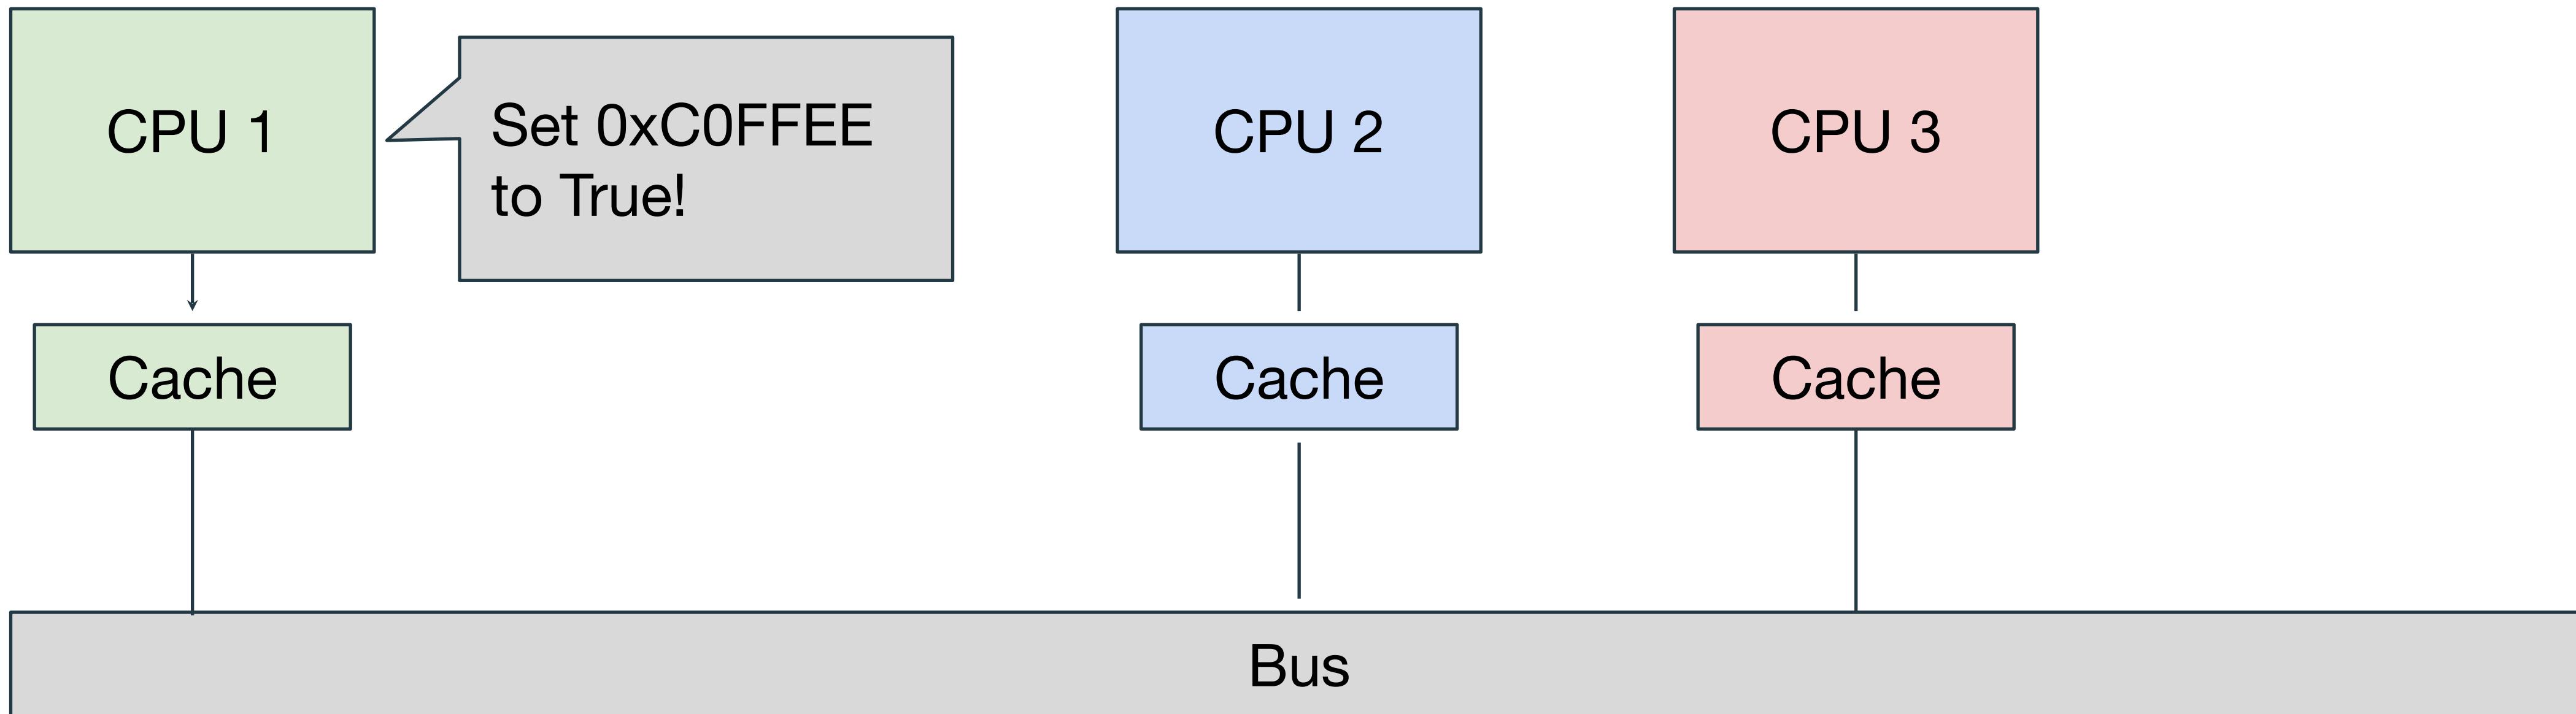


0xC0FFEE = False



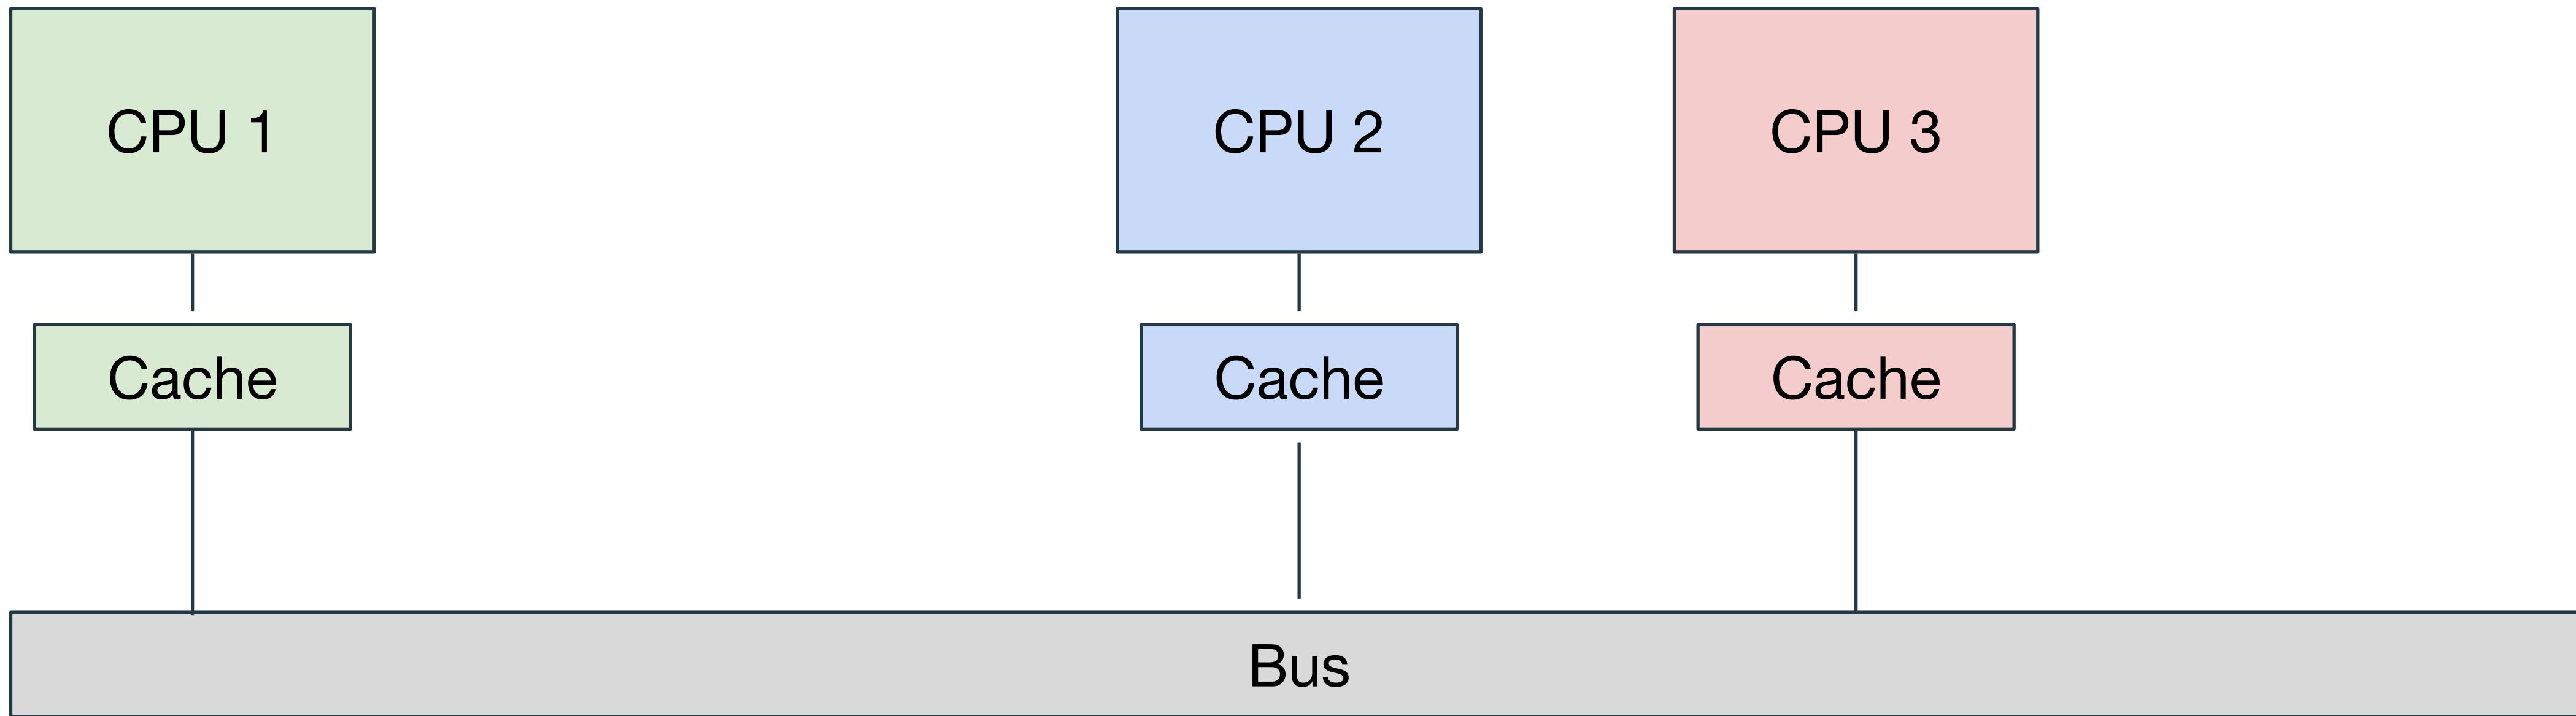
0xC0FFEE = False

0xC0FFEE = False



0xC0FFEE = True

0xC0FFEE = False



Hmm... These can't *both*  
be true! What happened?



# What happened?!

- Two CPU's on a bus (communication device!) both read the same data; each processor got a copy of the data and stored it in their cache.
- Then, one CPU performed a write; their cache is up to date, but the other processor has stale data (and it doesn't know its data is stale!)
- Things to think about:
  - How do we communicate when one processor changes the state of shared data?
    - Does every processor action cause data to change state?
    - Who should be responsible for providing the updated data?
    - What happens to memory while all of this is happening?

# Cat Break!

Computer Science 61C

McMahon & Weaver



# Keeping Multiple Caches Coherent

- Architect's job: shared memory  
=> keep cache values **coherent**
- Idea: When any processor has cache miss or writes, notify other processors via interconnection network
  - If only reading, many processors can have copies
  - If a processor writes, invalidate any other copies
- Write transactions from one processor, other caches “snoop” the common interconnect checking for tags they hold
  - Invalidate any copies of same address modified in other cache

# Basic Requirements: Keeping state & MSI

- Because we have multiple caches, we might have multiple copies of a piece of data floating around in our system (eg. the value of 0xC0FFEE in our previous example).
- We'll need some way to track the status of these copies; do they match what we have in memory? (dirty, or "modified"), do other processors have a copy? Is that copy up to date or stale?
- We'll keep state on a block-by-block basis; remember caches pull data in "blocksize" chunks, so all data within one block should have the same state.
  - We already have "Valid" and "Dirty", so lets add an additional bit for state: shared

# Additional Enhancements:

- We want to use **write-back** caches
  - A write-through cache design uses significantly more memory bandwidth
- We want to minimize writes overall
  - So "write-back" even in the case of shared cache blocks!  
Helps reduce the cost of coherency misses
- We can communicate by broadcasting requests
  - Shout to all other processors
  - The neighboring processor's cache is faster to get to than main memory!

# How Does HW Keep \$ Coherent?

- We already saw how to do this with just *valid* and *dirty*, but we can also think of it this way...
- Each cache tracks state of each *block* in cache:
  1. *Shared*: up-to-date data, other caches may have a copy  
(Valid Bit set, shared bit set)
  2. *Modified*: up-to-date data, changed (dirty), no other cache has a copy, OK to write, memory out-of-date (Valid and Dirty bit is set)

# Two Optional Performance Optimizations of Cache Coherency via New States

3. *Exclusive*: up-to-date data, no other cache has a copy, OK to write, memory up-to-date

- If any other cache reads this line the state becomes **shared**
  - Can provide the line if I'm faster than main memory
- If I write to this line, state becomes **modified** but I don't need to broadcast this when I do the write
- Only valid is set

4. *Owner*: up-to-date data, other caches may have a copy (they must be in Shared state), memory is **not up-to-date**

- But this cache now supplies data on read instead of going to memory:  
Saves the need for a write back when somebody else reads the line
- And now when you write you once again have to have the other caches invalidate
- Valid, dirty, and shared is set

# Name of This Common Cache Coherency Protocol: MOESI

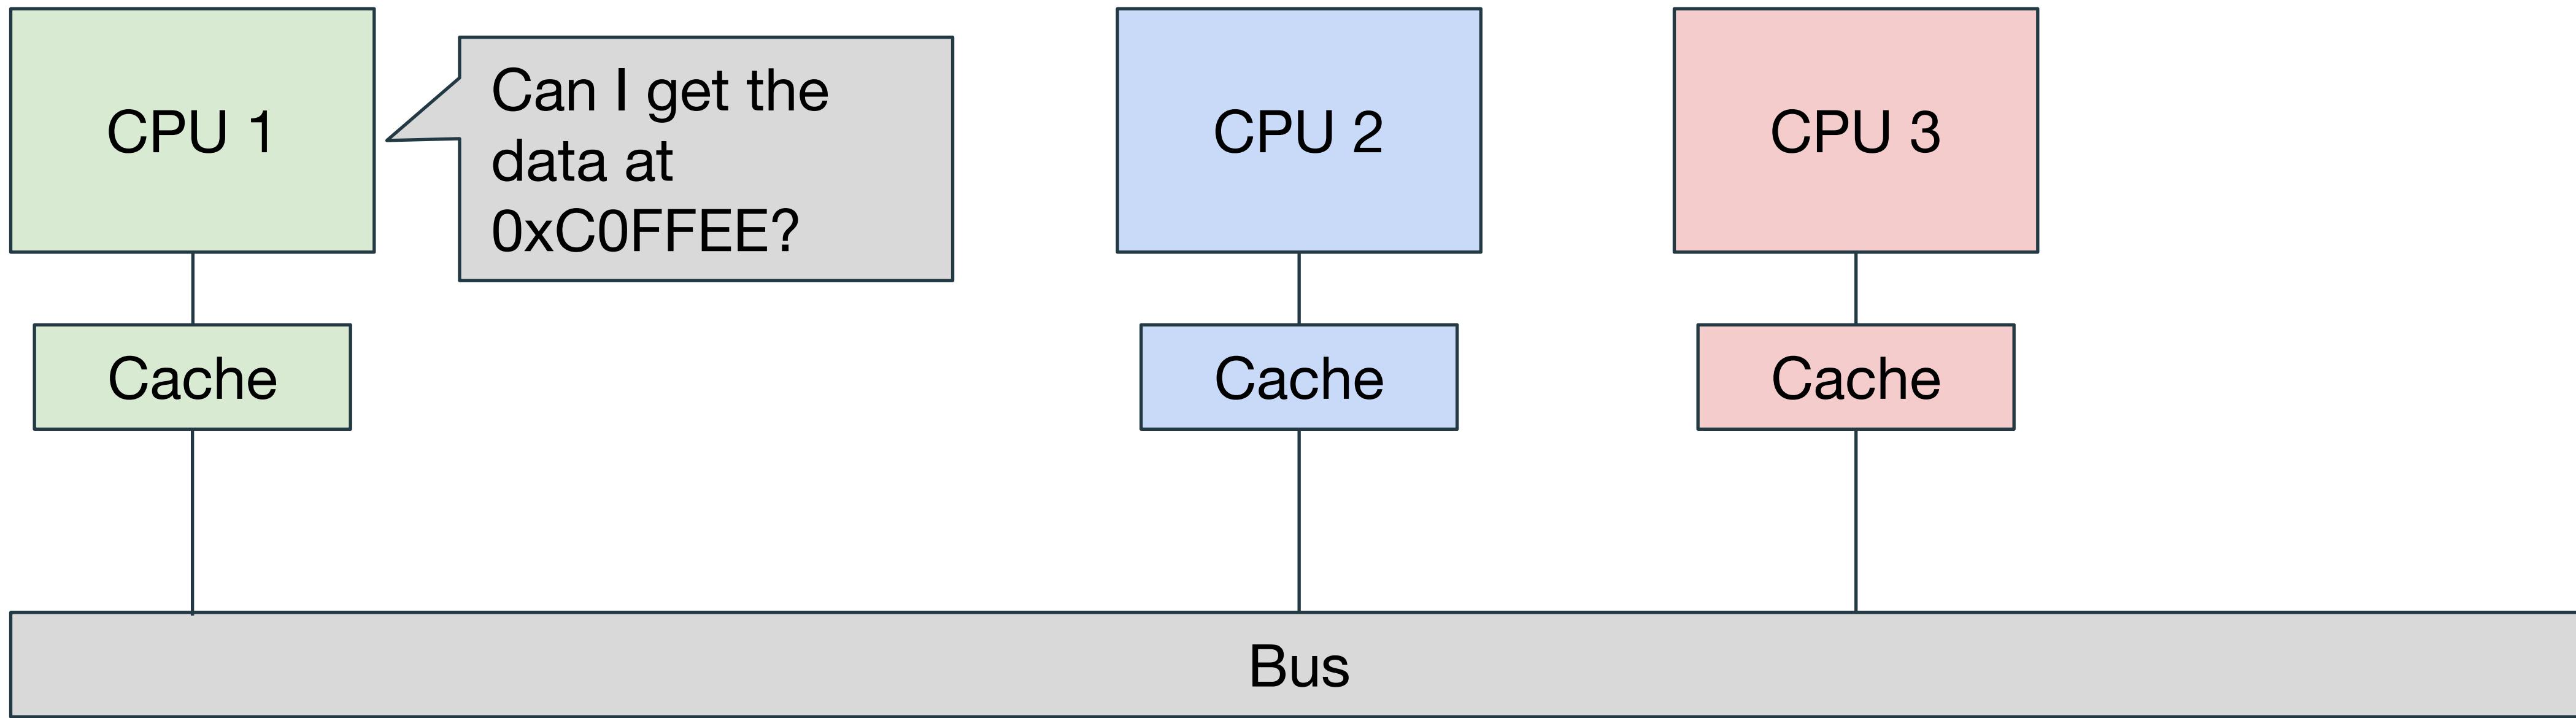
- Memory access to cache is either
  - Modified (in cache)
  - Owned (in cache)
  - Exclusive (in cache)
  - Shared (in cache)
  - Invalid (not in cache)



Snooping/Snoopy Protocols  
e.g., the Berkeley Ownership Protocol  
See [http://en.wikipedia.org/wiki/Cache\\_coherence](http://en.wikipedia.org/wiki/Cache_coherence)  
Berkeley Protocol is a wikipedia stub!

# Idea Behind States

- M: Modified
  - I have the only copy, and can write, and its dirty
  - If this gets evicted, I need to flush the entry
- O: Owned
  - I have the ***official*** copy, and can write, and its dirty
  - When writing, I have to tell everyone else I'm writing (and it now turns Modified)
- E: Exclusive
  - I have the ***only*** copy
- S: Shared
  - I have a copy and can read away...
- I: Invalid (duh)



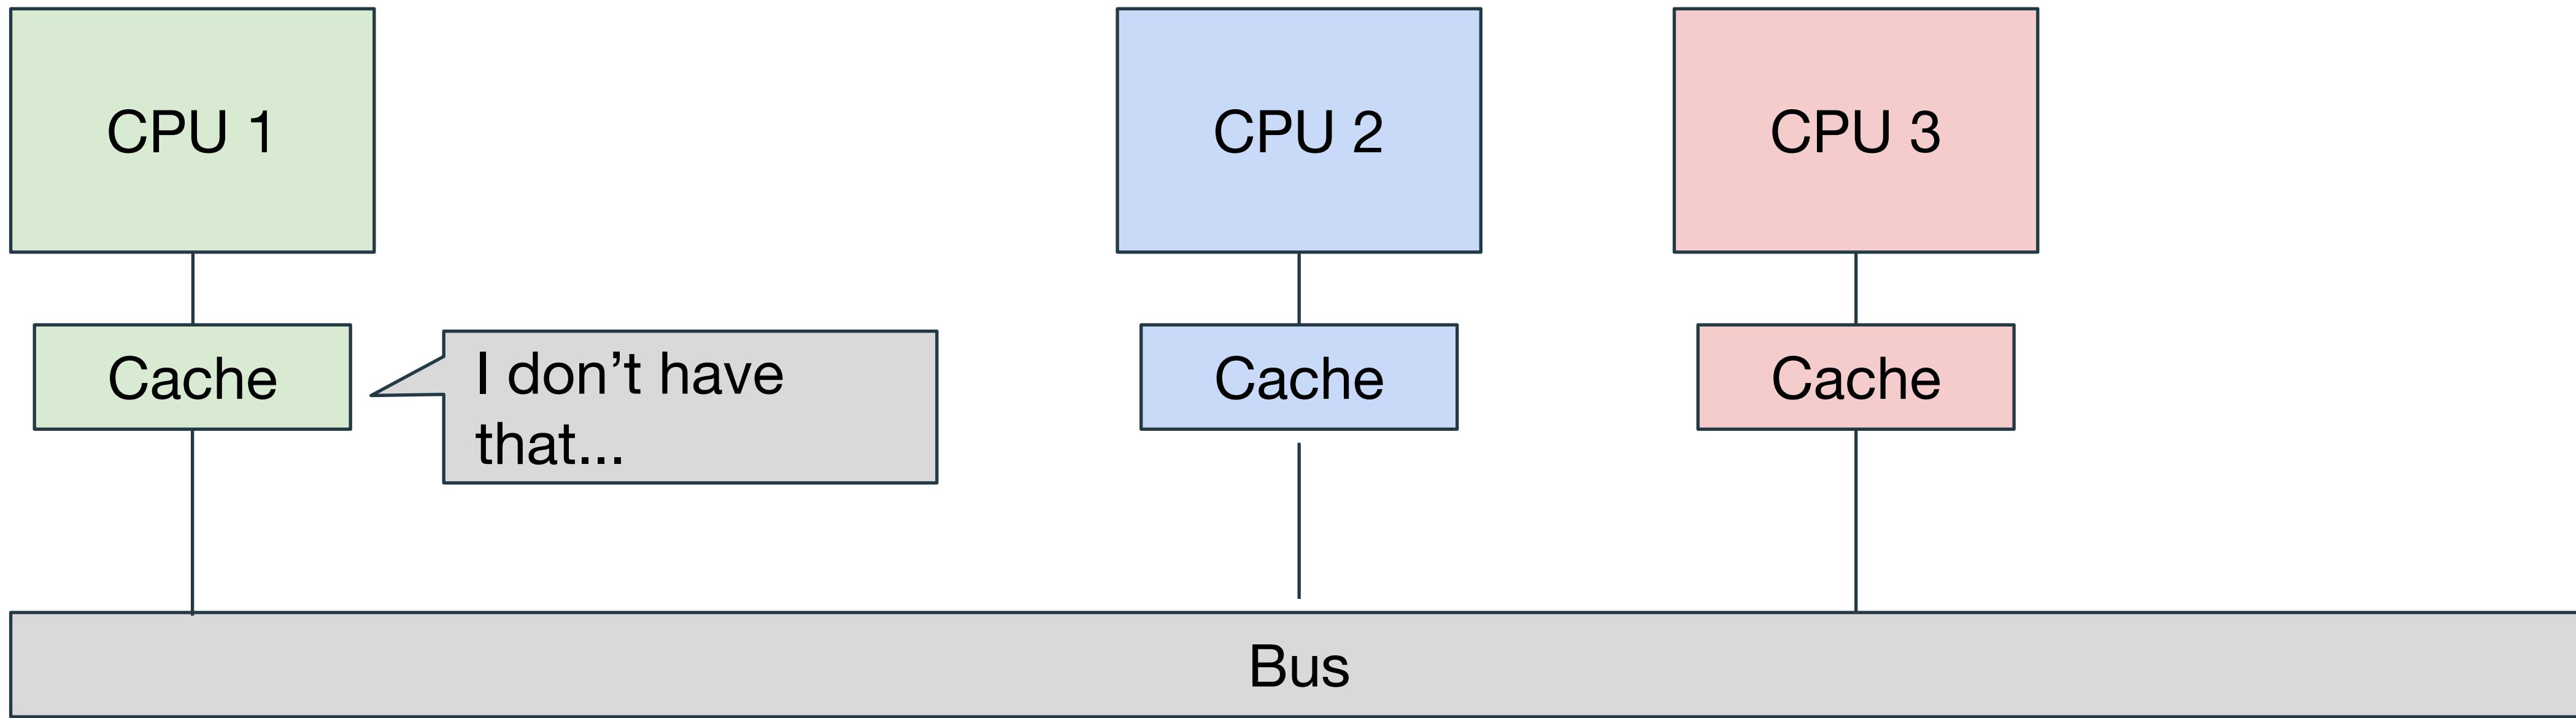
| Valid | Dirty | State |
|-------|-------|-------|
| 0     | 0     | I     |

| Valid | Dirty | State |
|-------|-------|-------|
| 0     | 0     | I     |



0xC0FFEE = ???

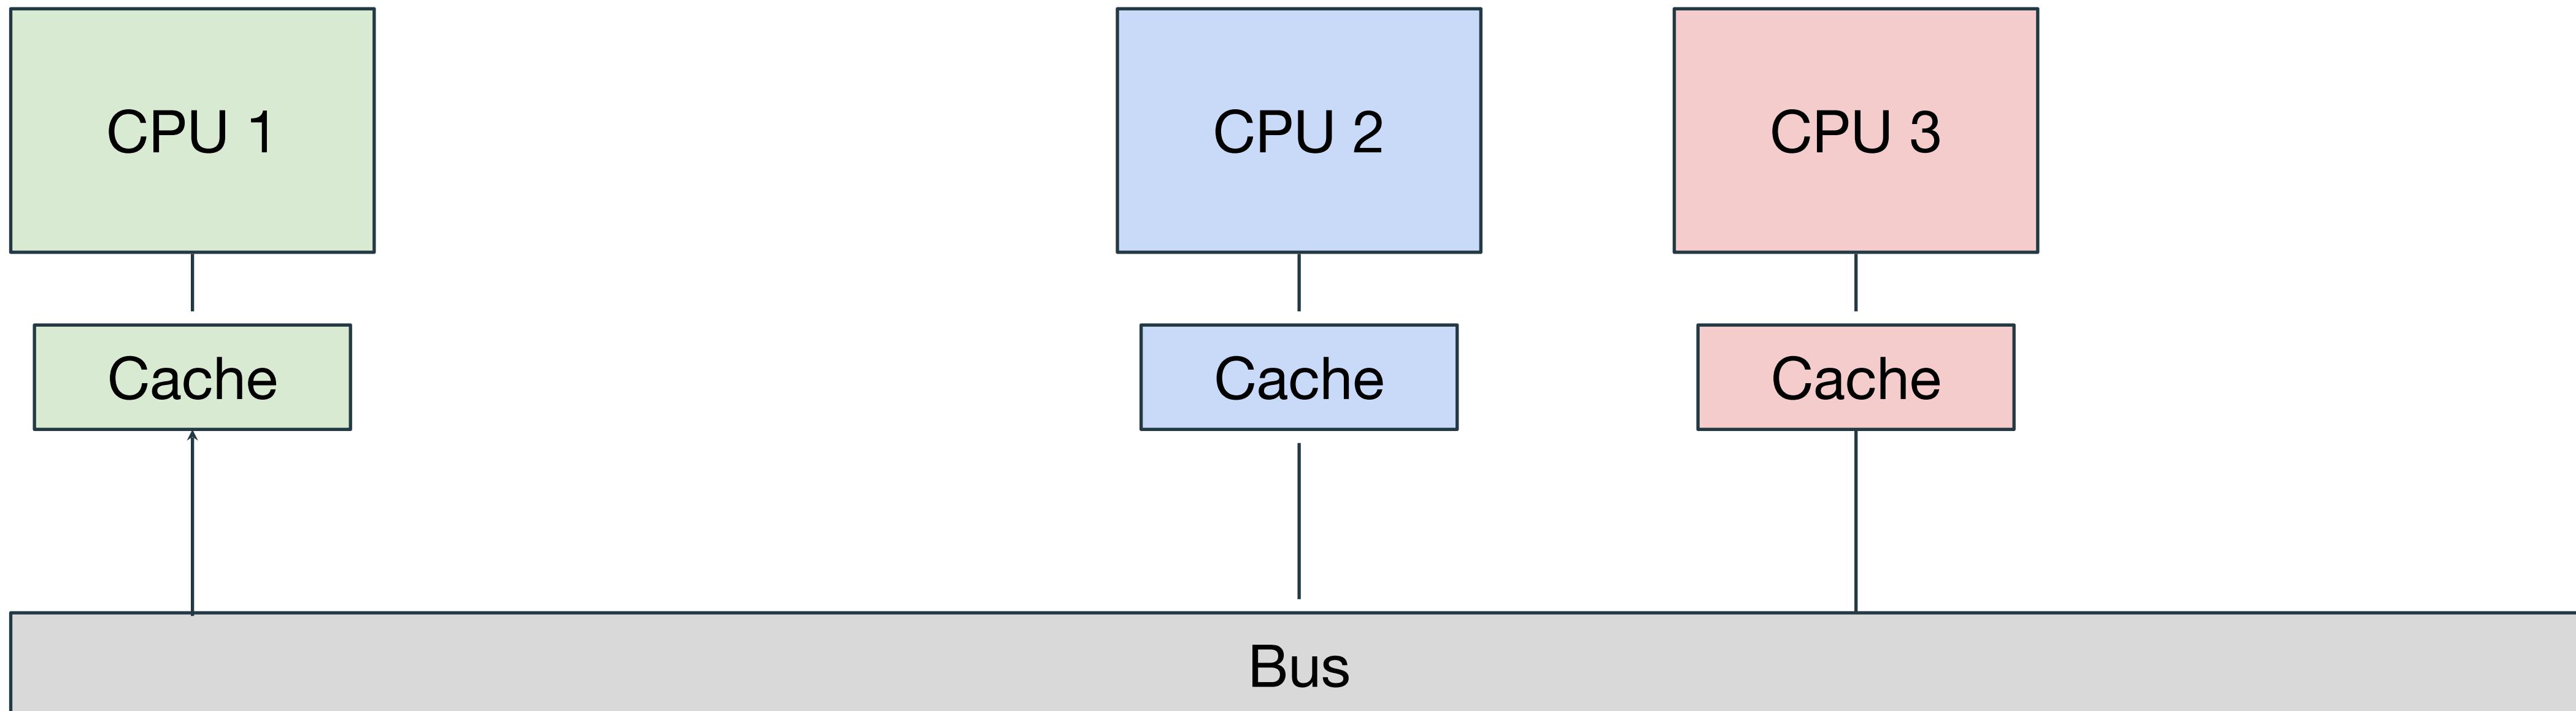
0xC0FFEE = ???



| Valid | Dirty | State |
|-------|-------|-------|
| 0     | 0     | I     |

| Valid | Dirty | State |
|-------|-------|-------|
| 0     | 0     | I     |



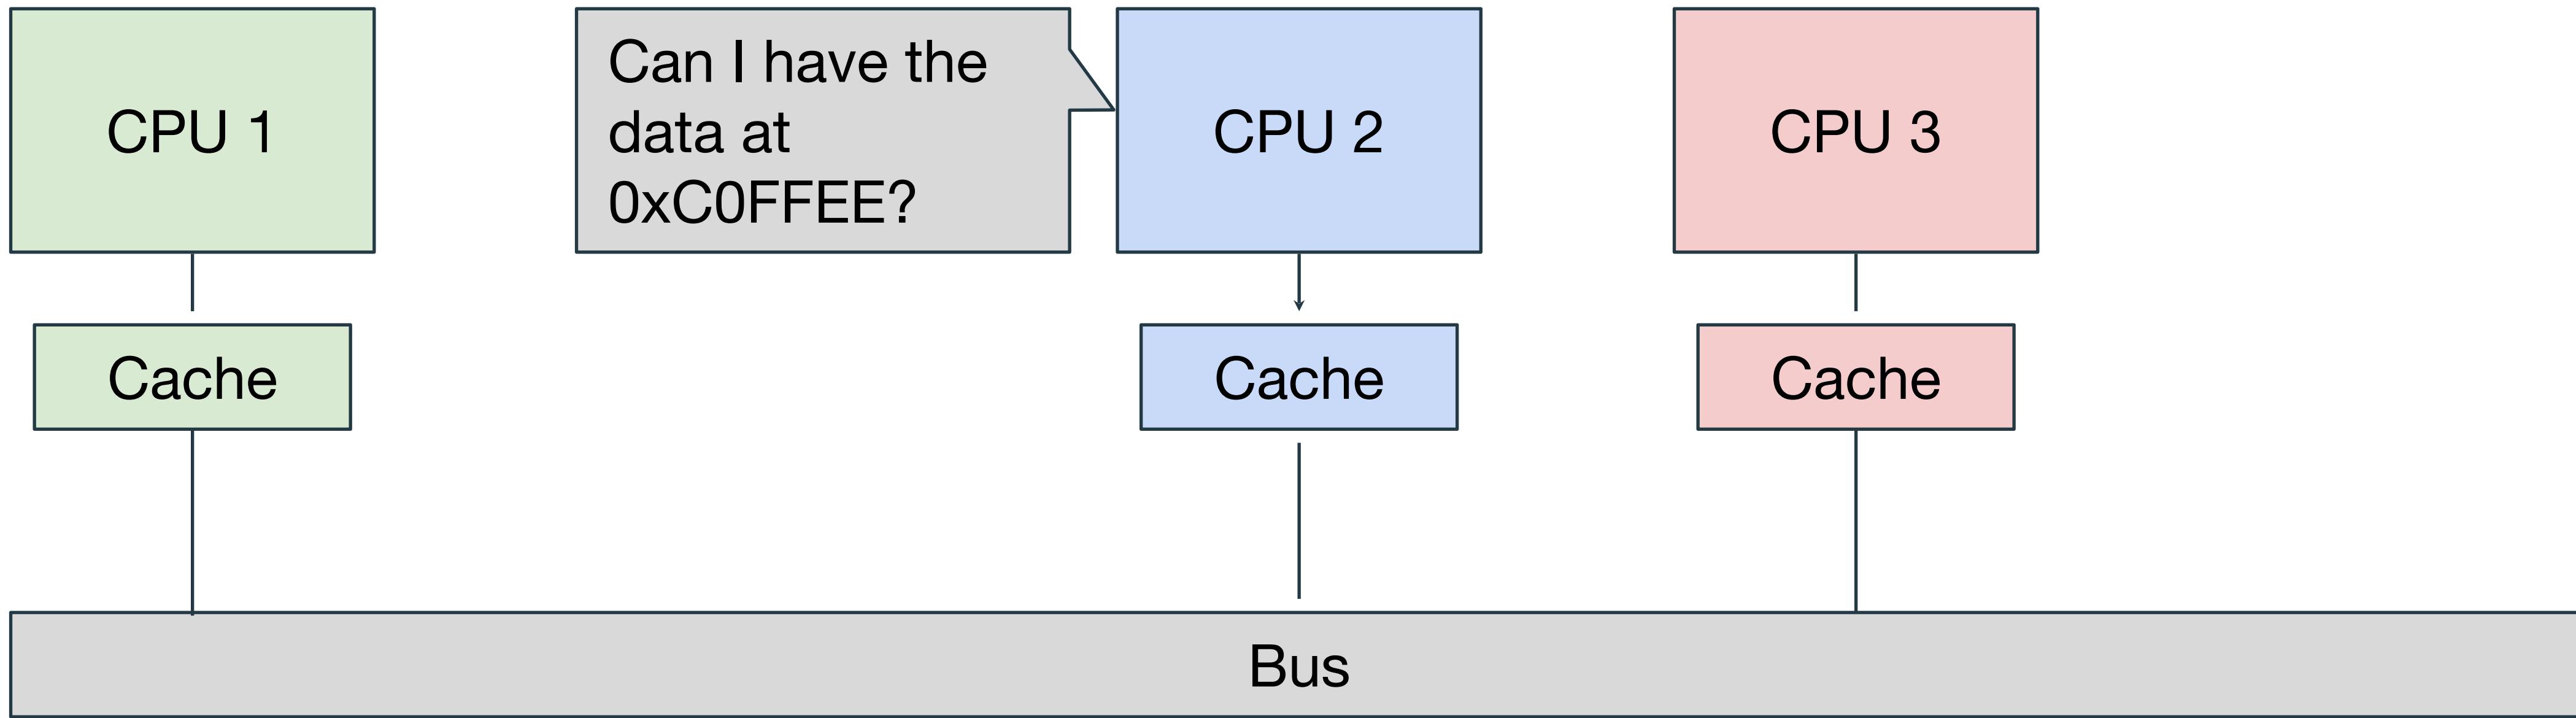


| Valid | Dirty | State |
|-------|-------|-------|
| 1     | 0     | E     |

| Valid | Dirty | State |
|-------|-------|-------|
| 0     | 0     | I     |

0xC0FFEE = False

0xC0FFEE = ???

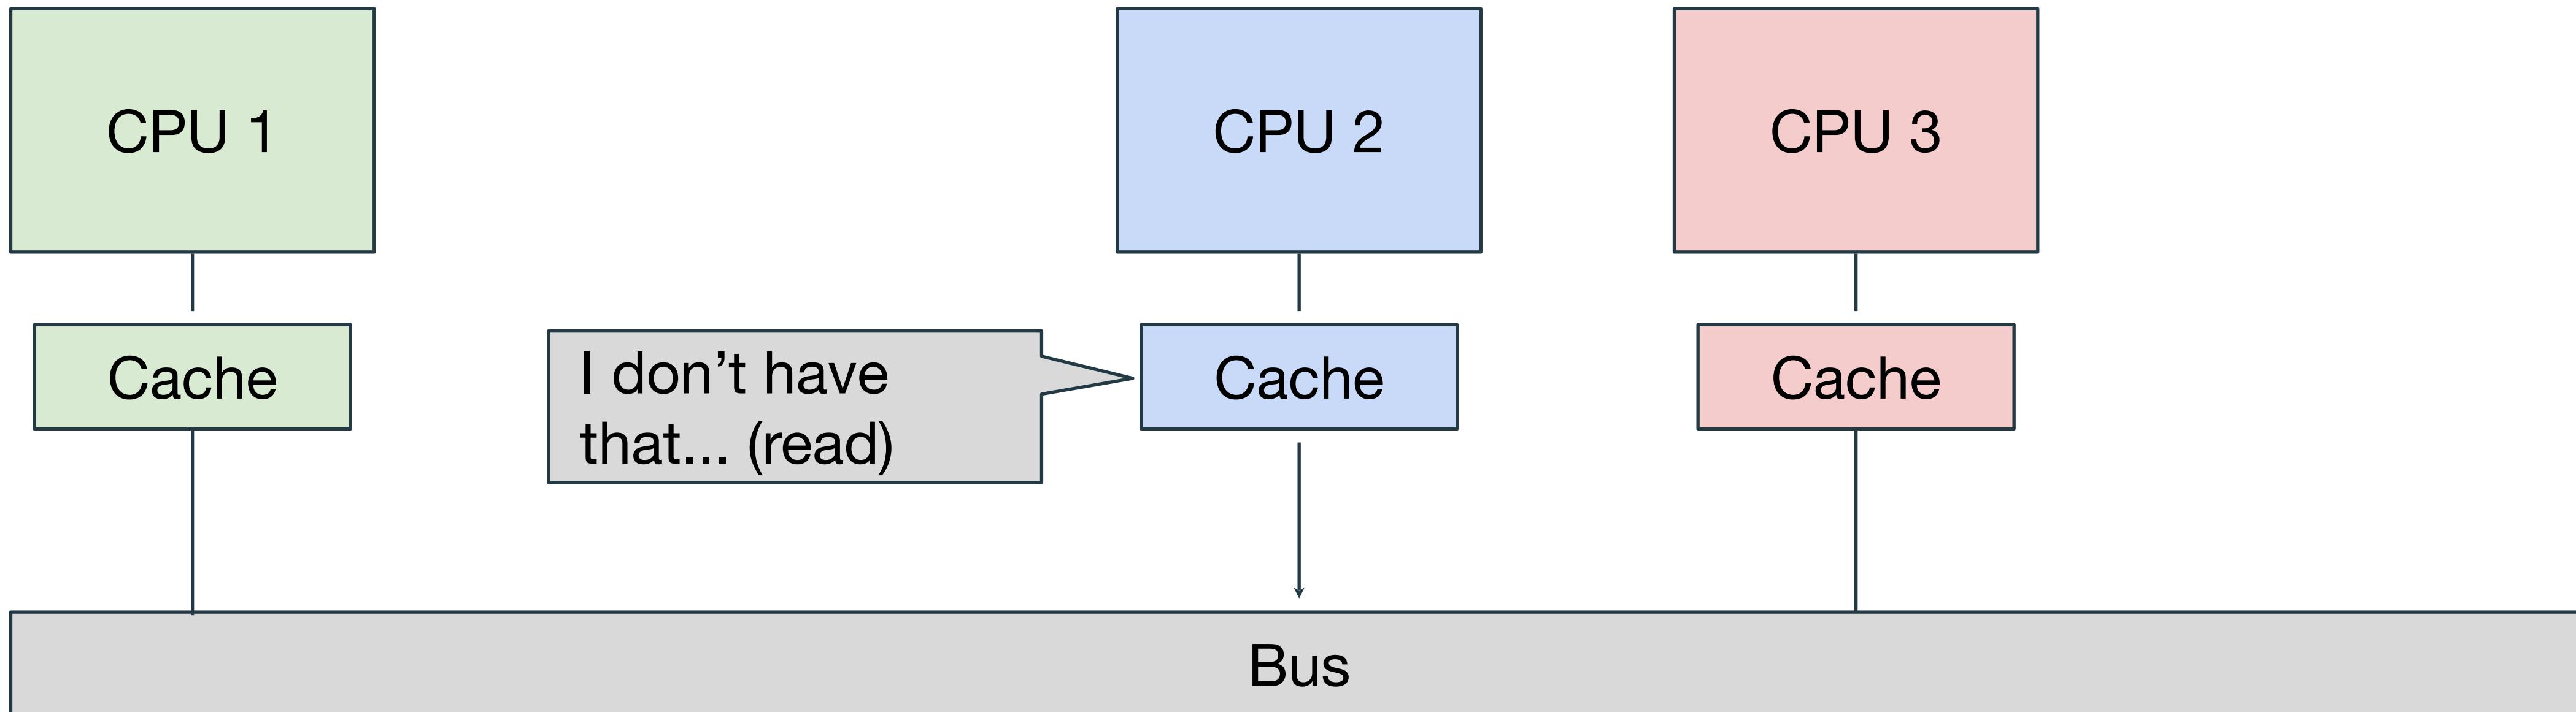


| Valid | Dirty | State |
|-------|-------|-------|
| 1     | 0     | E     |

| Valid | Dirty | State |
|-------|-------|-------|
| 0     | 0     | I     |

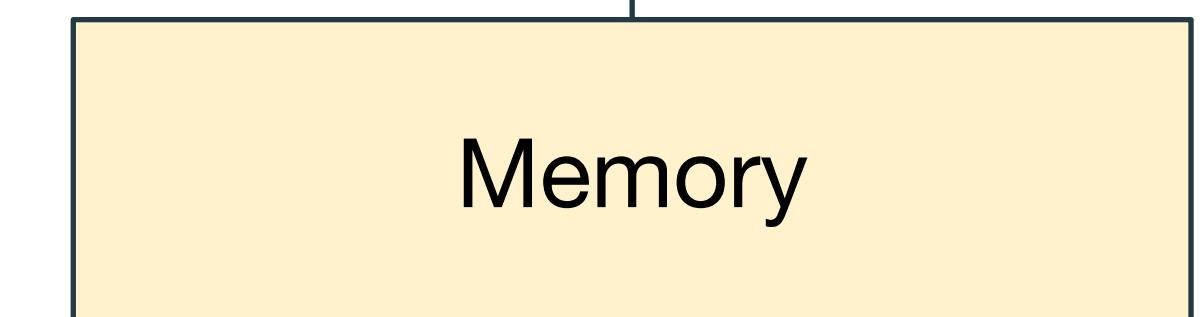
0xC0FFEE = False

0xC0FFEE = ???



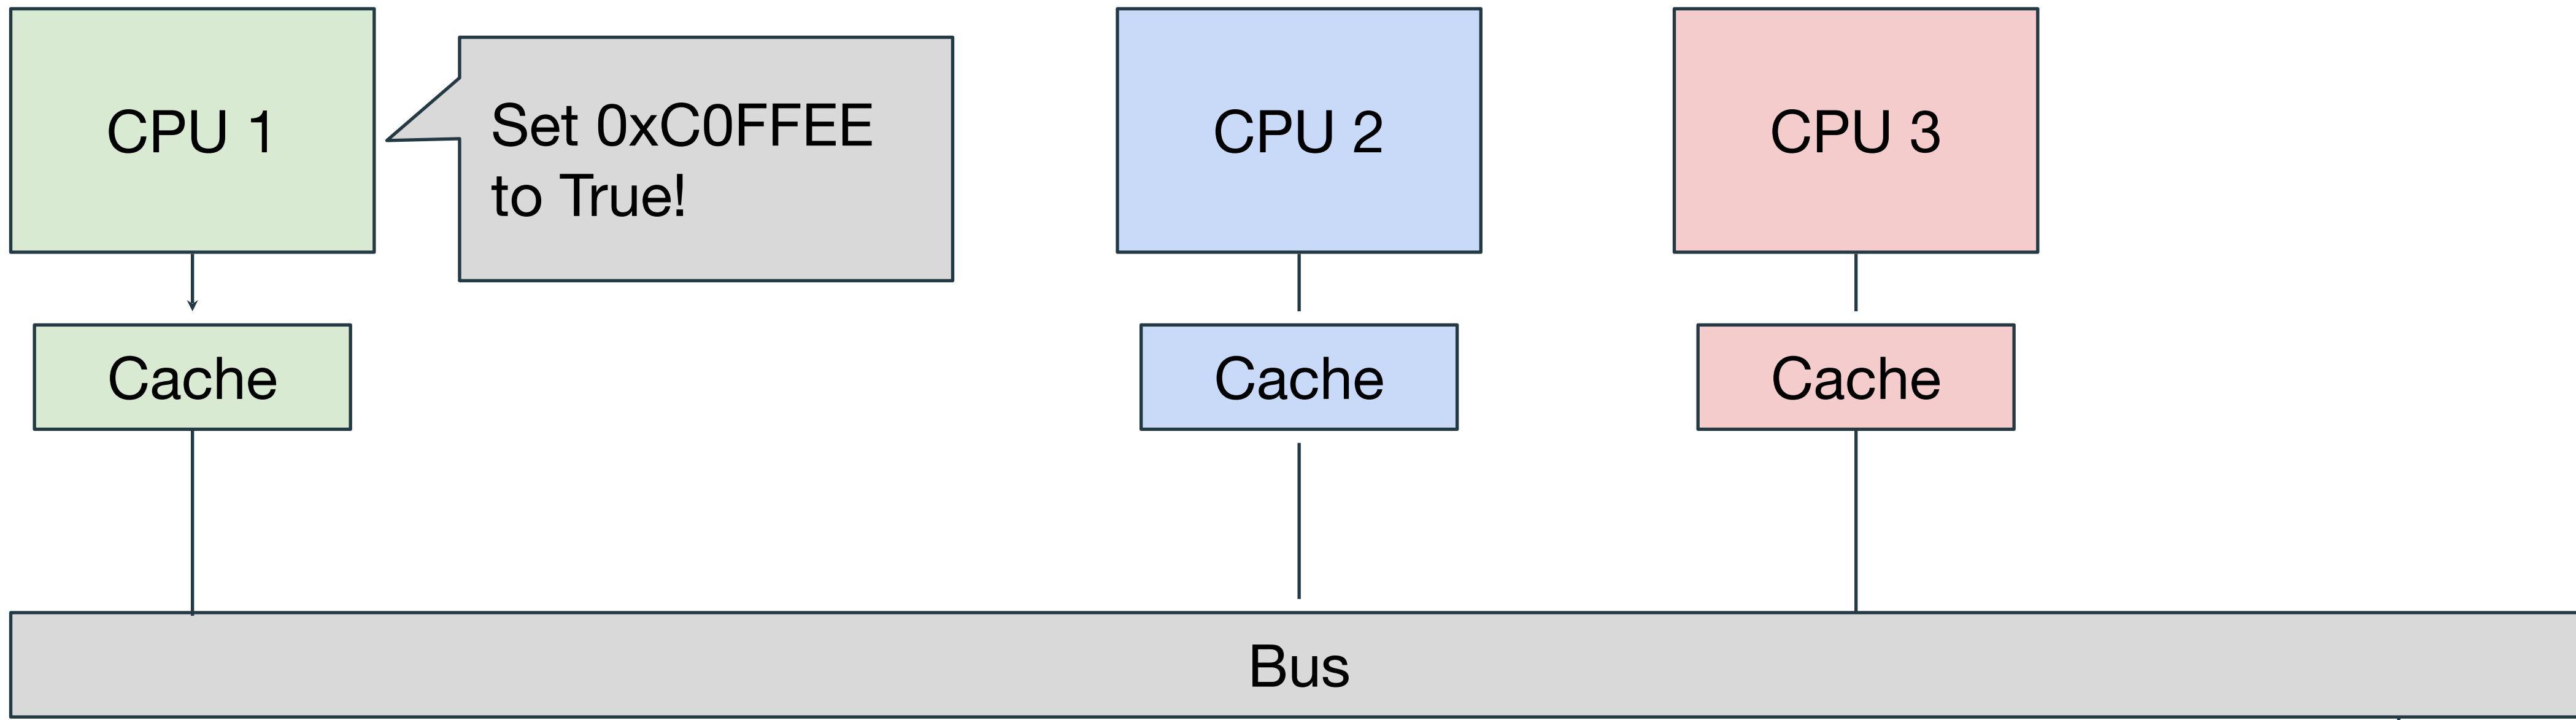
| Valid | Dirty | State |
|-------|-------|-------|
| 1     | 0     | S     |

| Valid | Dirty | State |
|-------|-------|-------|
| 1     | 0     | S     |



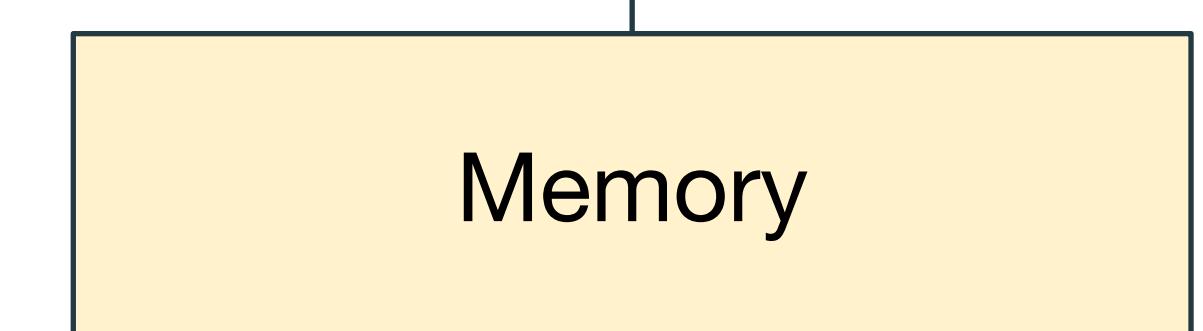
0xC0FFEE = False

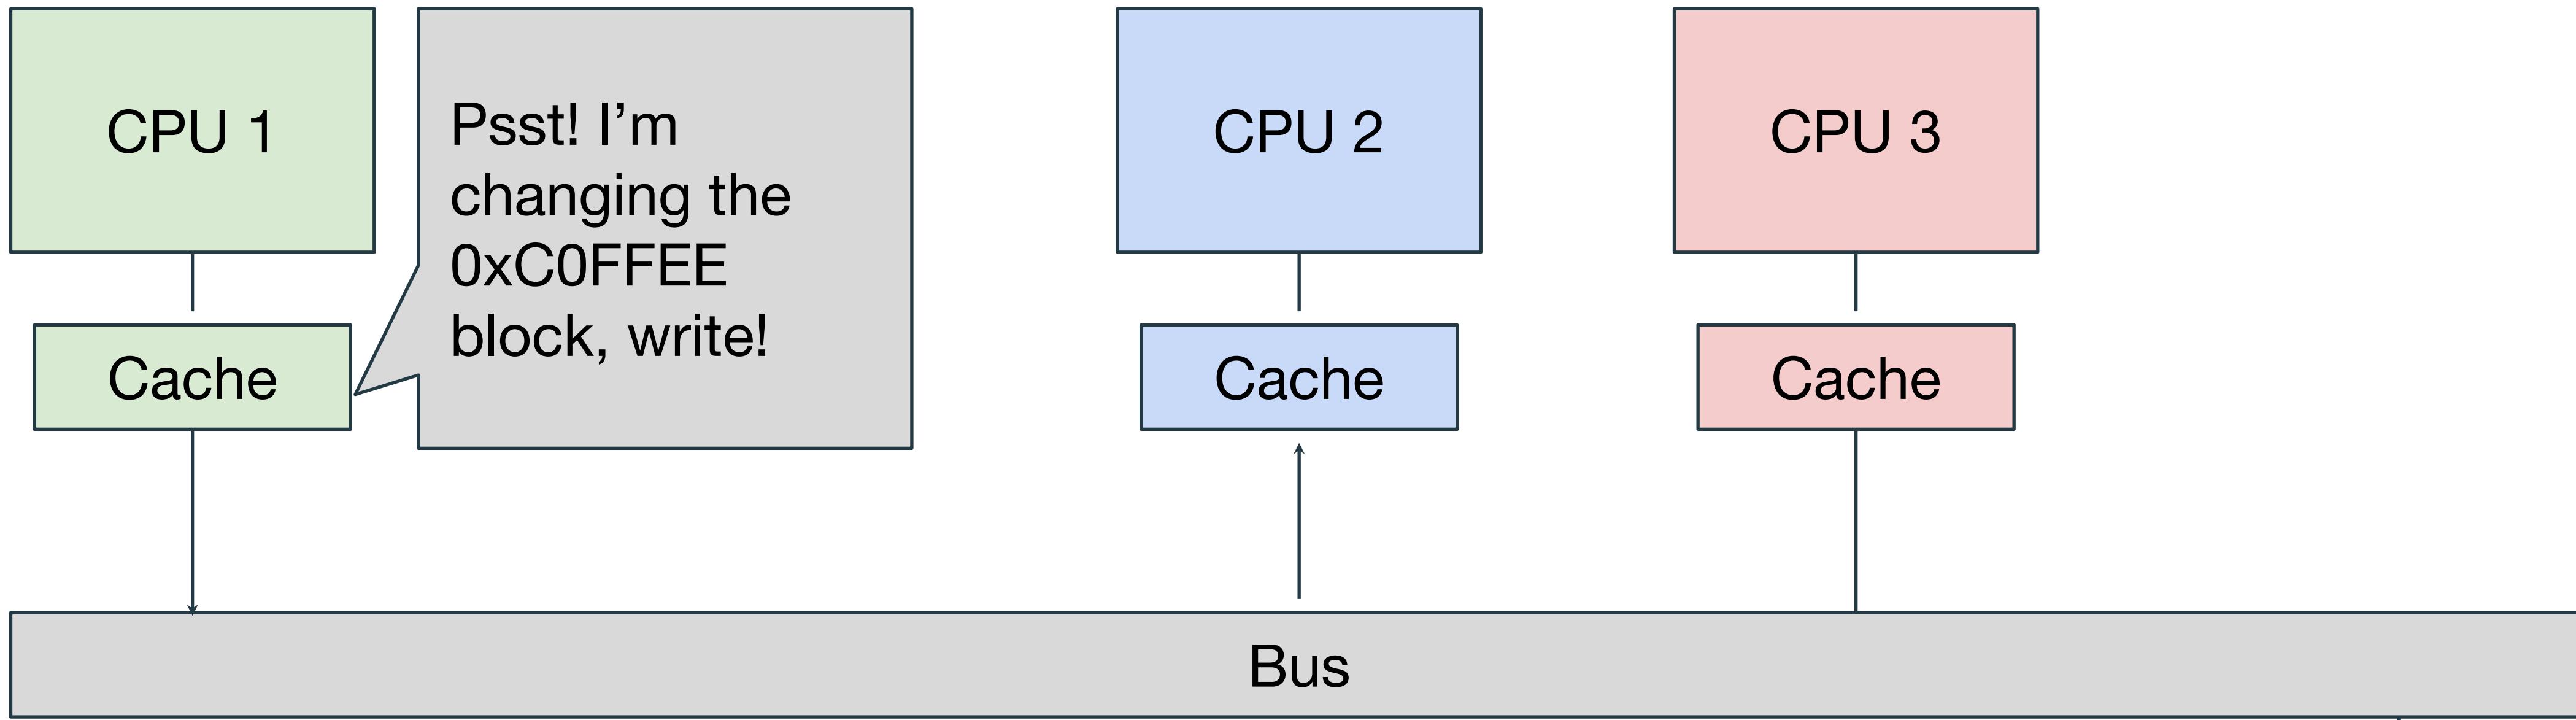
0xC0FFEE = False



| Valid | Dirty | State |
|-------|-------|-------|
| 1     | 0     | S     |

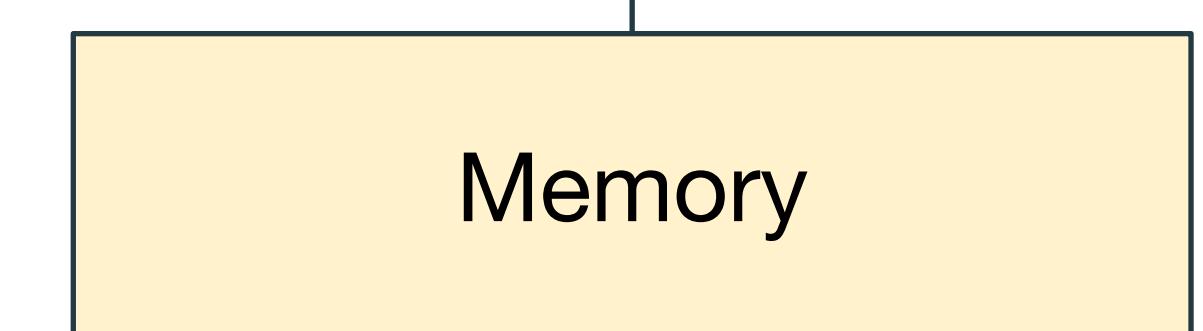
| Valid | Dirty | State |
|-------|-------|-------|
| ?     | 0     | ?     |

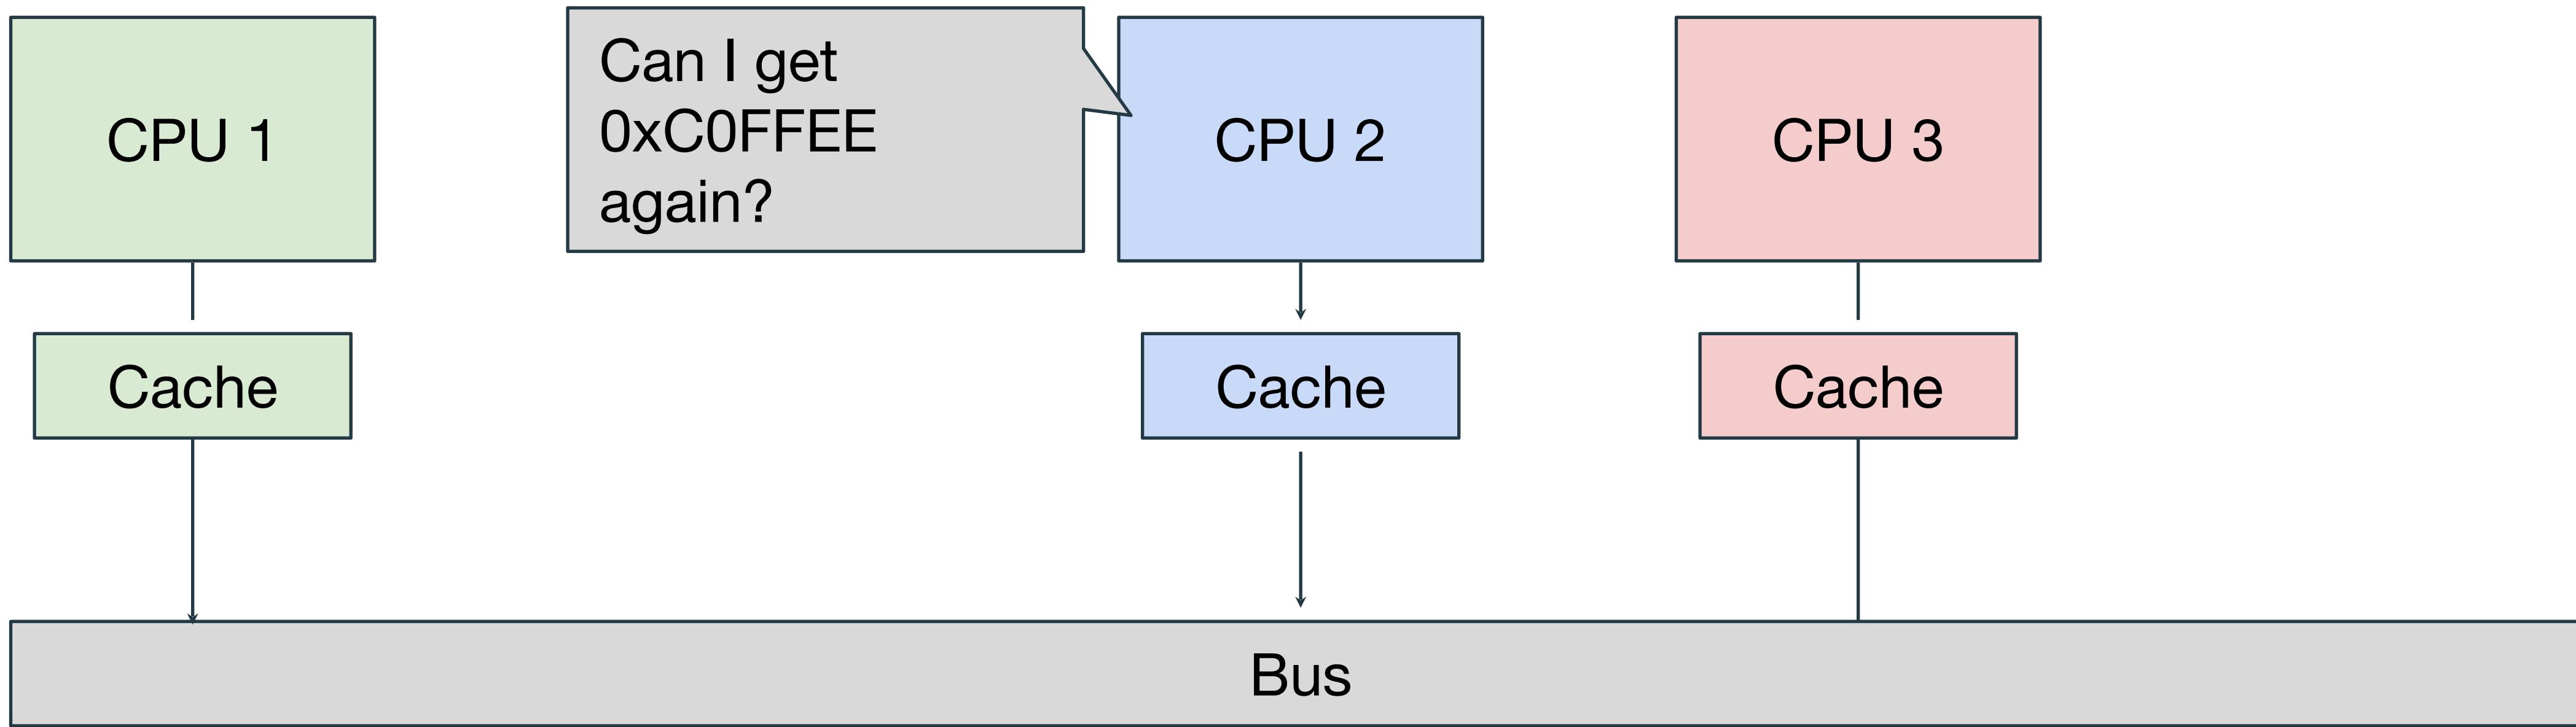




| Valid | Dirty | State |
|-------|-------|-------|
| 1     | 1     | M     |

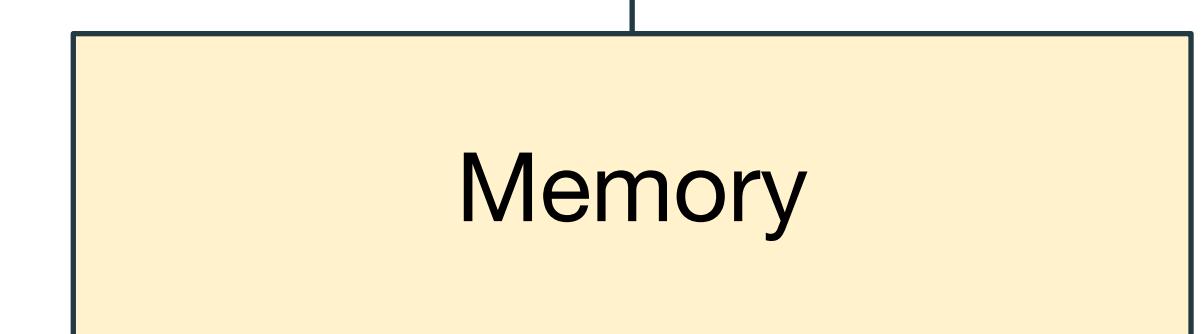
| Valid | Dirty | State |
|-------|-------|-------|
| 0     | 0     | I     |





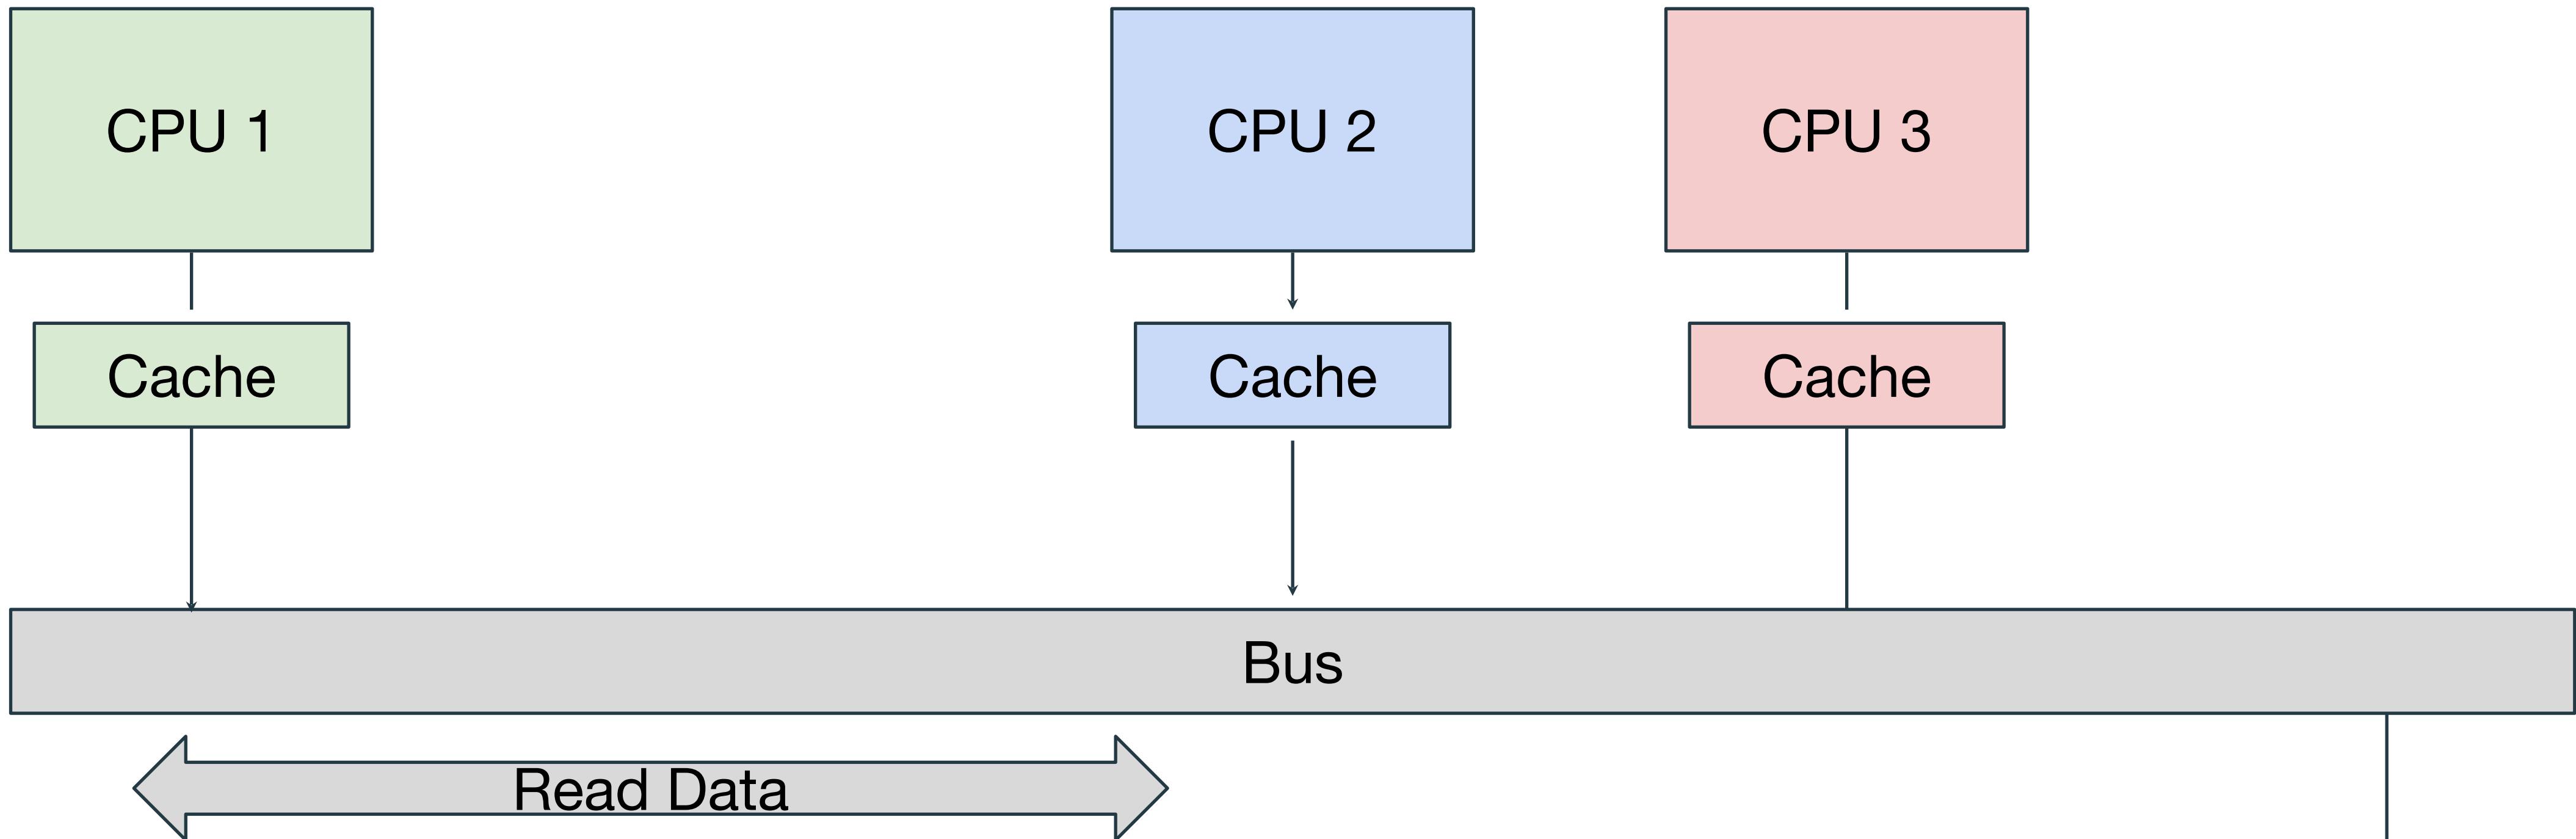
| Valid | Dirty | State |
|-------|-------|-------|
| 1     | 1     | M     |

| Valid | Dirty | State |
|-------|-------|-------|
| 0     | 0     | I     |



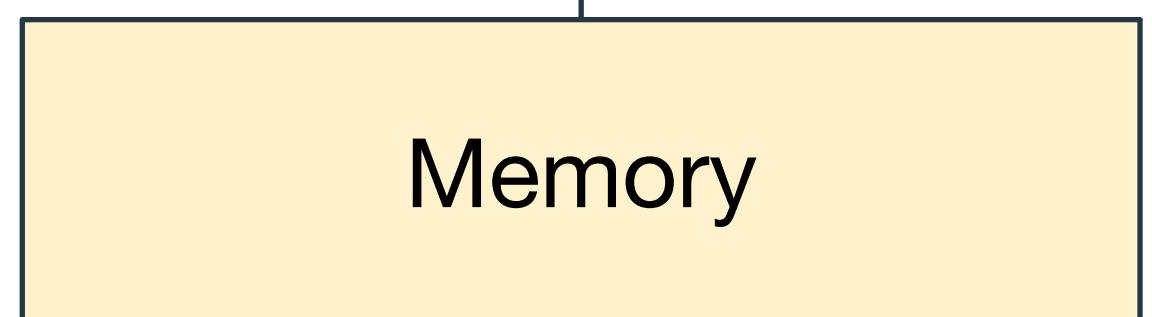
0xC0FFEE = True

0xC0FFEE = ???



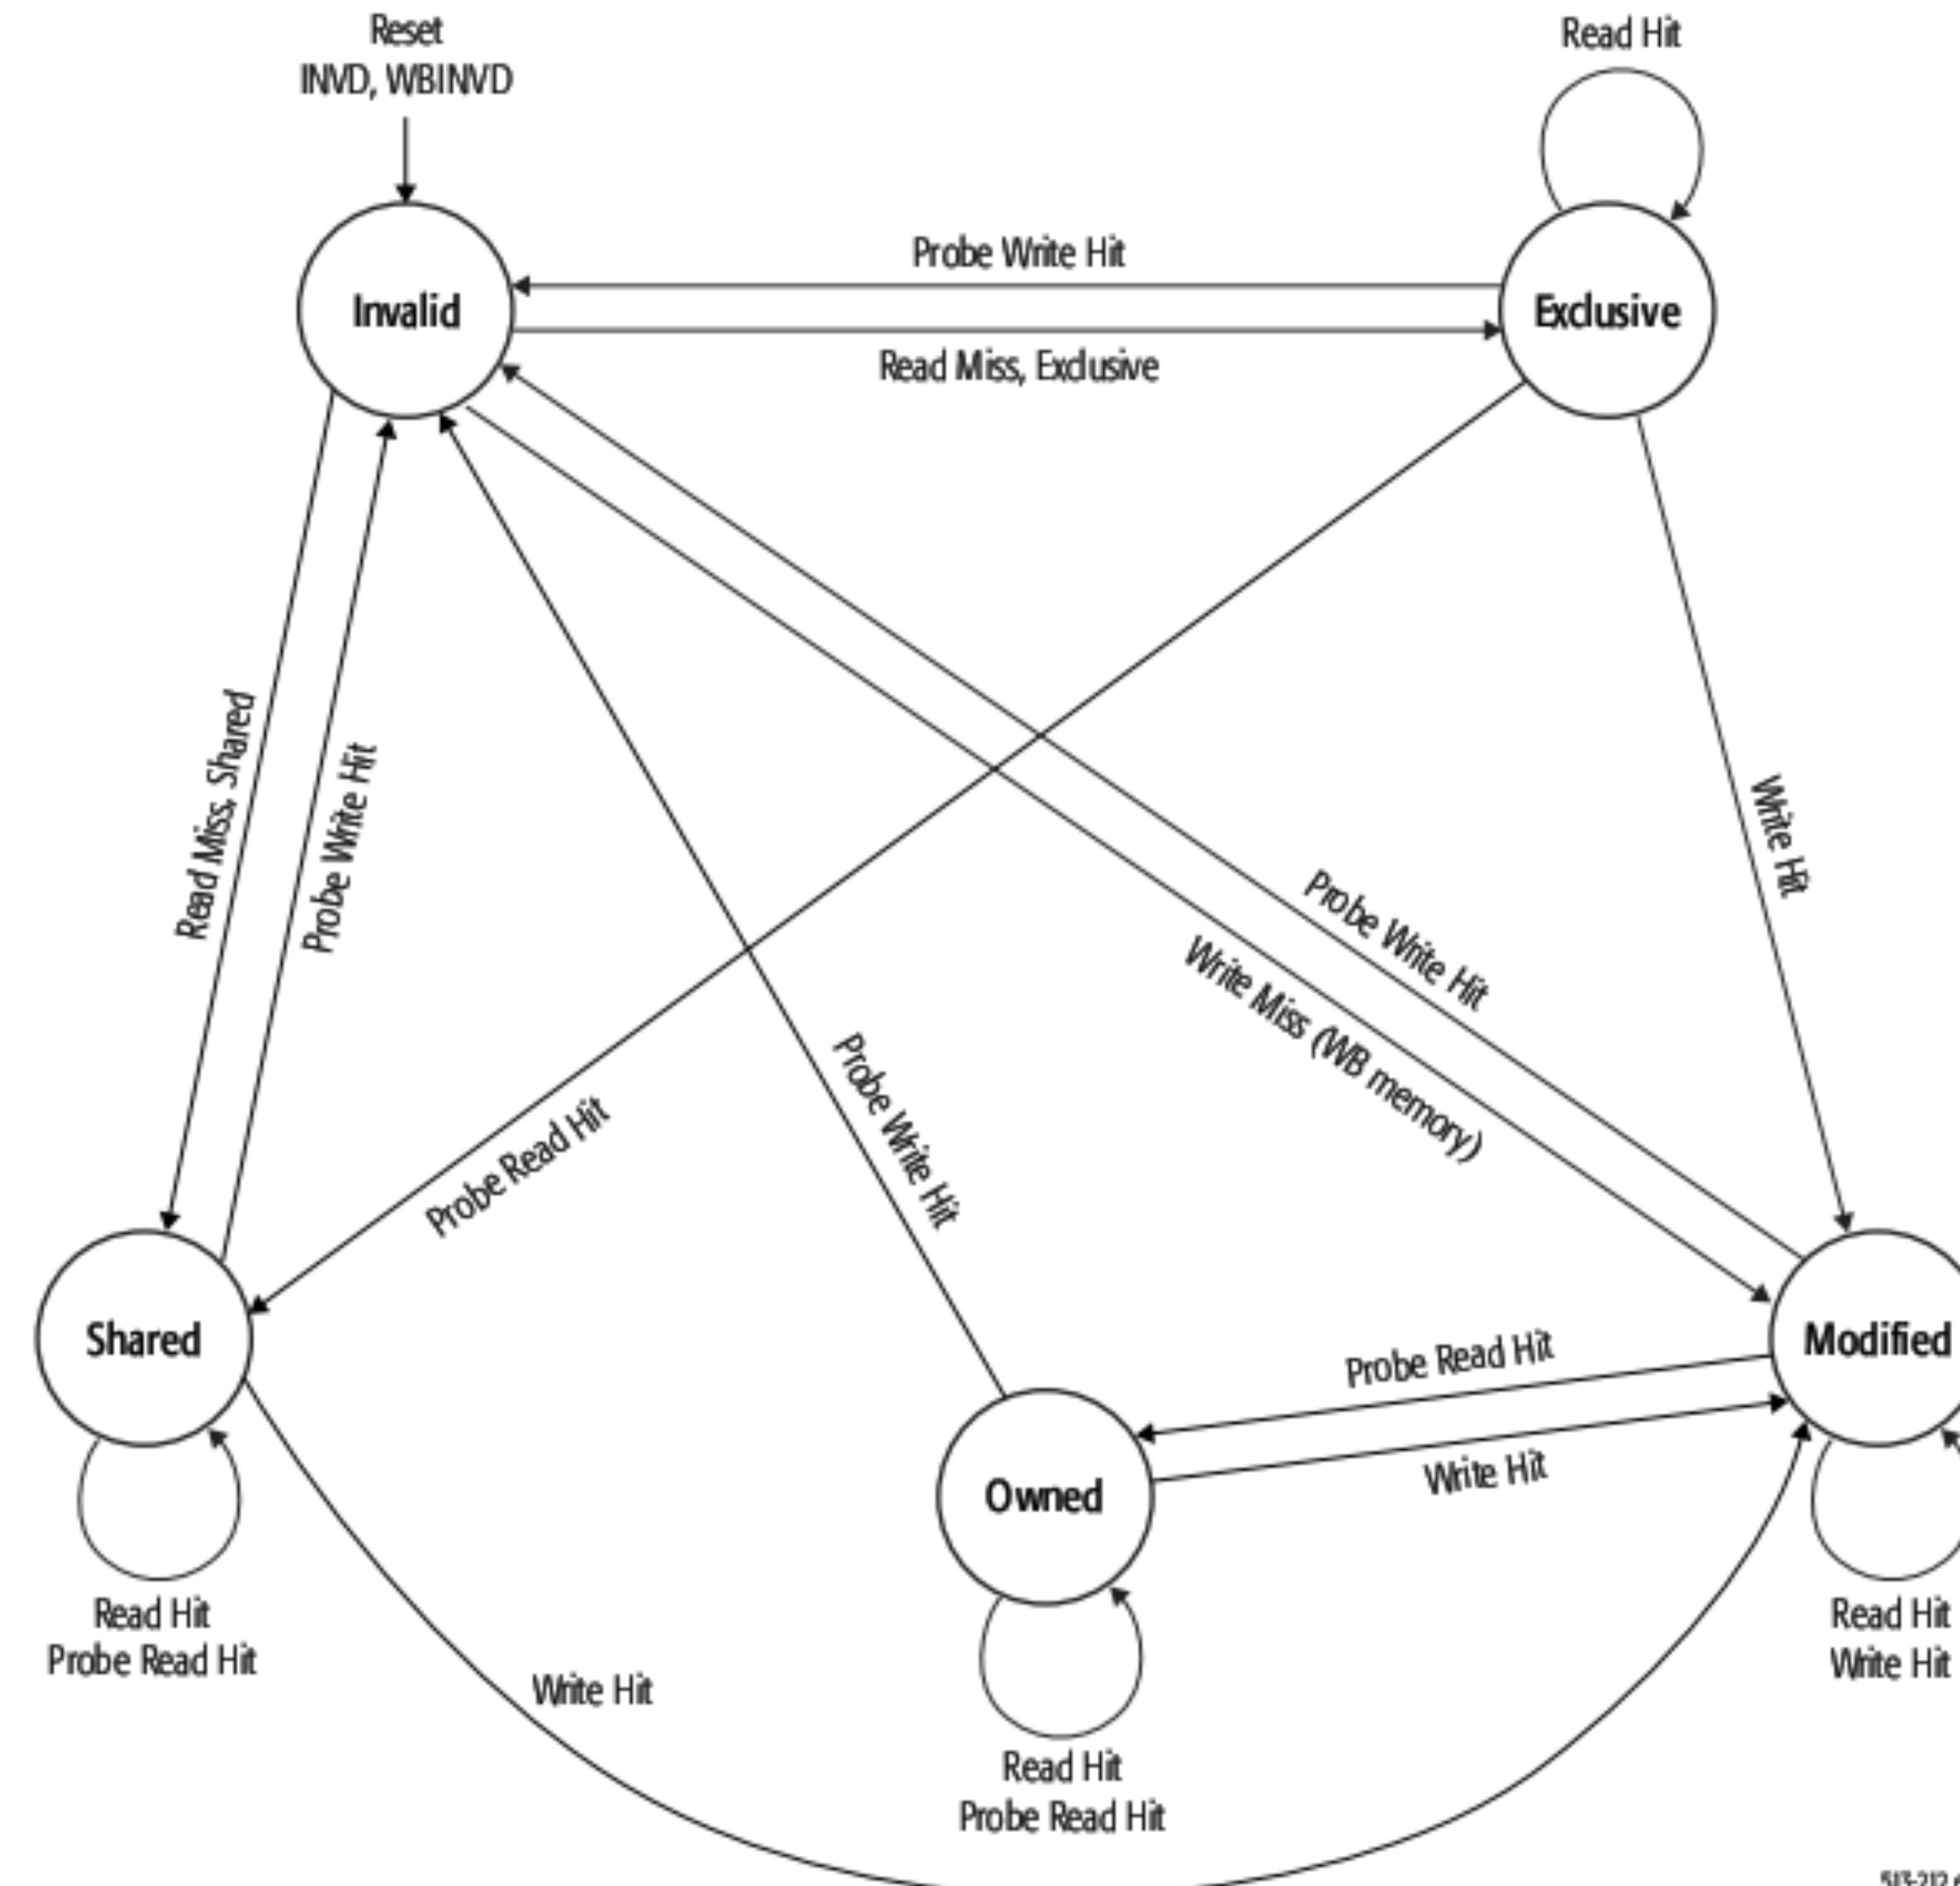
| Valid | Dirty | State |
|-------|-------|-------|
| 1     | 1     | O     |

| Valid | Dirty | State |
|-------|-------|-------|
| 1     | 0     | S     |



0xC0FFEE = True

0xC0FFEE = True



513-212.eps

# The Bits Needed

| Valid | Dirty | Shared | State     |
|-------|-------|--------|-----------|
| 0     | 0     | 0      | Invalid   |
| 0     | 0     | 1      | Invalid   |
| 0     | 1     | 0      | Invalid   |
| 0     | 1     | 1      | Invalid   |
| 1     | 0     | 0      | Exclusive |
| 1     | 0     | 1      | Shared    |
| 1     | 1     | 0      | Modified  |
| 1     | 1     | 1      | Owned     |

# And Implementing **l****r**/**sc**

- **lr:** Just do a load like normal...
  - Now it will be in the cache in E/S/M/O state...
- **sc:** Before storing, make sure it is still in the cache...
  - If E/M: Just write away...
  - If S/O: Broadcast a request to do the write
    - Need to handle the case if two caches do a write in the exact same cycle...
    - If successful, return success
      - If fail, return failure..
    - Which is why caches for multiprocessors are a huge beast to get right in practice
  - If I: Return failure...

# Well, that's the simple way...

- RISC-V offers some guarantees on **l****r**/**s****c**
  - Result is actually implementing this is even harder!
  - If you do a **l****r**/**s****c** pair with the following properties
    - The **l****r** is the last **l****r** issued
    - 16 instructions total, including recovery code
    - No branches except a back branch on failure of the **s****c**
    - Only arithmetic/logical instructions
  - You will make forward progress if multiple cores are doing this...
    - Allows e.g. more complex reduction operations without locks.

# The Problem: Sharing and False Sharing...

- If you share data between two processor cores:
  - Every time one does a write the other will take a cache miss on the next read or write...
  - New class of cache miss type: **coherency**
    - Goes along next to **compulsory**, **capacity**, and **conflict**
- If you share a cache **line** between two processor cores:
  - Every time one does a write the other will take a cache miss...
  - Even if you are writing to different **parts** of the cache line
- But as long as everyone is reading its all cool
  - And just make sure your writes are to different cache **lines** to minimize the cost of writes

# The Opposite Happens with Multithreaded cores

- If two hardware threads share the same physical core...
  - Both writing to and reading the same data?
    - No problem, doesn't cause Coherence misses
  - Both hardware threads using ***different data?***
    - Now we get a bunch of cache misses (?incoherence?)
    - Why initial hardware multithreading was often a performance loser
    - OS scheduler needs to be aware

# But An Alternate // Programming Paradigm... Communicating Sequential Processes

- OpenMP has ***very restrictive*** parallelism
  - Really only good for parallelizing loops with an optional reduction step
    - And Amdahl's law therefore quickly rears its ugly head
- Raw threads is ***very easy*** to get wrong
  - Deadlock situations
- Enter CSP:
  - A way for different threads to efficiently communicate
  - A good CSP language: Go (golang)
    - Really good for shared-memory multiprocessors

# What is Go

- Language created at Google starting in 2007
  - Primarily by a bunch of old Unix hands: Robert Griesemer, Rob Pike, and Ken Thompson
  - 1.0 released in March 2012
- Language continues to evolve, but a commitment to backwards compatibility (so far)
  - A correct program written today will still work tomorrow
    - I'm looking at you, ***python 3....***
- Mostly C-ish ***looking*** but...
  - Strong typing, no pointer arithmetic, lambdas, interfaces, garbage collection and...
  - Strong emphasis on concurrent computation

# Good Go Resources

- The Go website:
  - <https://golang.org/>
- Especially useful: Effective Go:
  - A cheatsheet of programming idioms. Several example from this lecture stolen from there
  - [https://golang.org/doc/effective\\_go.html](https://golang.org/doc/effective_go.html)
- When searching Google, ask for ***golang***, not go
  - The language may be Go, but golang refers to the language too
- If I'm starting new code and I need to care about performance, scalability, or maintainability, I use Go
  - Only exception is if the problem can't stand a garbage-collector pause (such as an OS, low level drone flight control, etc...), in that case I'll have to learn Rust

# So Think of Go as:

- C's general structure & concepts
  - But with implied ;s and a garbage collector
- A better typing system with interfaces, slices, and maps
  - No class inheritance, however
- Much more symmetric functions
  - Can return multiple values
- Scheme-like lexical scope
  - Lambdas and interior function declarations
- Communicating Synchronous Processes (CSP) concurrency
  - Multiple things at once in the same shared memory space:  
quite suitable for MIMD

# Coroutines, err, goroutine

- Conceptually, a goroutine is just a thread...
  - `go fubar() // Executes fubar as a new coroutine`
- But in practice it is designed to be ***much*** lighter weight
  - Threads (e.g. in OpenMP) relatively expensive to create:  
Operating System involvement is never cheap as ***any*** call to the OS is effectively an interrupt
- Go's runtime instead pre-creates a series of threads
  - And then schedules the active coroutines itself to the available threads
- Result is goroutines are ***cheap***
  - It is only slightly more expensive than a plain function call:  
new goroutine just has a small independent stack  
context switching between goroutines is very cheap

# Channels

- Channels are the primary synchronization mechanism
  - You have locks but why bother?
  - A typed and (optionally) buffered communication channel
  - `c := make(chan int)`
  - `e := make(chan *fubar, 100)`
- Writing data to a channel:
  - `c <- 32 // Writing blocks if unbuffered or full`
- Reading from a channel:
  - `var f = <- e // Reading blocks if no data`

# Channels as Synchronization Barriers

- Any writes in the code before writing to the channel complete first
  - `globalA = ...`
  - `d <- 1 // The write to globalA must complete just before this`
- Any reads in the code after reading from the channel do not start until channel-read takes place
  - `<- d`
  - `fubar(globalA) // Won't read globalA before channel read`
- Otherwise, compiler can reorder ***however it wants*** as long as sequential semantics are preserved for the sequential function
  - Go may have removed a lot of ways to shoot yourself in the foot...  
but unless you use channels you will easily blow it off with race conditions

# Without Synchronization Barriers, The Compiler Can Go To Town

- This doesn't work!
  - Compiler can reorder the writes between x and done safely
    - Similar variants also possible
    - Oh, and some processors (notably RISC-V) also allow the processor to recorder the writes to memory too!
- This is one of the two biggest pitfalls of Go:
  - Unless you explicitly synchronize, multiple processes can write in "weird" ways
  - The other is the abysmal error/exception handling mechanism

```
var x : string
var done : bool
func foo() {
    ...
    x = "something"
    done = true
}

func bar() {
    go foo()
    for !done { ... }
    y := x
}
```

# Using goroutines

- For things which may block or wait
  - EG, on input/output, waiting for stuff to happen, etc
  - Just create as many as you want!
    - Its cheap so why not: let the scheduler do useful work when another one is waiting
    - EG, if building a webserver, you launch a goroutine to handle each communication stream
    - This is what python threads are actually good for:  
And then use python Queues like go's channels (without the cool of select...)
- For performance tasks
  - Theoretically create as many as there are CPU cores
    - Otherwise you are wasting resources: but the wastage is low, so you don't necessarily need to Remember, Amdahl's law
  - But it should be much more efficient than OpenMP:
    - Thread creation is vastly more overhead in C/C++ than go

# EG a parallel loop in Go

```
• type Element struct {
    ...
    done chan(bool)
}

func bar(data Element, index int) {
    ...
    // Only if doing fork/join model
    data.done <- true
}

func foo(data []Element) {
    for x, i := range(foo) {
        go bar(x,i)
    }
    // Needed only if I want to do fork/join wait
    // until everything is done...
    for x, _ := range(foo) {
        <- x.done
    }
}
```

# Select

- Select allows you to wait on multiple channels

- ```
select {
    case c <- x:
        x, y = y, x+y
    case d := <- e:
        fmt.Println("Got %v", d)
    default:
        time.Sleep(50 * time.Millisecond)
}
```

- If can write or read to a channel, do so and ***then*** execute the associated case
  - If multiple cases are valid, chose one ***at random***
- If nothing is available, execute default (if any)
  - If no default, just block until you can write or read
  - Default enables non-blocking read & write

# Lots of other features for code correctness

- Compiler is, umm, persnickety
  - It is an error to declare but not use a variable or include but not use a package
- Designed to turn comments into documents
  - Including examples
- Libraries for building example and test routines
- Built in package management integrated into git
- “Single workspace” notion
  - Use common modules to prevent code drift

# So You are Starting A NEW Project?

- There are two modern languages you should first consider: Go and Rust
  - Both have good static typing, high performance compilers, wide support
  - Rust
    - **NO GARBAGE COLLECTOR!**  
Instead a complex memory model that gives malloc()/free() type performance with safety
      - Makes things better on tight embedded systems or problems where you need deterministic memory performance
      - Better exception handling syntactic sugar
  - Go
    - More graceful learning curve
    - Better parallel programming model with select

# And Finally... Nick Applies 61C to Build His New Desktop at the start of COVID

- In Spring 2020: being in the COVID-bunker, Nick decided he needed a new desktop system
  - Since he's always in a fixed place, the laptop's portability matters a lot less...
  - "Not My Money But Is My Money"
    - So sorta care about the budget, but not as much...
- 61C designed performance improvements
  - >2x improvement in sequential performance
    - That I can get this is *yuge*, and partially the reason for the expense
  - ~10x improvement in MIMD performance
    - Lots of cores
  - 10x improvement or more in SIMD performance

# Previous Desktop: Intel Skull Canyon NUC

- Core i7-6770HQ (6th generation)
  - 4 cores/8 threads
    - 2.6 GHz processor core, 6 uOP/cycle issue rate
    - L1I\$ = 32 KiB, 8-way set associative, 64B line size
    - L1D\$ = 32 KiB, 8-way set associative, 64B line size
    - L2\$ = 256 KiB, 4-way set associative, 64B line size
    - L3\$ = 6 MiB, 12-way set associative?, 64B line size
    - **L4\$ = 128 MiB** (DRAM in package/shared with graphics subsystem)
  - Iris Pro Graphics 580 GPU
    - 72 execution units, 1 Teraflop maximum single-precision theoretical output

# Step 1: Number Of Cores?

- Want  $>8$  cores to get the MIMD improvement
  - But don't want to sacrifice sequential performance!
- Decided on AMD over Intel
  - Brand A v Brand X a long subject of debate...
  - At the time, I liked the Zen-2 AMD cores slightly better
    - Today I'd get an Intel "Alder Lake" if building an x86 unless I needed  $>8$  performance cores...
- Zen 2 core
  - 32 KB, 8-way set associative L1 I and D caches
  - 512 KB, **8-way** set associative L2 caches
    - Double associativity is about as good as doubling the cache size!
  - Lots of other microarchitecture changes

# Two choices...

- Ryzen 9 3900x
  - 12 cores, 24 threads, 3.8 GHz
  - 64 MB L3 cache!
- Ryzen 9 3950x
  - 16 cores, 32 threads, 3.5 GHz
  - Still a 64 MB L3 cache
  - And more money!
- So go with the faster/fewer thread version
  - Amdahl's law is a @#)(@#()

# So Far...

- Somewhere between 2x and 3x sequential performance
  - AMD claims a ~17%+ gain in CPI over the Zen 1 core...  
Plus bigger caches...  
Plus much higher clock rate
    - Higher clock rate largely due to the ability to cool much more aggressively: the Skull Canyon was a small form factor so small fan vs huge-glowing-heat-sink
- Close to 10x on MIMD workloads
  - 3x the cores, with the cores being a lot better
- Now to select the SIMD solution...
  - Options are AMD vs Nvidia...  
Another Brand A/Brand N type situation
  - Decided on Nvidia: Both are pretty much the same for gaming, but Nvidia has an edge in use for compute...
  - Got one using the latest architecture:  
Worse on cost/performance but hey...

# Selecting High End Graphics/Compute...

- Nvidia GeForce RTX Super line
  - 2060 Super: 2176 CUDA cores, 6 TFLOP ~\$400
  - 2070 Super: 2560 CUDA cores, 8 TFLOP ~\$500
  - 2080 Super: 3072 CUDA cores, 10 TFLOP ~\$700
- Decided on the 2080...
  - Note: bought before the cryptocurrency space went insane (again), and not upgrading to a 3080...
- Result:
  - 2x-3x sequential
  - close to 10x MIMD
  - easily >10x SIMD
- For each of these:
  - 1/2 the performance from just going to newer technology
  - 1/2 the performance from selecting "High performance" rather than "cost/performance" on the design point